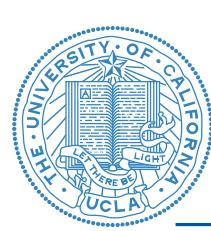




Lecture 9: Regularization and optimization

Announcements:

- HW #3 is due this Friday, February 7. **To submit your Jupyter Notebook, print the notebook to a pdf with your solutions and plots filled in. You must also submit your .py files as pdfs. Please make sure they are readable.**
- Midterm review session: Young Hall CS76 on February 14, from 4pm-6pm.
- The midterm is in just over 2 weeks on February 19, 2025.
 - You are allowed 4 cheat sheets (each an 8.5 x 11 inch paper). You can fill out both sides (8 sides total). You can put whatever you want on these cheat sheets.
 - It will cover material up to and including lecture on February 12, 2025.
 - Past exams are uploaded to Bruin Learn (under “Modules” —> “past exams”).
 - You may bring a calculator to the exam.
 - You **must do the exam in pen.**
- MSOL students will take the exam remotely on Saturday, February 22, from 2-3:50pm. We will send out details to the MSOL students closer to the date.
- We anticipate releasing the project by the end of this weekend (February 9, 2025).



Lecture 9: Regularization and optimization

ECE 239AS - Neural Networks & Deep Learning 2 enrollment.

- The course is offered as 239AS (Lecture 2), not C147B/C247B. As a result, undergraduates cannot enroll unless they have a PTE.
- We are therefore going to do enrollment for this course ONLY via PTE.
- Please fill out this link to sign up for and receive a PTE. Please read the instructions carefully: <https://forms.gle/5vSqKXGeZDTPzcNcA>
- Feel free to send the link to others.
- Note, MEng and MSOL students do not need a PTE. You can enroll as normal.
 - This is because these sections are counted as “Online” and do not count towards the classroom size.
 - I have been notified that even when this course is called “C247B” it will not be considered a “core” class for MSOL, it will still only be an elective. This is a blanket policy for any sequence courses in MSOL.



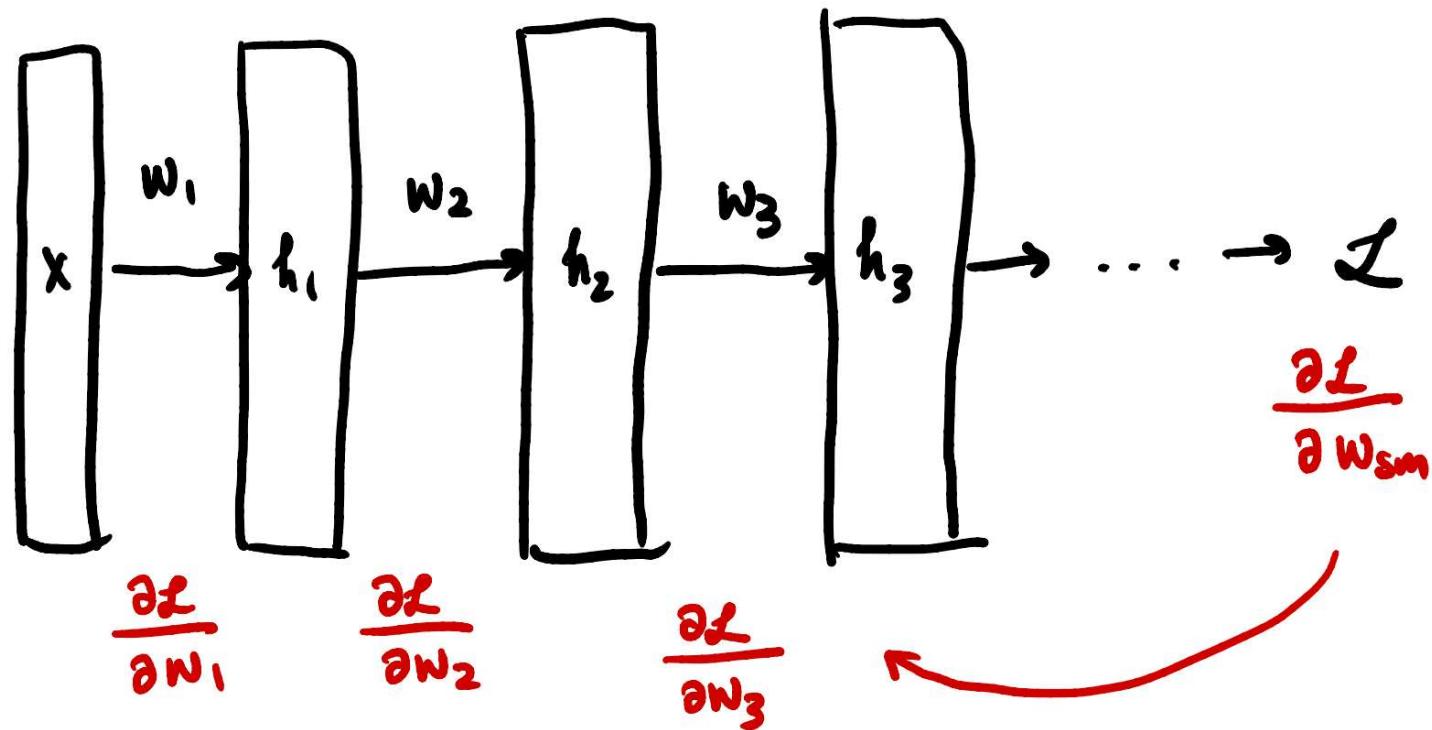
BATCH NORMALIZATION

Avoiding saturation, high variance activations, etc.

Ioffe, Szegedy, 2015

We will motivate batchnorm through internal covariate shift, though new research suggests that batchnorm works primarily due to making the loss surface smoother.

An obstacle to standard training is that the distribution of the inputs to each layer changes as learning occurs in previous layers. As a result, the unit activations can be very variable. Another consideration is that when we do gradient descent, we're calculating how to update each parameter *assuming the other layers don't change*. But these layers may change drastically.



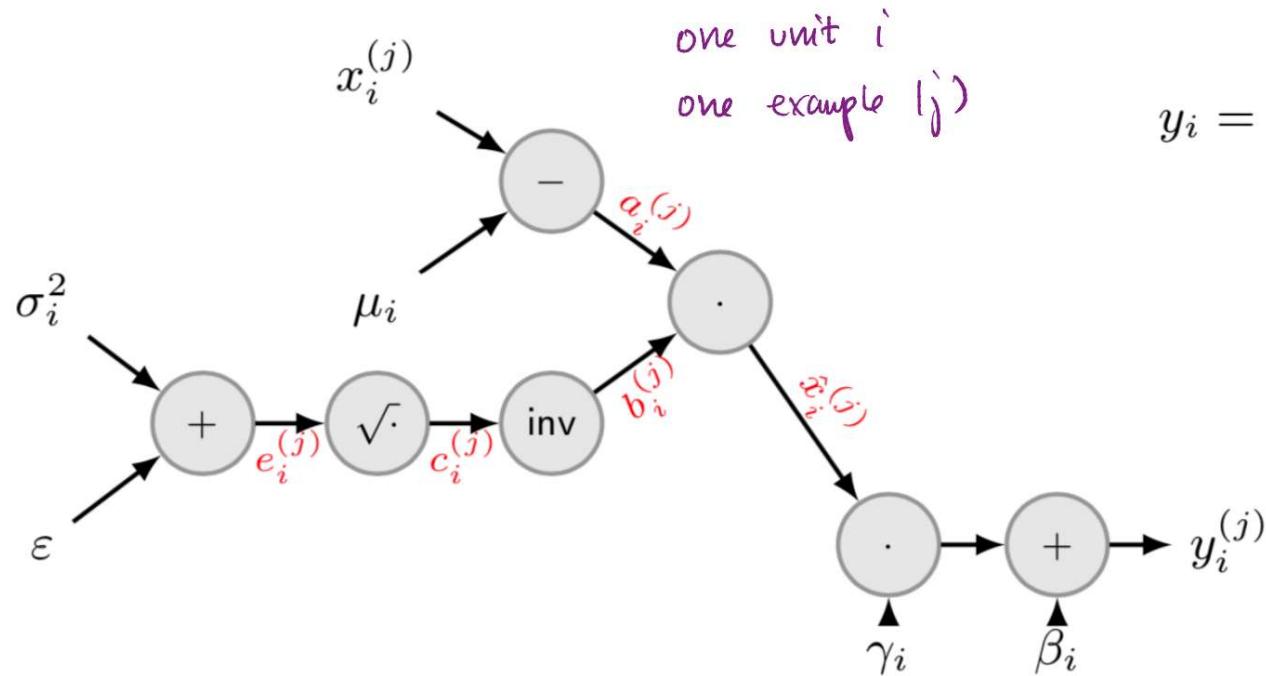


Batch normalization

Batch normalization computational graph

$$\hat{x}_i = \frac{x_i - \mu_i}{\sqrt{\sigma_i^2 + \varepsilon}}$$

$$y_i = \gamma_i \hat{x}_i + \beta_i$$



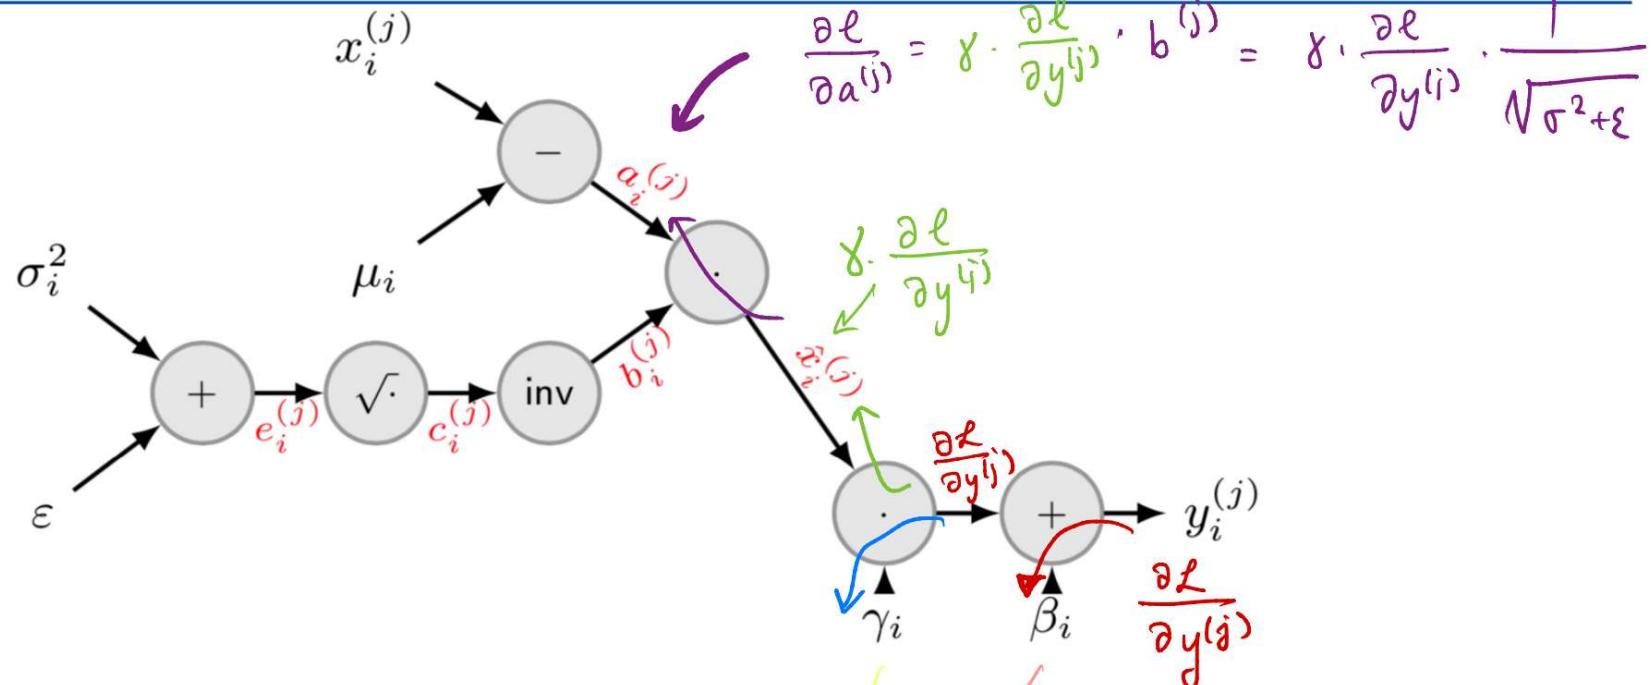
Gradients from backprop on next page.

We will drop the subscript i for convenience.



Batch normalization

$$b^{(j)} = \frac{1}{\sqrt{\sigma^2 + \epsilon}}$$

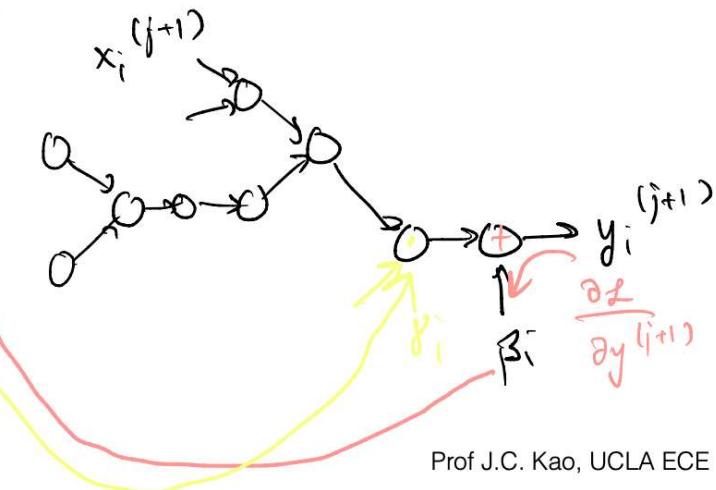


$$\left| \frac{\partial \ell}{\partial \beta} = \sum_{j=1}^m \frac{\partial \ell}{\partial y^{(j)}} \right.$$

$$\left| \frac{\partial \ell}{\partial \gamma} = \sum_{j=1}^m \frac{\partial \ell}{\partial y^{(j)}} \hat{x}^{(j)} \right.$$

$$\left| \frac{\partial \ell}{\partial \hat{x}^{(j)}} = \frac{\partial \ell}{\partial y^{(j)}} \gamma \right.$$

$$\left| \frac{\partial \ell}{\partial a^{(j)}} = \frac{1}{\sqrt{\sigma^2 + \epsilon}} \frac{\partial \ell}{\partial \hat{x}^{(j)}} \right.$$





$$\sigma^2 = \frac{1}{m} \sum_{j=1}^m (x^{(j)} - \mu)^2$$

$$\frac{\partial \sigma^2}{\partial \mu} = -\frac{2}{m} \sum_{j=1}^m (x^{(j)} - \mu) x_i^{(j)}$$

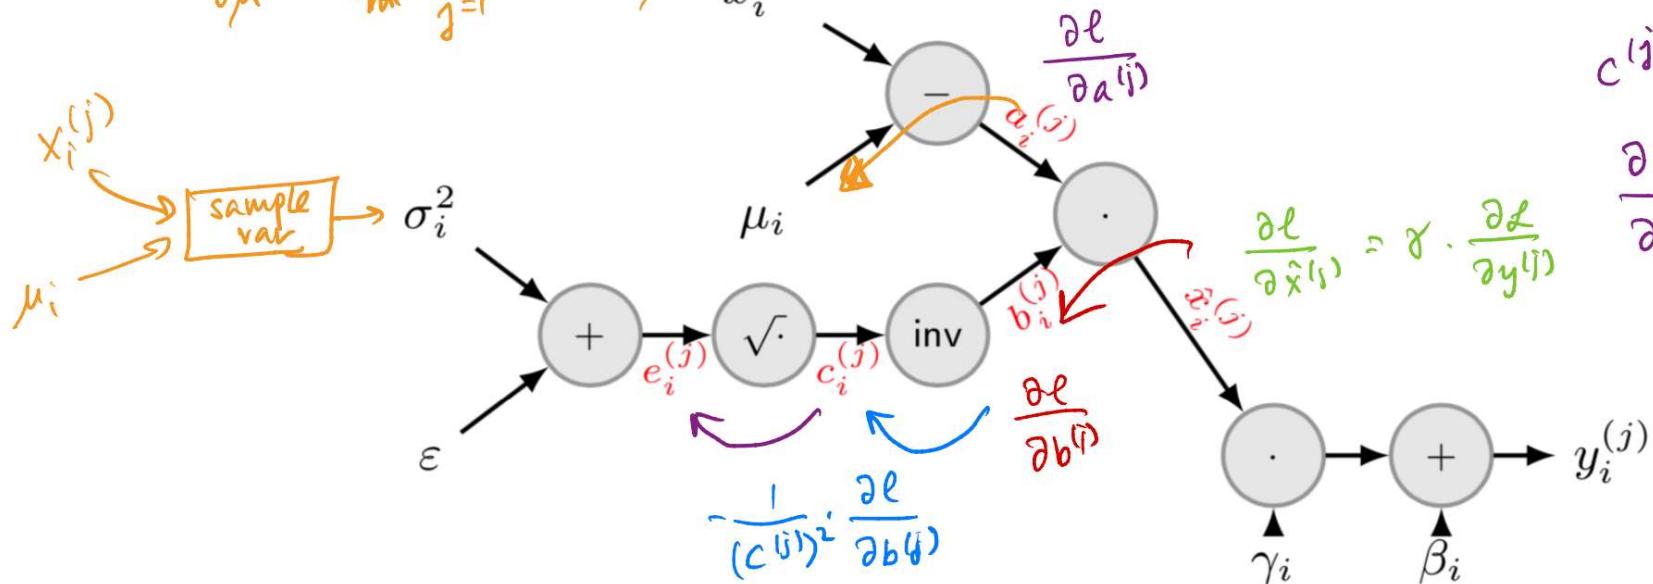
Batch normalization

$$b^{(j)} = \frac{1}{c^{(j)}}$$

$$\frac{\partial b^{(j)}}{\partial c^{(j)}} = -\frac{1}{(c^{(j)})^2}$$

$$c^{(j)} = \sqrt{e^{(j)}}$$

$$\frac{\partial c^{(j)}}{\partial e^{(j)}} = \frac{1}{2\sqrt{e^{(j)}}}$$



$$\frac{\partial \ell}{\partial \mu} = -\frac{1}{\sqrt{\sigma^2 + \epsilon}} \sum_{j=1}^m \frac{\partial \ell}{\partial \hat{x}^{(j)}} - \left(\frac{2}{m} \sum_{j=1}^m (x^{(j)} - \mu) \right) \frac{\partial \ell}{\partial \sigma^2}$$

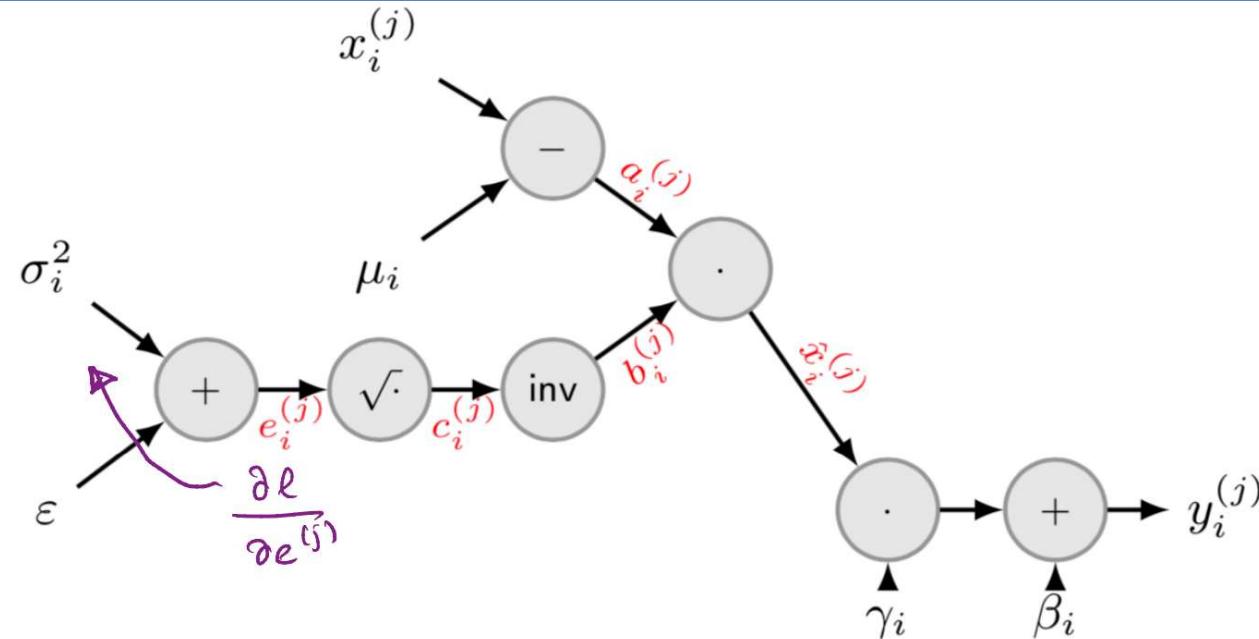
Red: $\frac{\partial \ell}{\partial b^{(j)}} = (x^{(j)} - \mu) \frac{\partial \ell}{\partial \hat{x}^{(j)}}$

Blue: $\frac{\partial \ell}{\partial c^{(j)}} = -\frac{1}{\sigma^2 + \epsilon} (x^{(j)} - \mu) \frac{\partial \ell}{\partial \hat{x}^{(j)}}$

Purple: $\frac{\partial \ell}{\partial e^{(j)}} = -\frac{1}{2} \frac{1}{(\sigma^2 + \epsilon)^{3/2}} (x^{(j)} - \mu) \frac{\partial \ell}{\partial \hat{x}^{(j)}}$



Batch normalization

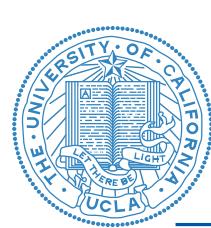


$$\frac{\partial \ell}{\partial a^{(j)}} = \frac{1}{\sqrt{\sigma^2 + \epsilon}} \frac{\partial \ell}{\partial \hat{x}^{(j)}}$$

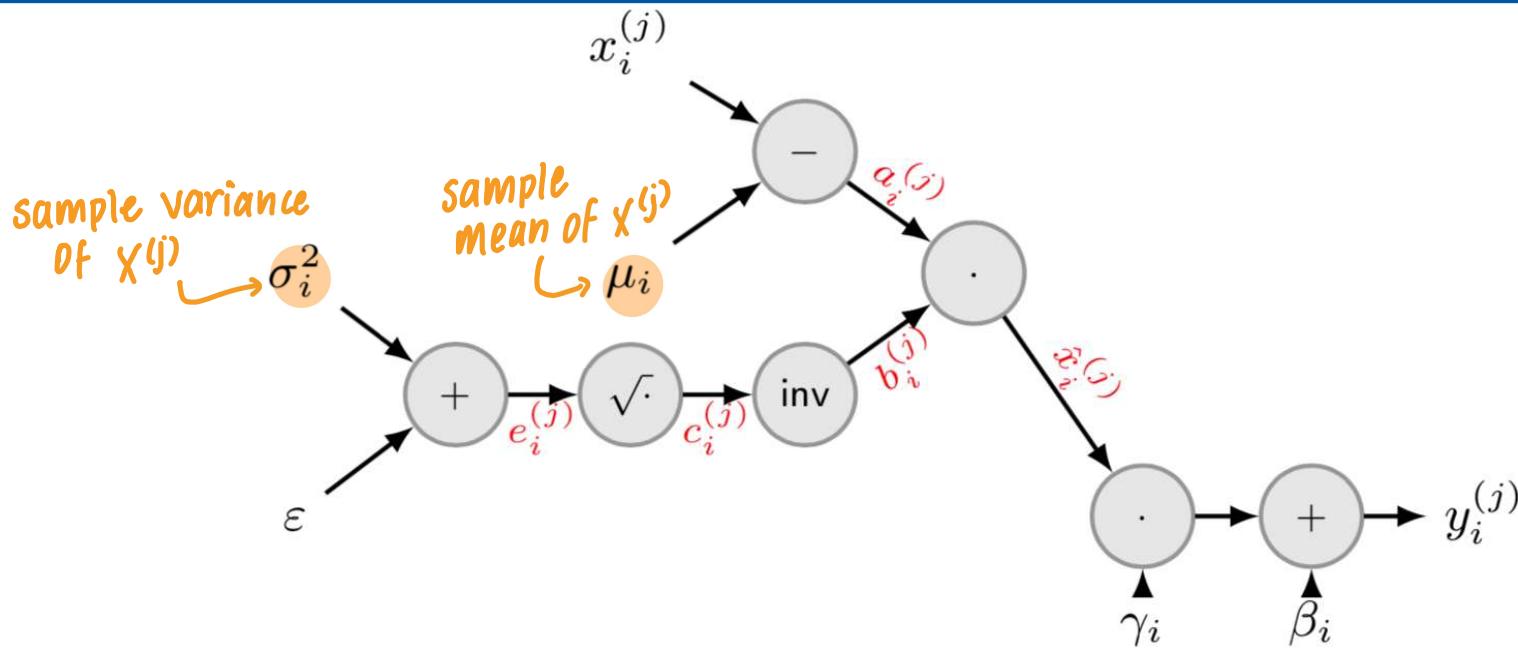
$$\frac{\partial \ell}{\partial e^{(j)}} = -\frac{1}{2} \frac{1}{(\sigma^2 + \epsilon)^{3/2}} (x^{(j)} - \mu) \frac{\partial \ell}{\partial \hat{x}^{(j)}}$$

$$\begin{aligned} \frac{\partial \ell}{\partial \sigma^2} &= \sum_{j=1}^m \frac{\partial \ell}{\partial e^{(j)}} \\ &= \sum_{j=1}^m -\frac{1}{2} \frac{1}{(\sigma^2 + \epsilon)^{3/2}} (x^{(j)} - \mu) \frac{\partial \ell}{\partial \hat{x}^{(j)}} \end{aligned}$$





Batch normalization

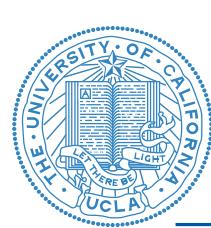


$$\frac{\partial \ell}{\partial a^{(j)}} = \frac{1}{\sqrt{\sigma^2 + \epsilon}} \frac{\partial \ell}{\partial \hat{x}^{(j)}}$$

$$\frac{\partial \ell}{\partial e^{(j)}} = -\frac{1}{2} \frac{1}{(\sigma^2 + \epsilon)^{3/2}} (x^{(j)} - \mu) \frac{\partial \ell}{\partial \hat{x}^{(j)}}$$

$$\begin{aligned} \frac{\partial \ell}{\partial x^{(j)}} &= \frac{\partial \ell}{\partial a^{(j)}} + \underbrace{\frac{\partial \sigma^2}{\partial x^{(j)}} \frac{\partial \ell}{\partial \sigma^2}}_{\text{from prior slide}} + \underbrace{\frac{\partial \mu}{\partial x^{(j)}} \frac{\partial \ell}{\partial \mu}}_{\text{from prior slide}} \\ &= \frac{1}{\sqrt{\sigma^2 + \epsilon}} \frac{\partial \ell}{\partial \hat{x}^{(j)}} + \underbrace{\frac{2(x^{(j)} - \mu)}{m} \frac{\partial \ell}{\partial \sigma^2}}_{\text{from prior slide}} + \underbrace{\frac{1}{m} \frac{\partial \ell}{\partial \mu}}_{\text{from prior slide}} \end{aligned}$$

$\left\{ \begin{array}{l} \mu = \frac{1}{m} \sum_{j=1}^m x^{(j)} \\ \frac{d\mu}{dx^{(j)}} = \frac{1}{m} \end{array} \right.$



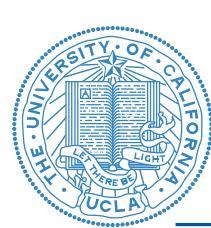
Batch normalization

Batch normalization layer

The batch normalization layer is typically placed right before the nonlinear activation. Hence, a layer of a neural network may look like:

$$\mathbf{h}_i = f(\text{batch-norm}(\mathbf{W}_i \mathbf{h}_{i-1} + \mathbf{b}_i))$$





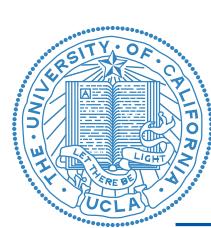
Batch normalization

Empirically, batch normalization allows higher learning rates to be used.

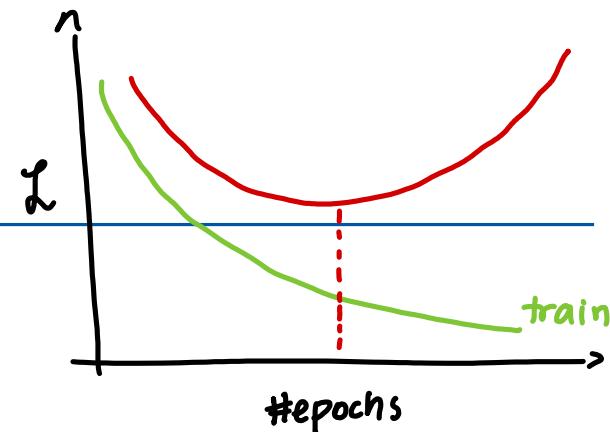
↓
prior to batch norm,
chose smaller learning
rate to mitigate internal
covariance shift.

↳ Batch norm makes
model less sensitive to
initialization.
(more robust)





What is regularization?



Regularizations

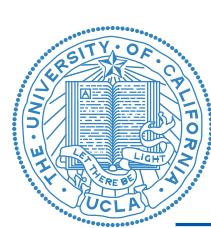
Regularizations are used to improve model generalization. Goodfellow, Bengio, and Courville define regularization in the following way (Deep Learning, p. 221):

[Regularization is] any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.

In this manner, regularization is used to **improve the generalizability** of the model. Other intuitions:

- Regularization tends to increase the estimator bias while reducing the estimator variance.
- Regularization can be seen as a way to prevent overfitting.
- A common problem is in picking the model size and complexity. It may be appropriate to simply choose a large model that is regularized appropriately.





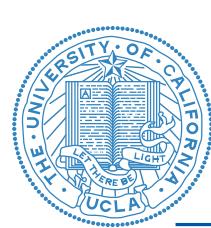
A simple example of regularization

A simple example: stopping early

One straightforward (and popular) way to regularize is to constantly evaluate the training and validation loss on each training iteration, and return the model with the lowest validation error.

- Requires caching the lowest validation error model.
- Training will stop when after a pre-specified number of iterations, no model has decreased the validation error.
- The number of training steps can be thought of as another hyperparameter.
- Validation set error can be evaluated in parallel to training.
- It doesn't require changing the model or cost function.
- The following is beyond the scope of the class – but in case curious, early stopping can be seen as a form of L^2 regularization (to be discussed in the next slides). See Goodfellow et al., *Deep Learning*, p. 242-5 for an indepth discussion.





A regularization familiar to you all: parameter norm penalties

You guys have already implemented some regularization on the HWs through parameter norm penalties.

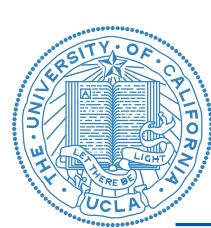
These are not specific to neural networks. These are commonly used, e.g., even in linear regression, where specific penalty norms have their own names (e.g., Tikhonov regularization / ridge regression).

$$\alpha \|\mathbf{w}\|_F$$

Froebinius norm.

↓
hyperparameter
that denotes how
much you care
about regularization.





Parameter norm penalties

Regularization via parameter norm penalties

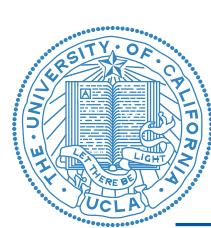
A common (and simple to implement) type of regularization is to modify the cost function with a *parameter norm penalty*. This penalty is typically denoted as $\Omega(\theta)$ and results in a new cost function of the form:

$$J(\theta; \mathbf{X}, \mathbf{y}) + \alpha\Omega(\theta)$$

with $\alpha \geq 0$. A few things to note:

- α is a hyperparameter that weights the contribution of the norm penalty. When $\alpha = 0$, there is no regularization. When $\alpha \rightarrow \infty$, the cost function is irrelevant and the model will set the parameters to minimize $\Omega(\theta)$.
- The choice of α can strongly affect generalization performance.





L2 regularization

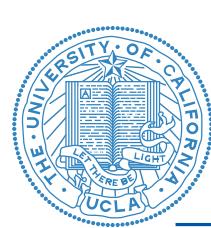
L^2 regularization

A common form of parameter norm regularization is to penalize the size of the weights (L^2 regularization). This is also commonly called “ridge regression” or “Tikhonov regularization.” This promotes models with parameters that are closer to 0 (and hence, colloquially speaking, “simpler”). If \mathbf{w} are the model parameters to be regularized, then L^2 regularization sets:

$$\Omega(\theta) = \frac{1}{2} \mathbf{w}^T \mathbf{w} = \frac{1}{2} \|\mathbf{w}\|_2^2$$

Intuitively, to prevent $\Omega(\theta)$ from getting large, L^2 regularization will cause the weights \mathbf{w} to have small norm.





L2 regularization

L^2 regularization (cont)

More formally, when using L^2 regularization, the new cost function is:

$$\tilde{J}(\mathbf{w}; \mathbf{X}, \mathbf{y}) = J(\mathbf{w}; \mathbf{X}, \mathbf{y}) + \frac{\alpha}{2} \mathbf{w}^T \mathbf{w}$$

No param norm penalty:
 $\mathbf{w} \leftarrow \mathbf{w} - \epsilon \nabla_{\mathbf{w}} \tilde{J}$

with corresponding gradient:

$$\nabla_{\mathbf{w}} \tilde{J}(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \alpha \mathbf{w} + \nabla_{\mathbf{w}} J(\mathbf{w}; \mathbf{X}, \mathbf{y})$$

↑ self-reg
↑ weights

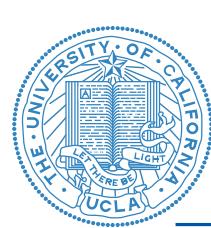
The gradient step is:

$$\mathbf{w} \leftarrow (1 - \epsilon \alpha) \mathbf{w} - \epsilon \nabla_{\mathbf{w}} J(\mathbf{w}; \mathbf{X}, \mathbf{y})$$

This formalizes the intuition that L^2 regularization will shrink the weights, \mathbf{w} , before performing the usual gradient update.

with param norm:
 $\mathbf{w} \leftarrow \mathbf{w} - \epsilon \nabla_{\mathbf{w}} \tilde{J}$
 $\mathbf{w} \leftarrow \mathbf{w} - \epsilon \nabla_{\mathbf{w}} (J + \frac{\alpha}{2} \mathbf{w}^T \mathbf{w})$
 $\mathbf{w} \leftarrow \mathbf{w} - \epsilon \nabla_{\mathbf{w}} J - \epsilon \alpha \mathbf{w}$
$$\boxed{\mathbf{w} \leftarrow (1 - \epsilon \alpha) \mathbf{w} - \epsilon \nabla_{\mathbf{w}} J}$$

weight decay = L2 Reg.
(encourages weights
to be smaller)



Other places you may have seen L2 regularization (**NOT tested**)

Other equivalent statements of L^2 regularization

While we won't discuss these at length in class, it may be worthwhile to work out these equivalences:

- L^2 regularization is equivalent to maximum a-posteriori inference, where the prior on the parameters has a unit Gaussian distribution, i.e.,

$$\mathbf{w} \sim \mathcal{N}(0, \frac{1}{\alpha} \mathbf{I})$$

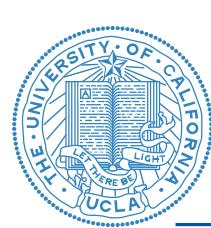
- When performing L^2 regularization, the component of \mathbf{w} aligned with the i th eigenvector of the Hessian is rescaled by a factor

$$\frac{\lambda_i}{\lambda_i + \alpha}$$

- In linear regression, the least squares solution $\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$ becomes:

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

scaling the variance of each input feature. The dimensions of $\mathbf{X}^T \mathbf{X}$ that are large (i.e., high variance) aren't affected as much, while those dimensions where $\mathbf{X}^T \mathbf{X}$ is small are.



L2 regularization

Extensions of L^2 regularization

Other related forms of regularization include:

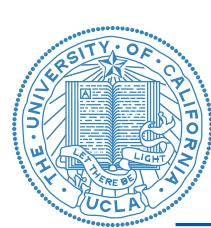
- Instead of a soft constraint that \mathbf{w} be small, one may have prior knowledge that \mathbf{w} is close to some value, \mathbf{b} . Then, the regularizer may take the form:

$$\Omega(\theta) = \|\mathbf{w} - \mathbf{b}\|_2$$

- One may have prior knowledge that two parameters, $\mathbf{w}^{(1)}$ and $\mathbf{w}^{(2)}$, ought be close to each other. Then, the regularizer may take the form:

$$\Omega(\theta) = \left\| \mathbf{w}^{(1)} - \mathbf{w}^{(2)} \right\|_2$$





L1 regularization

L^1 regularization

L^1 regularization defines the parameter norm penalty as:

$$\begin{aligned}\Omega(\theta) &= \|\mathbf{w}\|_1 \\ &= \sum_i |w_i|\end{aligned}$$

many weights
are pushed to 0.
(sparse weights).

gradients
get smaller
closer to 0,
less learning
happens

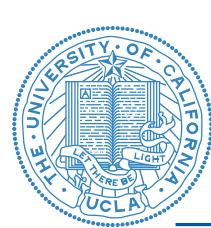


Intuitively, this penalty also causes the weights to be small. However, because the subgradient of $\|\mathbf{w}\|_1$ is $\text{sign}(\mathbf{w})$, the gradient is the same regardless of the size of \mathbf{w} . (Contrast this to L^2 regularization, where the size of \mathbf{w} matters.) Empirically, this typically results in sparse solutions where $w_i = 0$ for several i .

- This may be used for *feature selection*, where features corresponding to zero weights may be discarded.
- L^1 regularization is equivalent to maximum a-posteriori inference where the prior on the parameters has an isotropic Laplace distribution, i.e.,

$$w_i \sim \text{Laplace} \left(0, \frac{1}{\alpha} \right)$$

NOT
TESTED



L1 regularization on the units

Sparse representations

Instead of having sparse parameters (i.e., elements of w being sparse), it may be appropriate to have sparse *representations*. Imagine a hidden layer of activity, $\mathbf{h}^{(i)}$. To achieve a sparse representation, one may set:

$$\Omega(\mathbf{h}^{(i)}) = \|\mathbf{h}^{(i)}\|_1$$

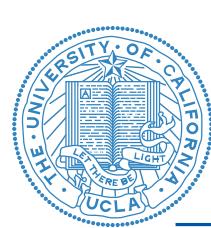




Dataset augmentation

Original image:





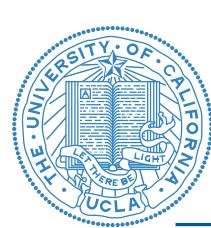
Dataset augmentation

Original image:



Flipped image:





Dataset augmentation

Original image:



Flipped image:



Cropped image:





Dataset augmentation

Original image:



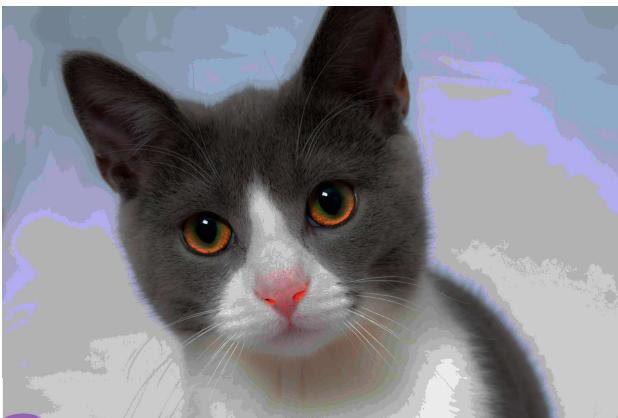
Flipped image:



Cropped image:



Adjust brightness





Dataset augmentation

Original image:



Flipped image:



Cropped image:



Adjust brightness



Lens correction





Dataset augmentation

Original image:



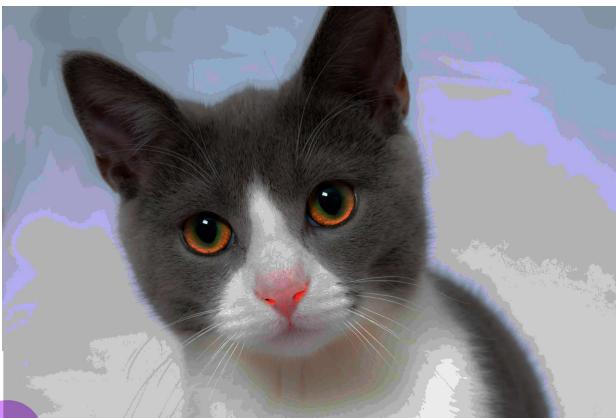
Flipped image:



Cropped image:



Adjust brightness

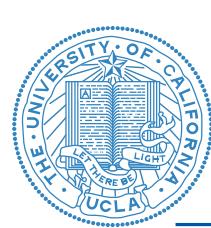


Lens correction



Rotate





Dataset augmentation

There are various heuristics to keeping the input size the same.

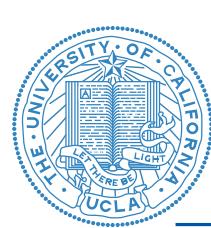
You can be creative for dataset augmentation.

- During testing, we adopted a more aggressive cropping approach than that of Krizhevsky et al. [9]. Specifically, we resized the image to 4 scales where the shorter dimension (height or width) is 256, 288, 320 and 352 respectively, take the left, center and right square of these resized images (in the case of portrait images, we take the top, center and bottom squares). For each square, we then take the 4 corners and the center 224×224 crop as well as the square resized to 224×224 , and their mirrored versions. This leads to $4 \times 3 \times 6 \times 2 = 144$ crops per image. A similar ap-

Number of models	Number of Crops	Cost	Top-5 error	compared to base
1	1	1	10.07%	base
1	10	10	9.15%	-0.92%
1	144	144	7.89%	-2.18%
7	1	7	8.09%	-1.98%
7	10	70	7.62%	-2.45%
7	144	1008	6.67%	-3.45%

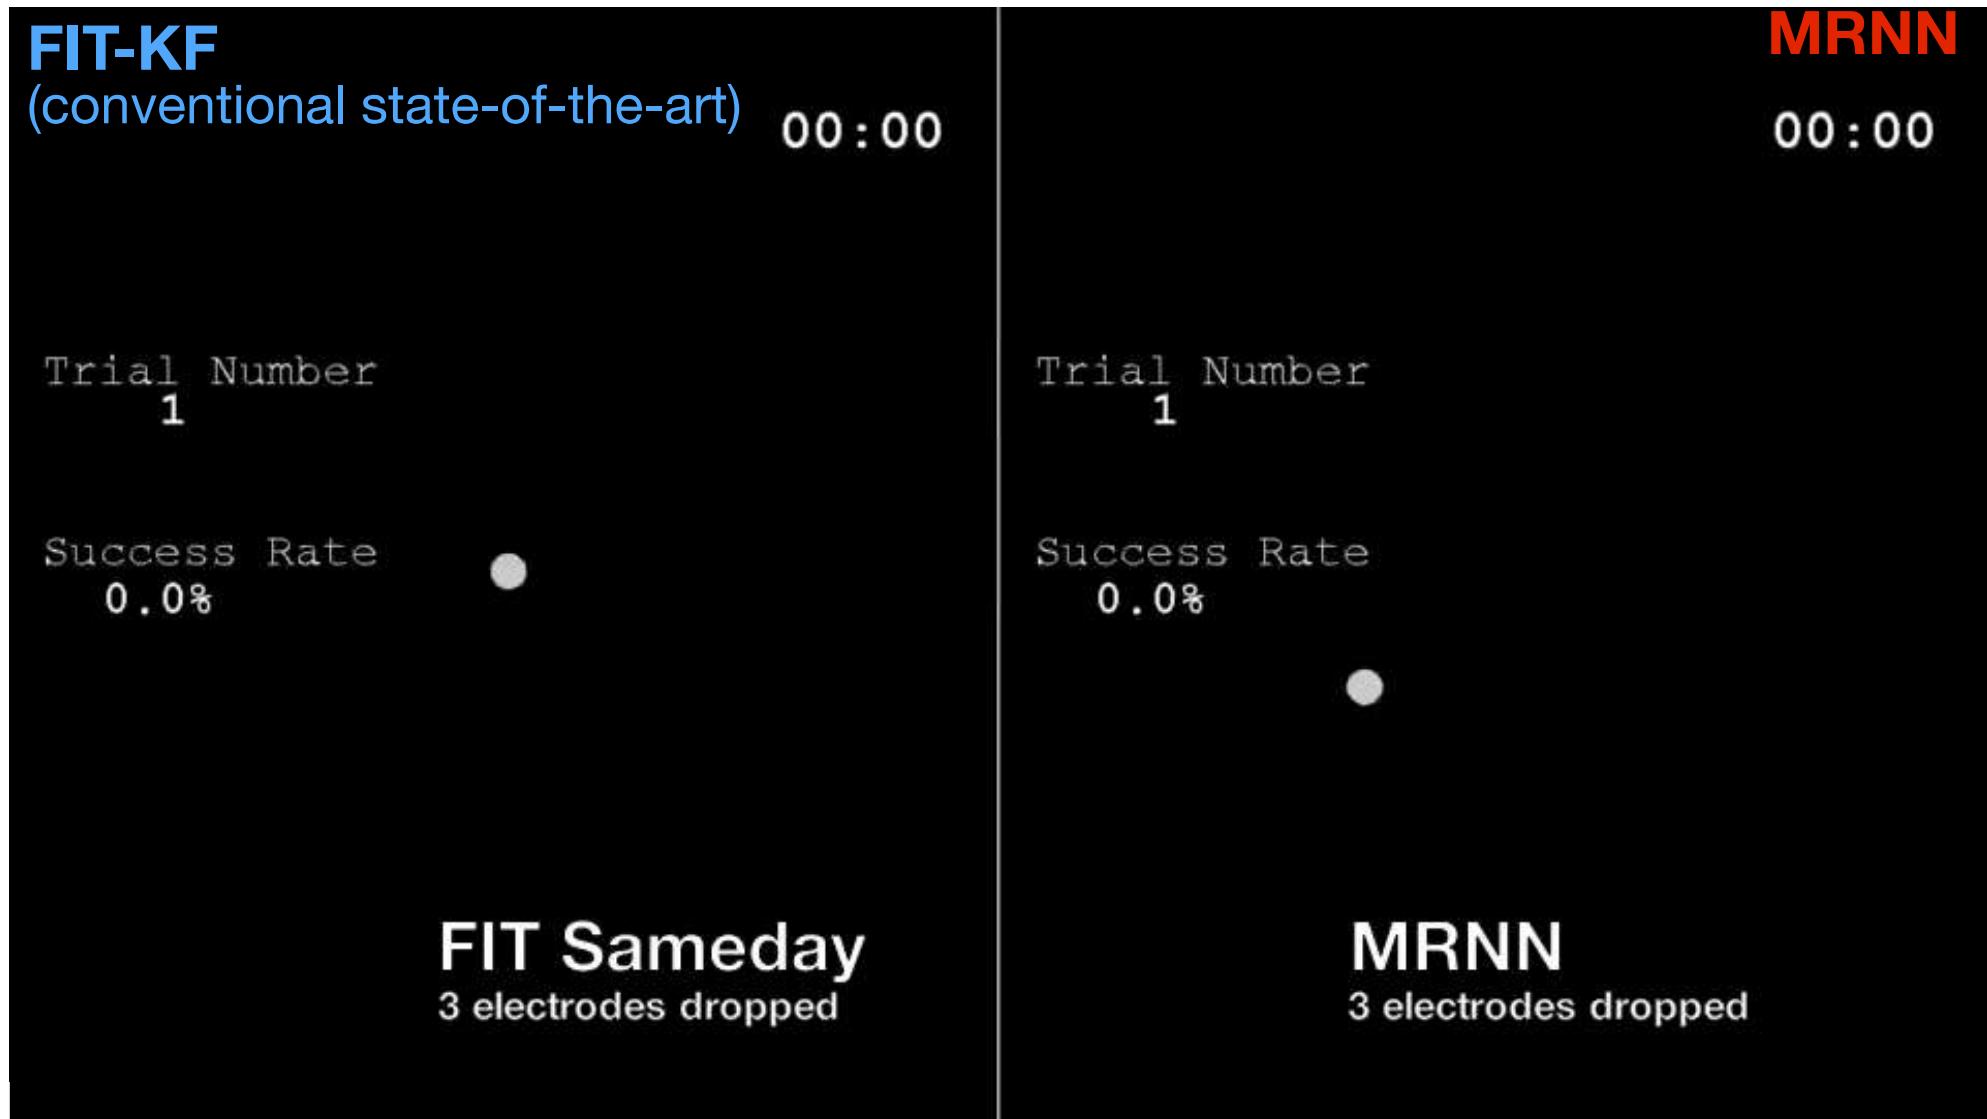
Szegedy et al., CVPR 2015



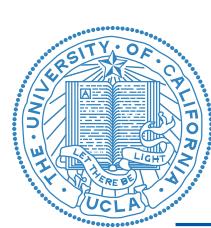


Dataset augmentation

Example from brain-machine interfaces:

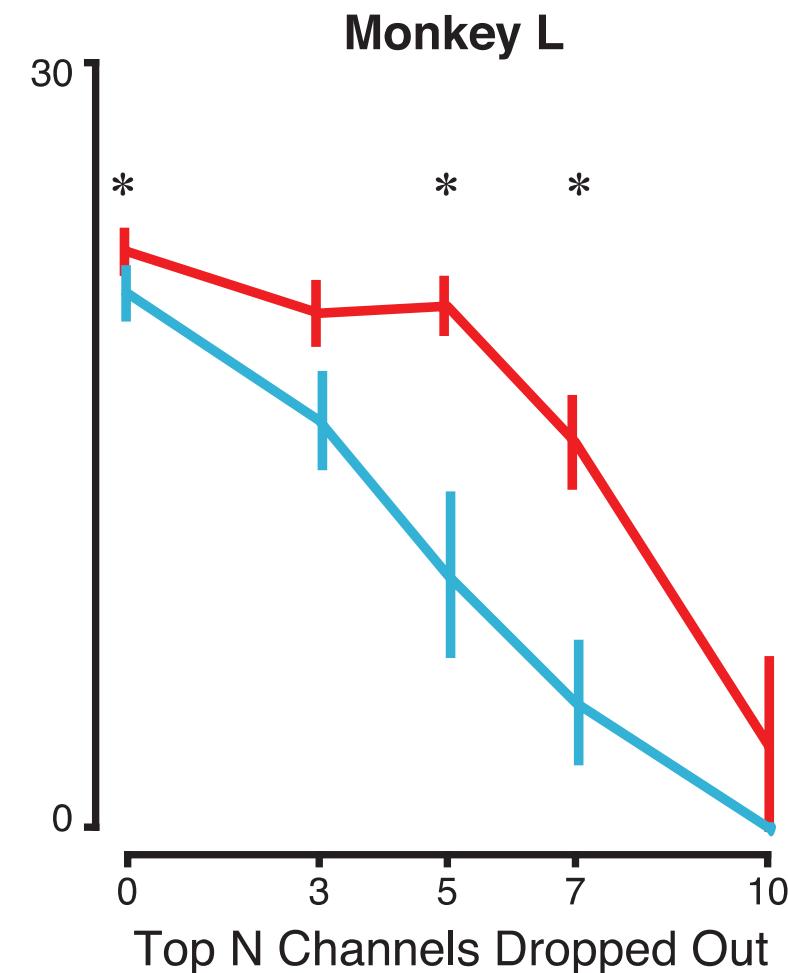
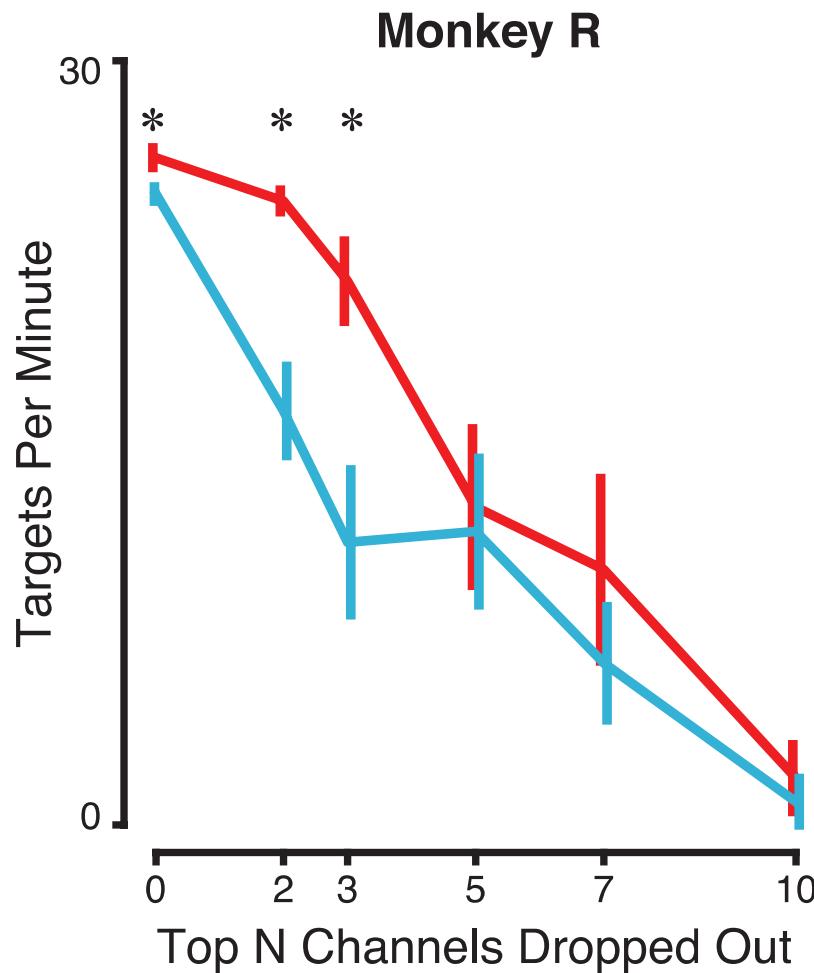


(back-to-back comparisons during the same experiment)



Dataset augmentation

Example from brain-machine interfaces:





Dataset augmentation

Other types of dataset augmentation:

- Inject noise into the network
- Label smoothing

main idea:
injecting some noise
improves generalizability.

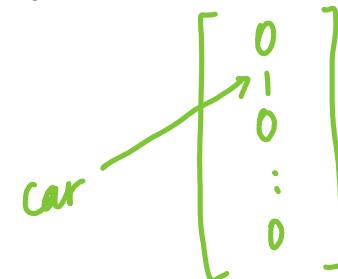
in this case, you are
saying that your outputs
may be noisy.

Smaller batches = noisy gradients = regularization effect.

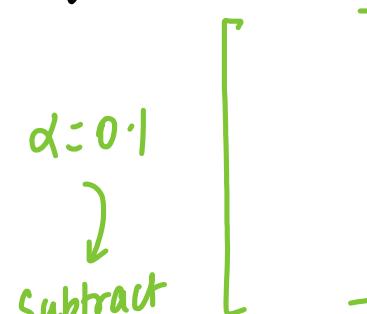
Network	Top-1 Error	Top-5 Error	Cost Bn Ops
GoogLeNet [20]	29%	9.2%	1.5
BN-GoogLeNet	26.8%	-	1.5
BN-Inception [7]	25.2%	7.8	2.0
Inception-v2	23.4%	-	3.8
Inception-v2	23.1%	6.3	3.8
RMSProp	23.1%	6.3	3.8
Inception-v2	22.8%	6.1	3.8
Label Smoothing	22.8%	6.1	3.8
Inception-v2	21.6%	5.8	4.8
Factorized 7×7	21.6%	5.8	4.8
Inception-v2	21.2%	5.6%	4.8
BN-auxiliary			

Szegedy et al., arXiv 2015

one hot representation

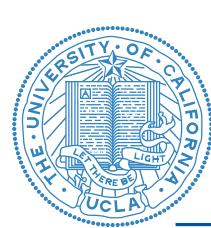


Label smoothing



Subtract this from the correct class, and distribute equally across other classes.

)
acts as a kind of regularization.

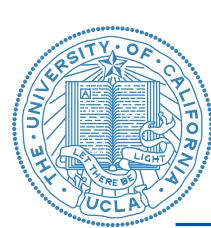


Multitask learning

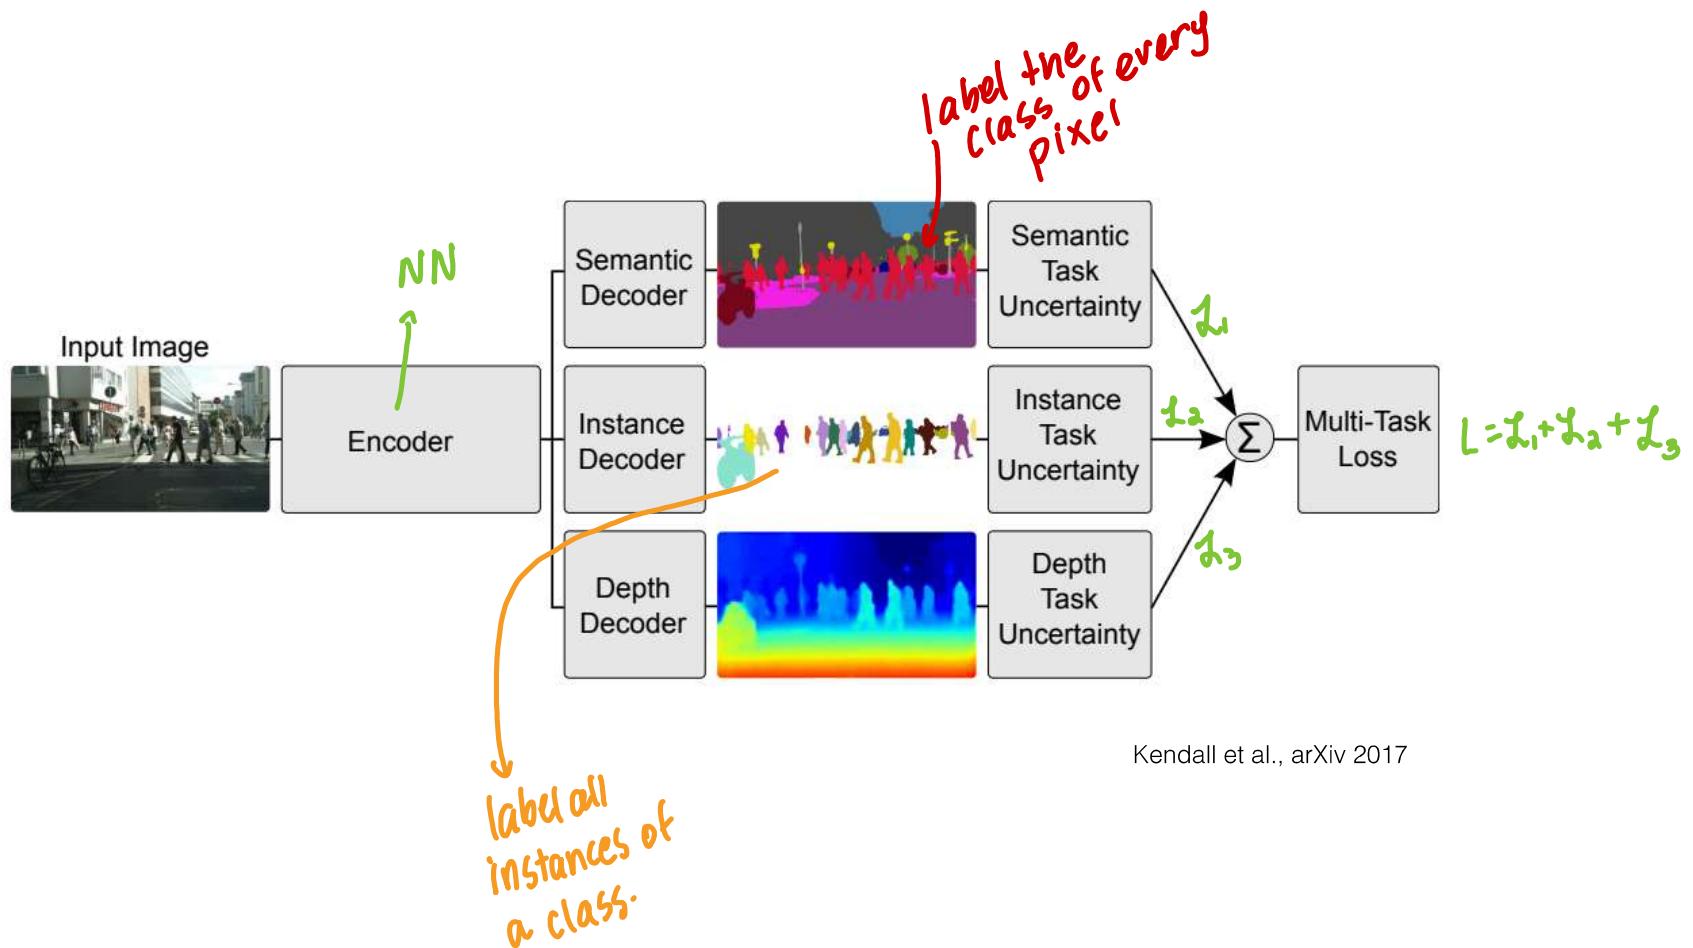
Multitask learning

Another way to improve generalization is by having the model be trained to perform multiple tasks. This represents the prior belief that multiple tasks share common factors to explain variations in the data.



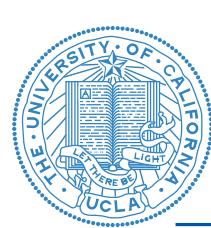


Multitask learning



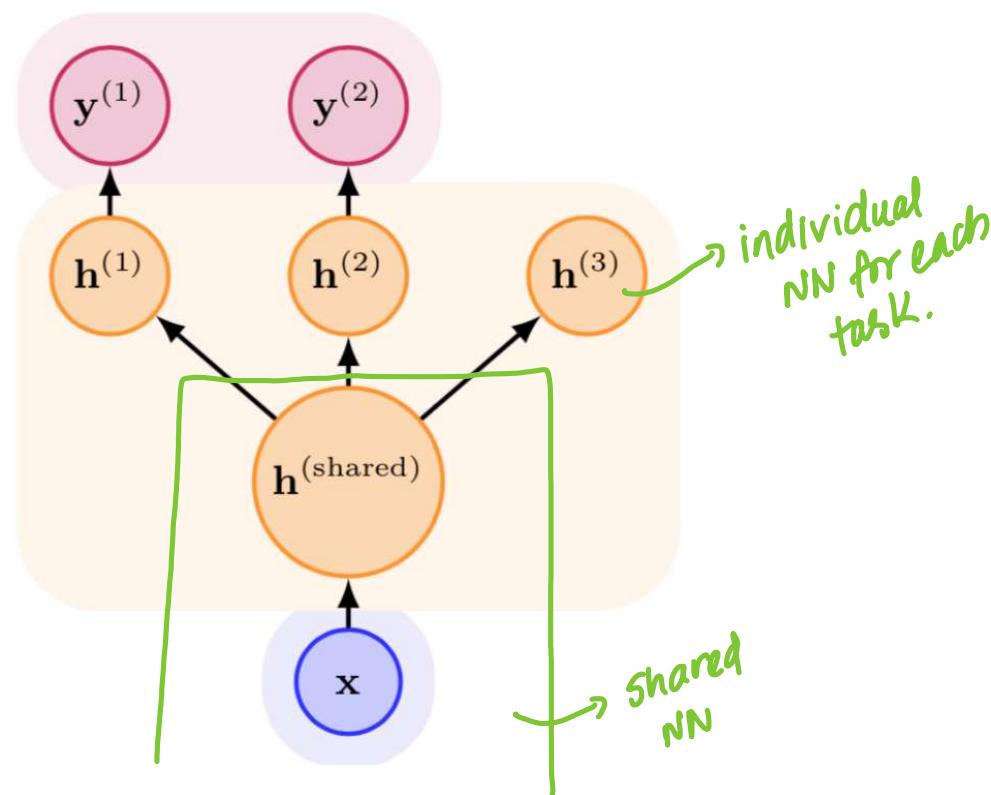
Kendall et al., arXiv 2017

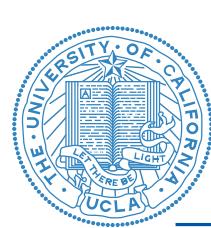




Multitask learning

- The entire model need not be shared across different tasks.
- Here, $\mathbf{h}^{\text{shared}}$ captures common features that are then used by task-specific layers to predict $\mathbf{y}^{(1)}$ and $\mathbf{y}^{(2)}$.
- $\mathbf{h}^{(3)}$ could represent a feature for unsupervised learning.



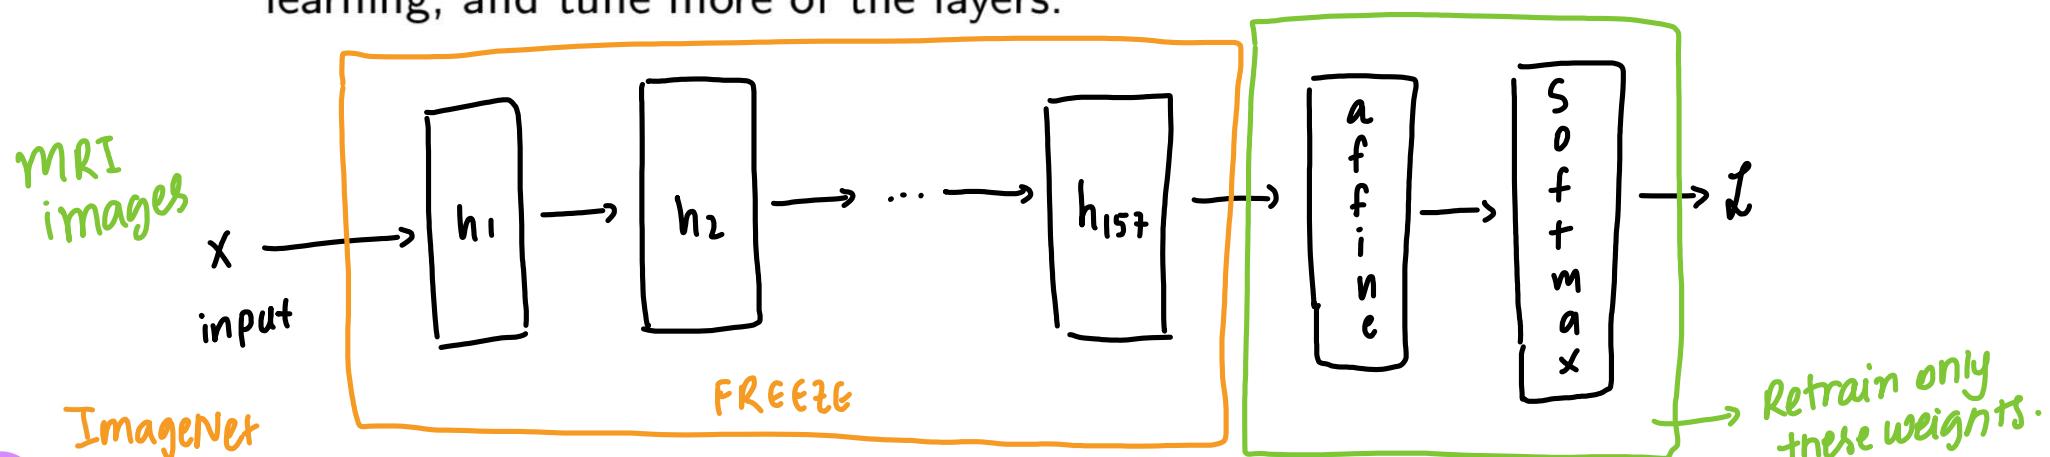


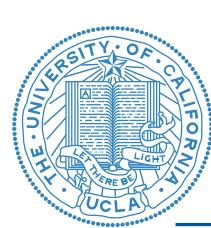
Transfer learning

Transfer learning

We'll discuss this more in the convolutional neural networks lecture, but a related idea is to take neural networks trained in one context and use them in another with little additional training.

- The idea is that if the tasks are similar enough, then the features at later layers of the network ought to be good features for this new task.
- If little training data is available to you, but the tasks are similar, all you may need to do is train a new linear layer at the output of the pre-trained network.
- If more data is available, it may still be a good idea to use transfer learning, and tune more of the layers.





Ensemble methods

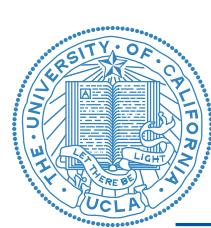
One way to get a boost in performance for very little cognitive work is to use ensemble methods.

The approach is:

1. Train multiple different models
2. Average their results together at test time.

This almost always increases performance by substantial amounts (e.g., a few percentage improvement in testing).





Ensemble methods

- The basic intuition between ensemble methods is that if models are independent, they will usually not all make the same errors on the test set.
- With k independent models, the average model error will decrease by a factor $\frac{1}{k}$. Denoting ϵ_i to be the error of model i on an example, and assuming $\mathbb{E}\epsilon_i = 0$ as well as that the statistics of this error is the same across all models,

$$\mathbb{E}[\epsilon_i, \epsilon_j] = \mathbb{E}(\epsilon_i)\mathbb{E}(\epsilon_j)$$

$$\begin{aligned}\mathbb{E} \left[\left(\frac{1}{k} \sum_{i=1}^k \epsilon_i \right)^2 \right] &= \frac{1}{k^2} \sum_{i=1}^k \mathbb{E} \epsilon_i^2 \quad \text{when errors are independent.} \\ &= \frac{1}{k} \mathbb{E} \epsilon_i^2\end{aligned}$$

- If the models are not independent, it can be shown that:

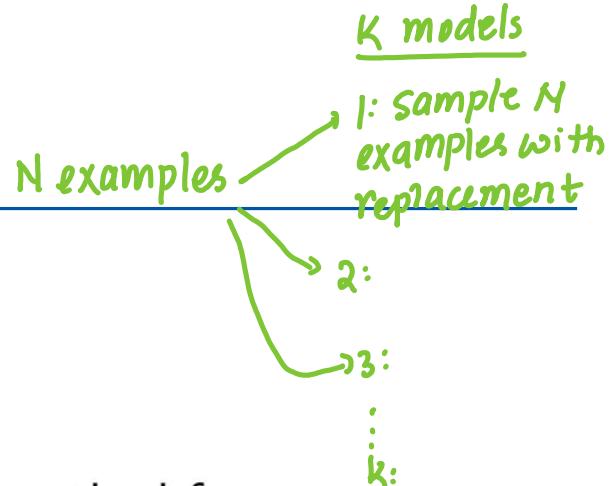
$$\frac{1}{k} \mathbb{E} \epsilon_i^2 + \frac{k-1}{k} \mathbb{E}[\epsilon_i \epsilon_j]$$

which is equal to $\mathbb{E} \epsilon_i^2$ only when the models are perfectly correlated.

NOT
TESTED



Ensemble methods



One implementation of ensemble methods is via bagging.

Bagging

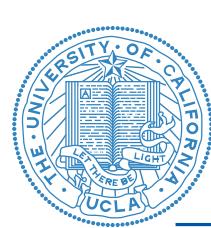
Bagging stands for bootstrap aggregating. It is an ensemble method for regularization. The procedure is as follows:

- Construct k datasets using the bootstrap (i.e., set a data size, N , and draw with replacement from the original dataset to get N samples; do this k times).
- Train k different models using these k datasets.

A few notes:

- In practice, neural networks reach a wide variety of solutions given different initializations, hyperparameters, etc., and so in practice even if they are trained from the same dataset, they tend to produce partially independent errors.
- While model averaging is powerful, it is expensive for neural networks, since the time to train models can be very large.

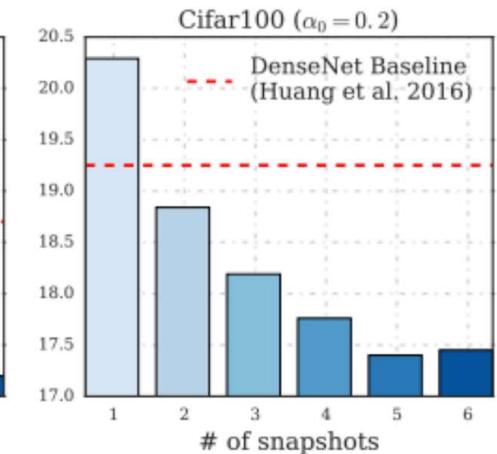
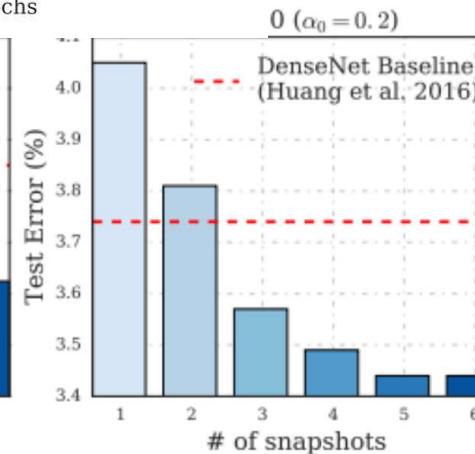
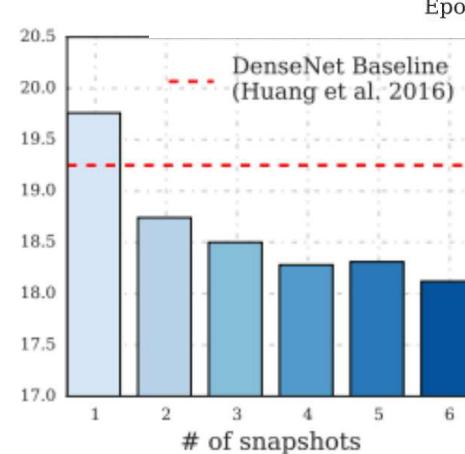
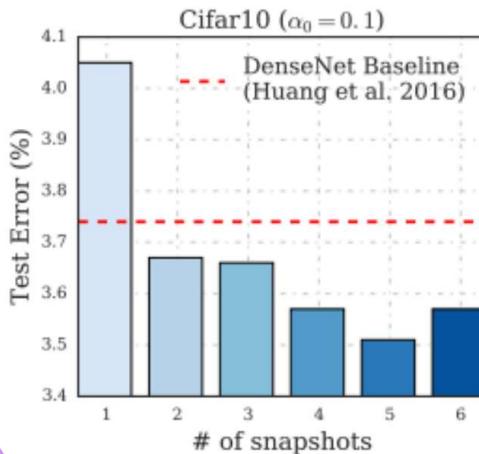
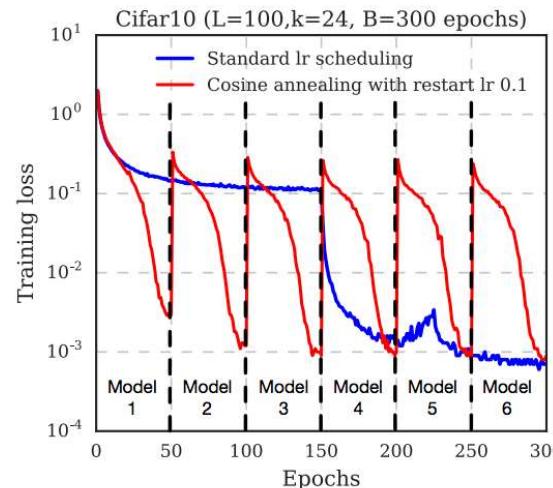


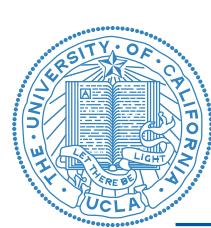


Ensemble methods

Ways to get around computational expense?

You could take snapshots of the model at different local minima and average them together. See Huang, Li et al., ICLR 2017.





Dropout

Dropout

Dropout is a computationally inexpensive yet effective method for generalization. It can be viewed as approximating the bagging procedure for exponentially many models, while only optimizing a single set of parameters. The following steps describe the dropout regularizer:

- On a given training iteration, sample a binary mask (i.e., each element is 0 or 1) for all input and hidden units in the network.
 - The Bernoulli parameter, p , is a hyperparameter.
 - Typical values are that $p = 0.8$ for input units and $p = 0.5$ for hidden units.
- Apply (i.e., multiply) the mask to all units. Then perform the forward pass and backwards pass, and parameter update.
- In *prediction*, multiply each hidden unit by the parameter of its Bernoulli mask, p .

h
100 units

Draw 100 Bernoulli R.Vs.

$w \cdot p \cdot P$

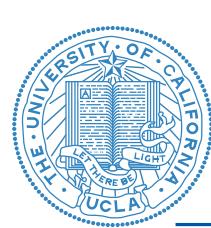
$w \cdot p \cdot 1-p$

$m = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}$

$h \odot m$

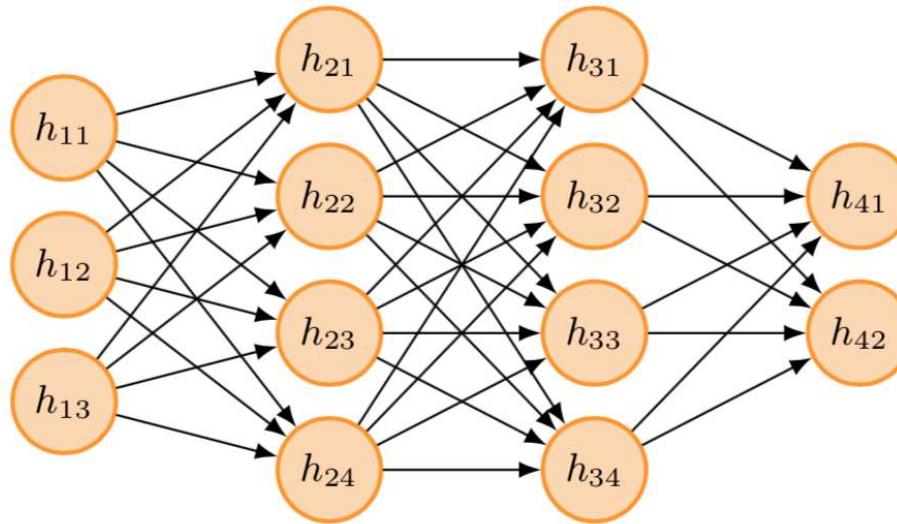
Check:

The diagram shows a vertical vector labeled m with 100 entries. Above the first two entries, arrows point from the text "w.p. P" and "w.p. 1-p". The vector m has entries 1, 0, 1, followed by several intermediate entries, and ending with 0. To the right of the vector, there is handwritten text "h ⊙ m" with an orange checkmark, and below it, the word "check".

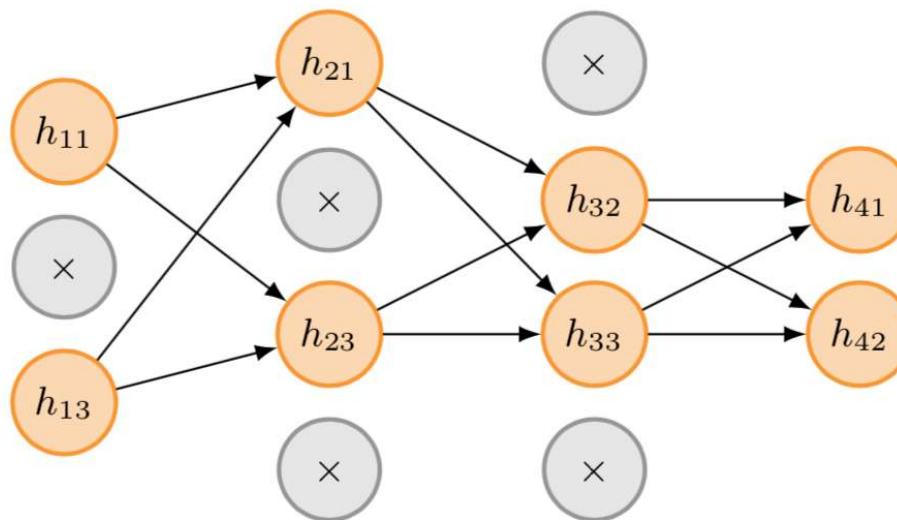


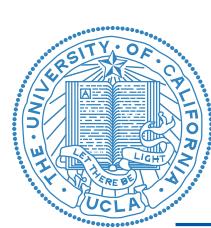
Dropout

Hidden layers of network to be trained



Application of random mask on iteration i





Dropout

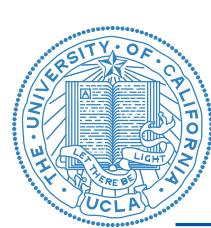
Dropout in code.

```
p = 0.5 # probability of dropping out
relu = lambda x: x * (x > 0)

def forward(X):
    H1 = relu(np.dot(W1, X) + b1) # First hidden layer activations
    M1 = np.random.rand(*H1.shape) < p # Sample random mask
    H1 *= M1 # Dropout on first hidden layer

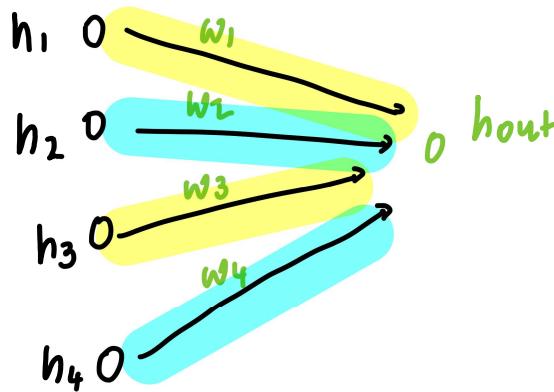
    H2 = relu(np.dot(W2, H1) + b2) # Second hidden layer activations
    M2 = np.random.rand(*H2.shape) < p # Sample random mask
    H2 *= M2 # Dropout on second hidden layer

    Z = np.dot(W3, H2) + b3
```



Dropout

How about during test time? What configuration do you use?



$$p = \underline{0.5}$$

↓
dropping units
~50% each time

Training iter = 1

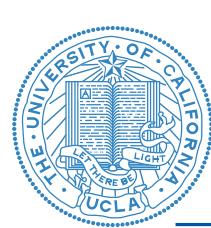
$$m = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix} \implies h_{\text{out}} = \text{relu}(w_1 h_1 + w_3 h_3)$$

Training iter = 2

$$m = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \end{bmatrix} \implies h_{\text{out}} = \text{relu}(w_2 h_2 + w_4 h_4)$$

Test:
$$\underline{h_{\text{out}}} = \text{relu}((w_1 h_1 + w_2 h_2 + w_3 h_3 + w_4 h_4) \cdot p)$$

over many iterations, the contribution
of $w_i h_i$ to h_{out} is actually $p \cdot w_i \cdot h_i$



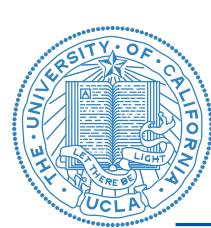
Dropout

How about during test time? What configuration do you use?

We call this approach the **weight scaling inference rule**. There is not yet any theoretical argument for the accuracy of this approximate inference rule in deep nonlinear networks, but empirically it performs very well.

In this class, instead of scaling the weights, we'll scale the activations.

) putting p outside
relu in formula on
prev. slide.



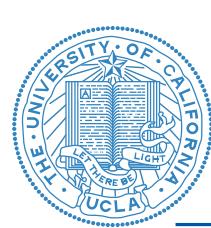
Dropout

```
p = 0.5 # probability of dropping out
relu = lambda x: x * (x > 0)

def test(X):
    H1 = relu(np.dot(W1, X) + b1) * p
    H2 = relu(np.dot(W1, H1) + b2) * p
    Z = np.dot(W3, H2) + b3
```

Note: an additional pro of dropout is that in testing time, there is no additional complexity. With m ensemble models, our test time evaluation would scale $O(m)$.





Inverted dropout

A common way to implement dropout is *inverted dropout* where the scaling by $1/p$ is done in training. This causes the output to have the same expected value as if dropout was never been performed.

Thus, testing looks the same irrespective of if we use dropout or not. See code below:

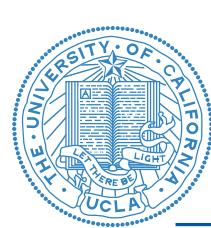
```
p = 0.5 # probability of dropping out
relu = lambda x: x * (x > 0)

def train_forward(X):
    H1 = relu(np.dot(W1, X) + b1) # First hidden layer activations
    M1 = (np.random.rand(*H1.shape) < p) / p # Sample random mask AND normalization by p
    H1 *= M1 # Dropout on first hidden layer

    H2 = relu(np.dot(W2, H1) + b2) # Second hidden layer activations
    M2 = (np.random.rand(*H2.shape) < p) / p # Sample random mask AND normalization by p
    H2 *= M2 # Dropout on first hidden layer

    z = np.dot(W3, H2) + b3

def test(X):
    H1 = relu(np.dot(W1, X) + b1)
    H2 = relu(np.dot(W2, H1) + b2)
    z = np.dot(W3, H2) + b3
```



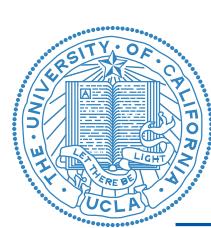
Dropout

How is this a good idea?

- 1) Dropout approximates bagging, since each mask is like a different model. For a model with N hidden units, there are 2^N different model configurations.

Each of these configurations must be good at predicting the output.

- 2) You can think of dropout as regularizing each hidden unit to work well in many different contexts.
- 3) Dropout may cause units to encode redundant features (e.g., to detect a cat, there are many things we look for, e.g., it's furry, it has pointy ears, it has a tail, a long body, etc.).



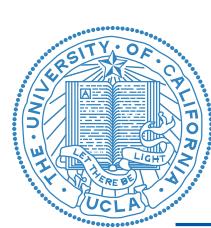
Lecture summary

Here, we've covered tricks that we can do in initialization, regularization, and data augmentation to improve the performance of neural networks.

But what about the optimizer, stochastic gradient descent? Can we improve this for deep learning?

That's the topic of our next lecture.



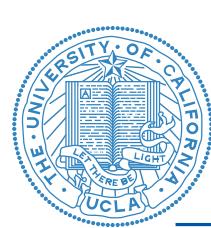


Optimization for neural networks

In this lecture, we'll talk about specific techniques in optimization that aid in training neural networks.

- Stochastic gradient descent
- Momentum and Nesterov momentum
- Adaptive gradients
- RMSProp
- Adaptive moments (**Adam**)
- Overview of second order methods
- Challenges of gradient descent

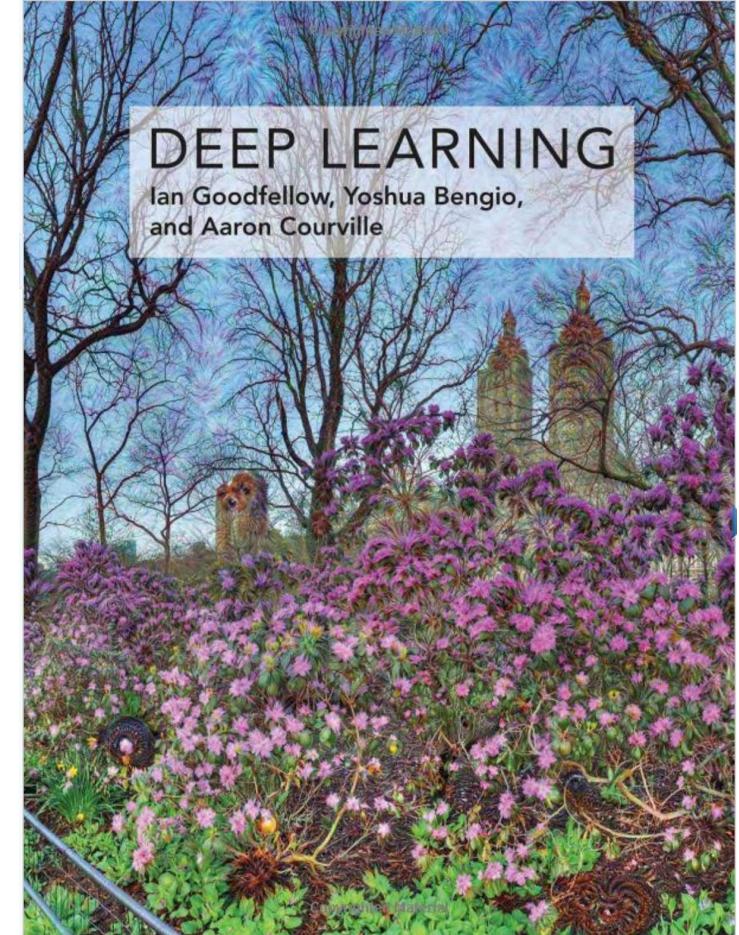


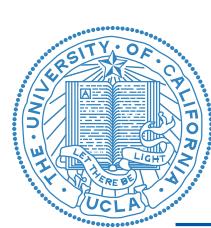


Reading

Reading:

Deep Learning, Chapter 8 (intro), 8.1, 8.2, 8.3,
8.4, 8.5, 8.6 (skim)





Where we are now

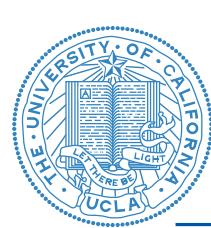
At this point, we know:

- Neural network architectures.
- Hyperparameters and cost functions to use for neural networks.
- How to calculate gradients of the loss w.r.t. all parameters in the neural network.
- How to initialize the weights and regularize the network in ways to improve the training of the network.

We do know how to optimize these networks with stochastic gradient descent. But can it be improved?

In this lecture, we talk about how to make optimization more efficient and effective.





Gradient descent

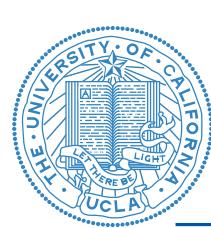
A refresher on gradient descent.

- Cost function: $J(\theta)$ -Loss.
- Parameters: θ

Then, the gradient descent step is:

$$\theta \leftarrow \theta - \epsilon \nabla_{\theta} J(\theta)$$





Stochastic gradient descent

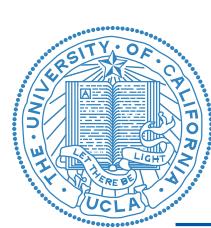
Calculating the gradient exactly is expensive, because it requires evaluating the model on all m examples in the dataset. This leads to an important distinction.

- Batch algorithm: uses all m examples in the training set to calculate the gradient.
- Minibatch algorithm: approximates the gradient by calculating it using k training examples, where $m > k > 1$.
- Stochastic algorithm: approximates the gradient by calculating it over one example.

It is typical in deep learning to use minibatch gradient descent. Note that some may also use minibatch and stochastic gradient descent interchangeably.

A note: small batch sizes can be seen to have a regularization effect, perhaps because they introduce noise to the training process.





Stochastic gradient descent

Stochastic gradient descent

Stochastic gradient descent proceeds as follows.

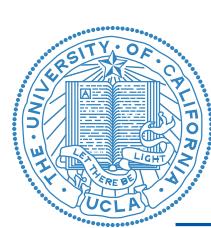
Set a learning rate ε and an initial parameter setting θ . Set a minibatch size of m examples. Until the stopping criterion is met:

- Sample m examples from the training set, $\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}\}$ and their corresponding outputs $\{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(m)}\}$.
- Compute the gradient estimate:

$$\mathbf{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} J(\theta)$$

- Update parameters:

$$\theta \leftarrow \theta - \varepsilon \mathbf{g}$$

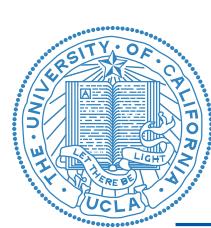


Stochastic gradient descent

```
while last_diff > tol:  
    cost, g = func(x)  
    x -= eps*g  
    last_diff = np.linalg.norm(x - path[-1])
```

softmax.log-and-grad()

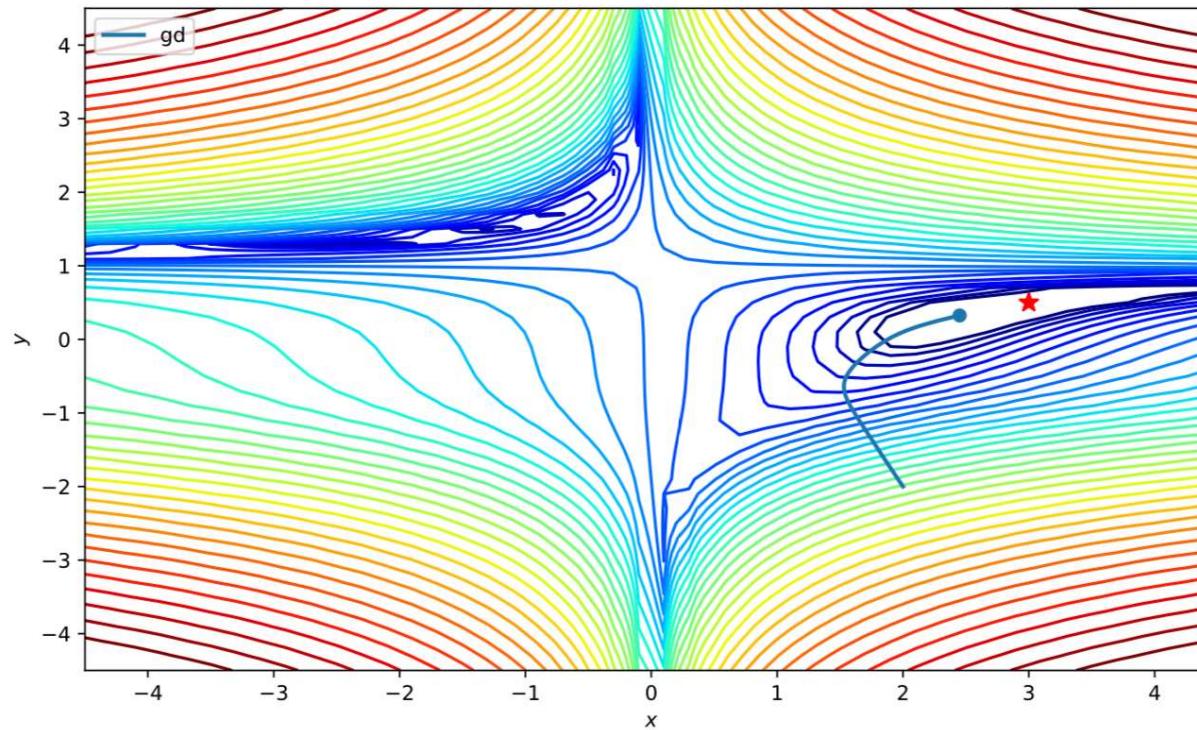




Stochastic gradient descent

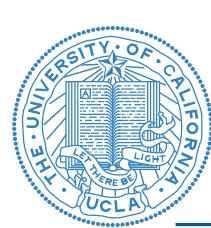
Stochastic gradient descent (opt 1)

The following shows gradient descent applied to Beale's function, for two initializations at $(2, -2)$ and $(-3, 3)$. The two initializations are shown because later on we'll contrast to other techniques. The iteration count is capped at 10,000 iterations, so gradient descent does not get to the minimum.

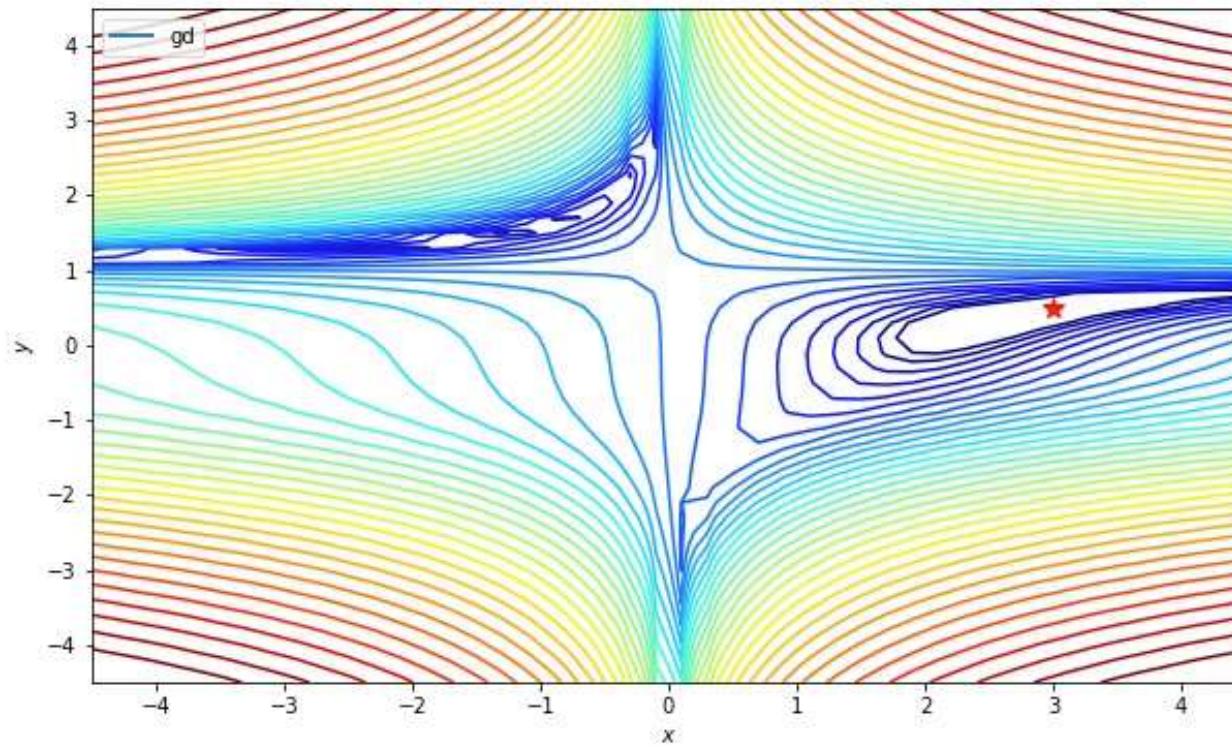


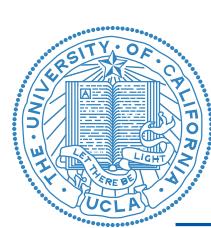
Video: http://seas.ucla.edu/~kao/opt_anim/1gd.mp4

Animation help thanks to: <http://louistiao.me/note/visualizing-and-animating-optimization-algorithms-with-matplotlib/>



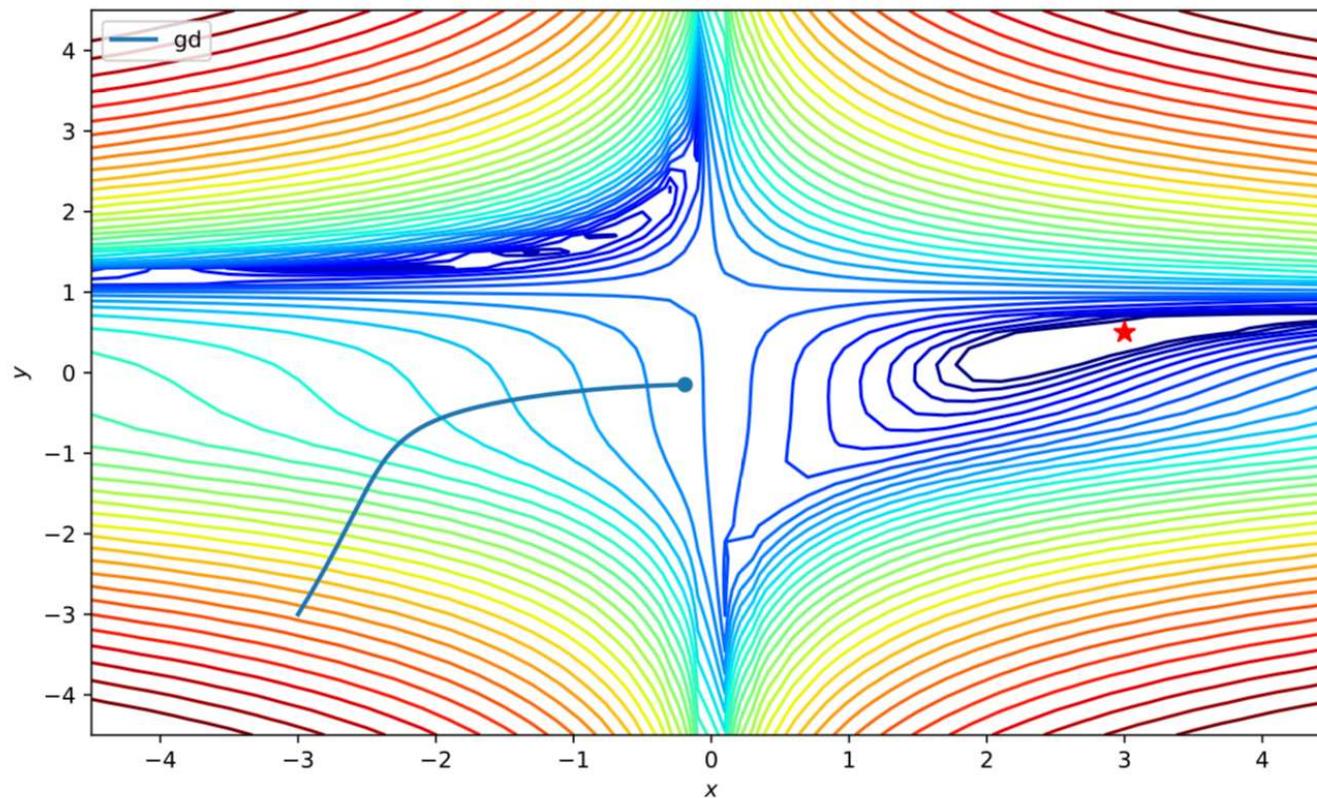
Stochastic gradient descent



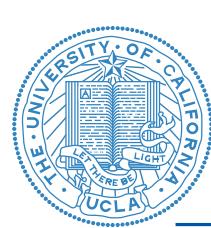


Stochastic gradient descent

Stochastic gradient descent (opt 2)

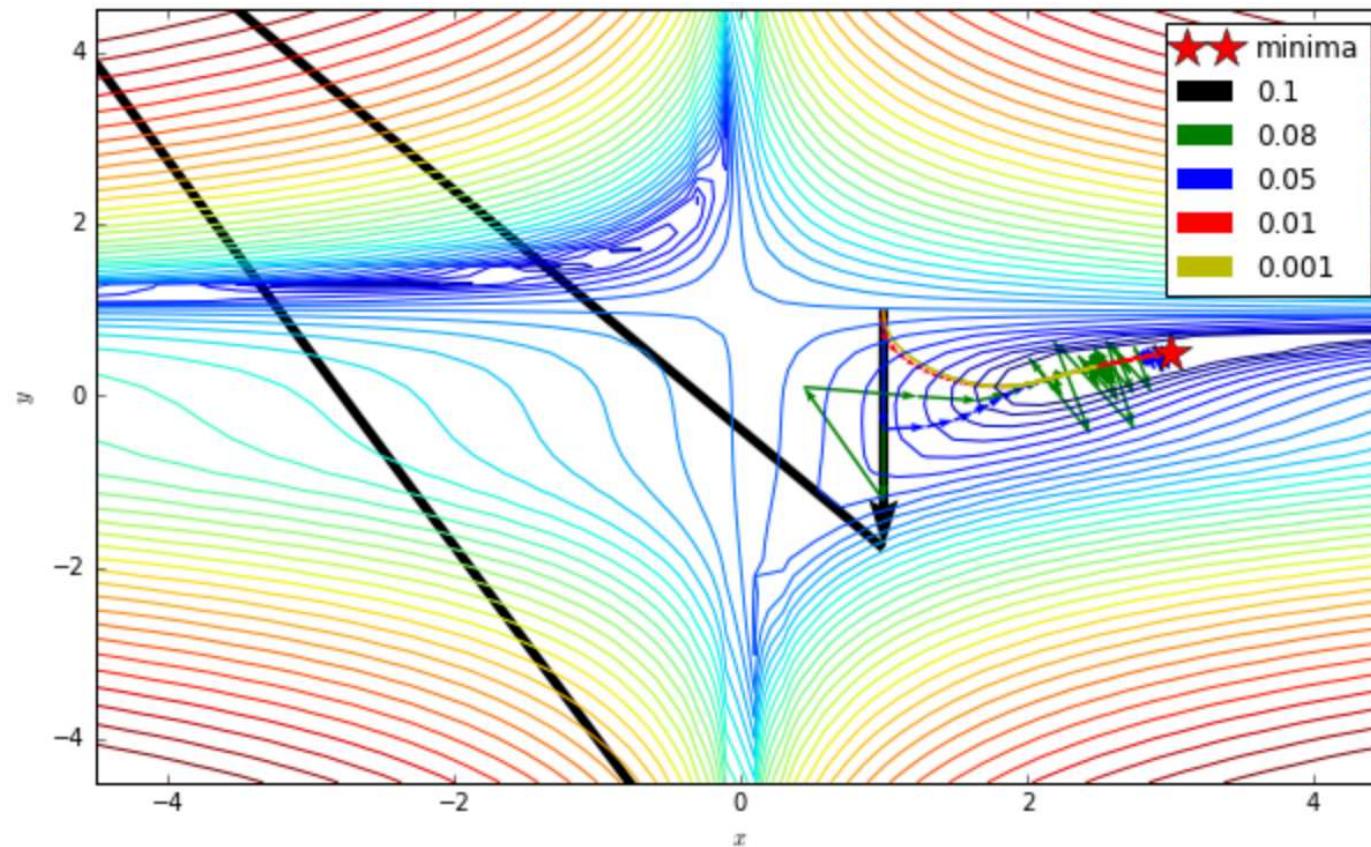


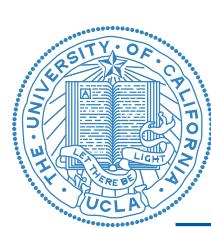
Video: http://seas.ucla.edu/~kao/opt_anim/2gd.mp4



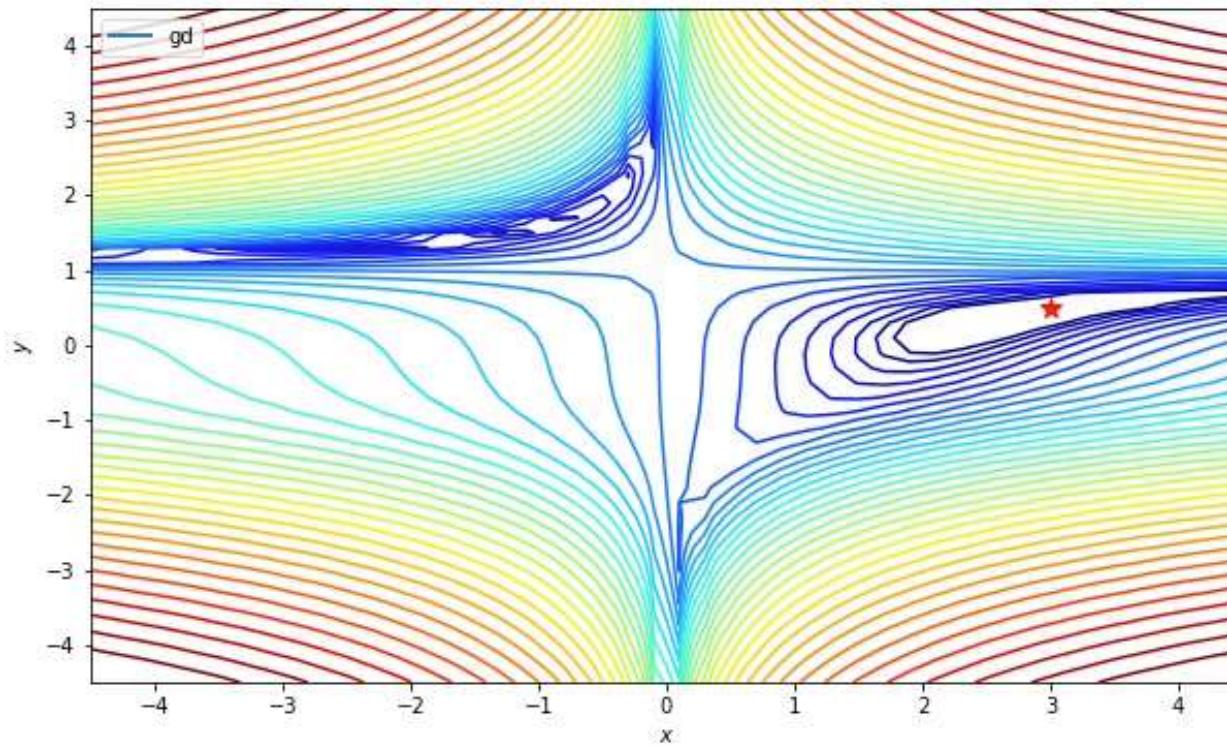
Finding the optimal weights through gradient descent

Varying the learning rate:





Stochastic gradient descent

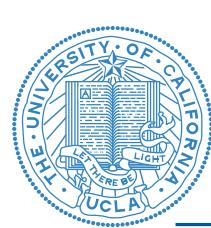




Lecture 10: Optimization + convolutional neural networks

Announcements:

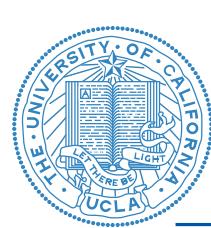
- HW #4 is due this Friday, February 14. **To submit your Jupyter Notebook, print the notebook to a pdf with your solutions and plots filled in. You must also submit your .py files as pdfs. Please make sure they are readable.**
- Midterm review session: Young Hall CS76 on February 14, from 4pm-6pm.
- The midterm is in 9 days on February 19, 2025.
 - You are allowed 4 cheat sheets (each an 8.5 x 11 inch paper). You can fill out both sides (8 sides total). You can put whatever you want on these cheat sheets.
 - It will cover material up to and including lecture on February 12, 2025.
 - Past exams are uploaded to Bruin Learn (under “Modules” —> “past exams”).
 - You may bring a calculator to the exam.
 - You **must do the exam in pen.**
 - MSOL students will take the exam remotely on Saturday, February 22, from 2-3:50pm. We will send out details to the MSOL students closer to the date.
- If you have CAE accommodations, you must request CAE to provide them.
- We will release the project before Wednesday’s lecture (and go over it at the start).



Optimization for neural networks

In this lecture, we'll talk about specific techniques in optimization that aid in training neural networks.

- Stochastic gradient descent
- Momentum and Nesterov momentum
- Adaptive gradients
- RMSProp
- Adaptive moments
- Overview of second order methods
- Challenges of gradient descent



Stochastic gradient descent

Stochastic gradient descent

Stochastic gradient descent proceeds as follows.

Set a learning rate ε and an initial parameter setting θ . Set a minibatch size of m examples. Until the stopping criterion is met:

- Sample m examples from the training set, $\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}\}$ and their corresponding outputs $\{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(m)}\}$.
- Compute the gradient estimate:

$$\mathbf{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} J(\theta)$$

- Update parameters:

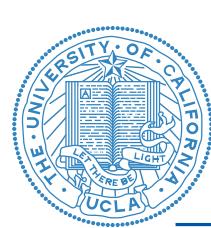
$$\theta \leftarrow \theta - \varepsilon \mathbf{g}$$



Stochastic gradient descent

```
while last_diff > tol:  
    cost, g = func(x)  
    x -= eps*g  
    last_diff = np.linalg.norm(x - path[-1])
```

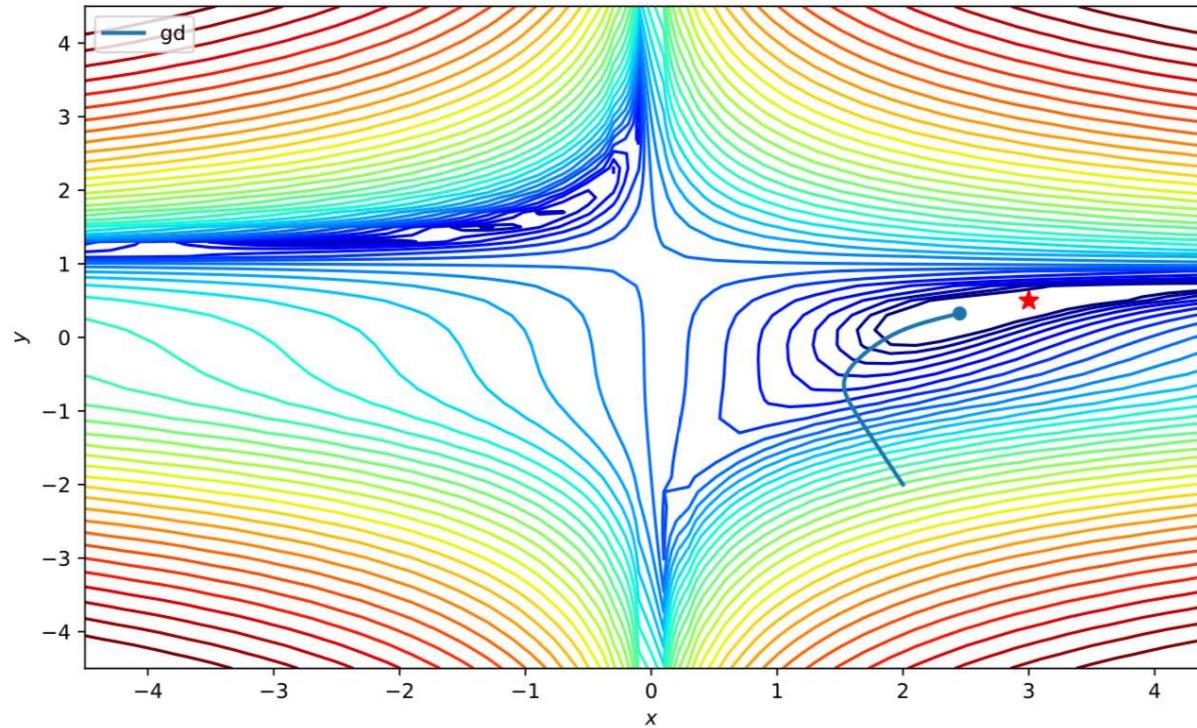
softmax-loss-and-gradient()



Stochastic gradient descent

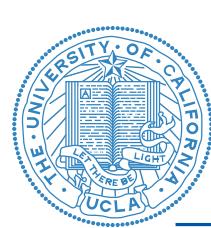
Stochastic gradient descent (opt 1)

The following shows gradient descent applied to Beale's function, for two initializations at $(2, -2)$ and $(-3, 3)$. The two initializations are shown because later on we'll contrast to other techniques. The iteration count is capped at 10,000 iterations, so gradient descent does not get to the minimum.

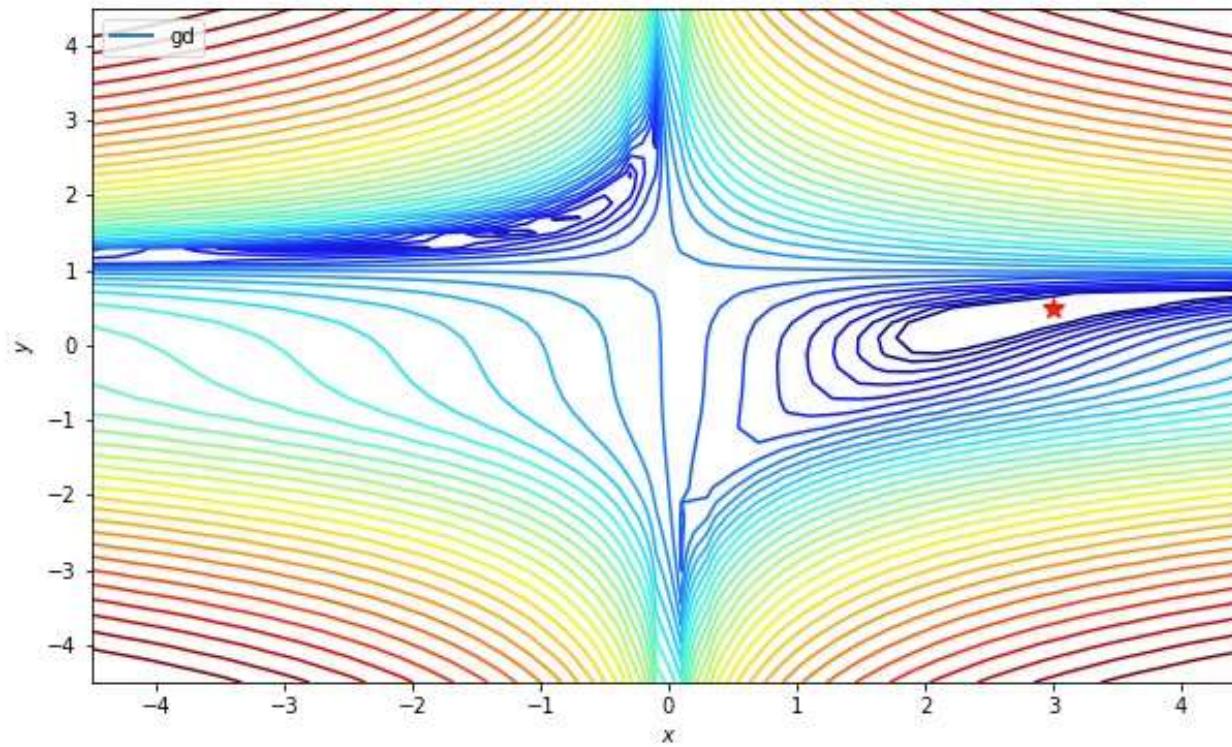


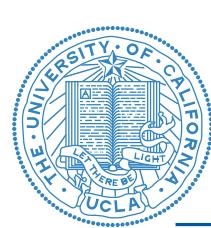
Video: http://seas.ucla.edu/~kao/opt_anim/1gd.mp4

Animation help thanks to: <http://louistiao.me/note/visualizing-and-animating-optimization-algorithms-with-matplotlib/>



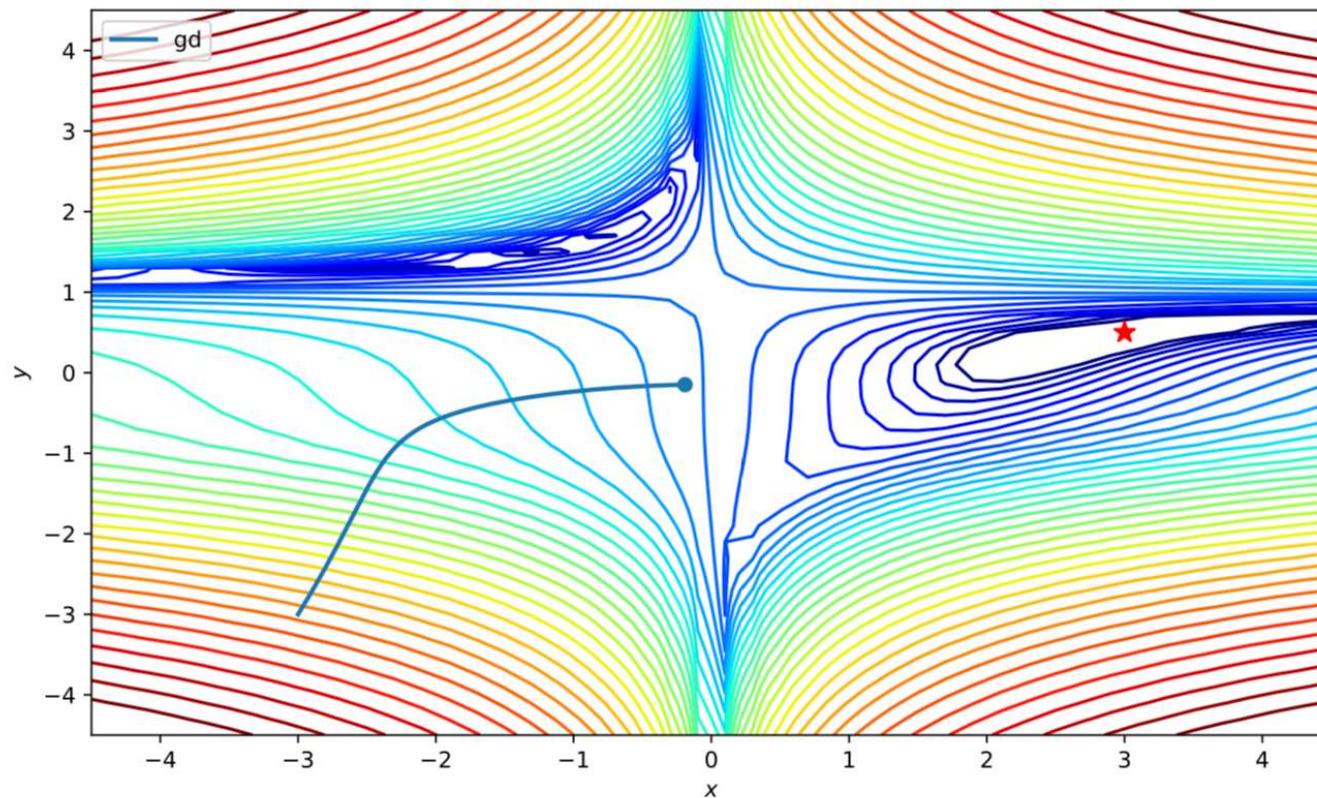
Stochastic gradient descent





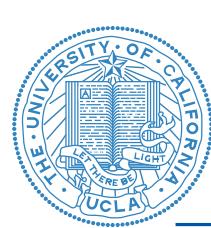
Stochastic gradient descent

Stochastic gradient descent (opt 2)



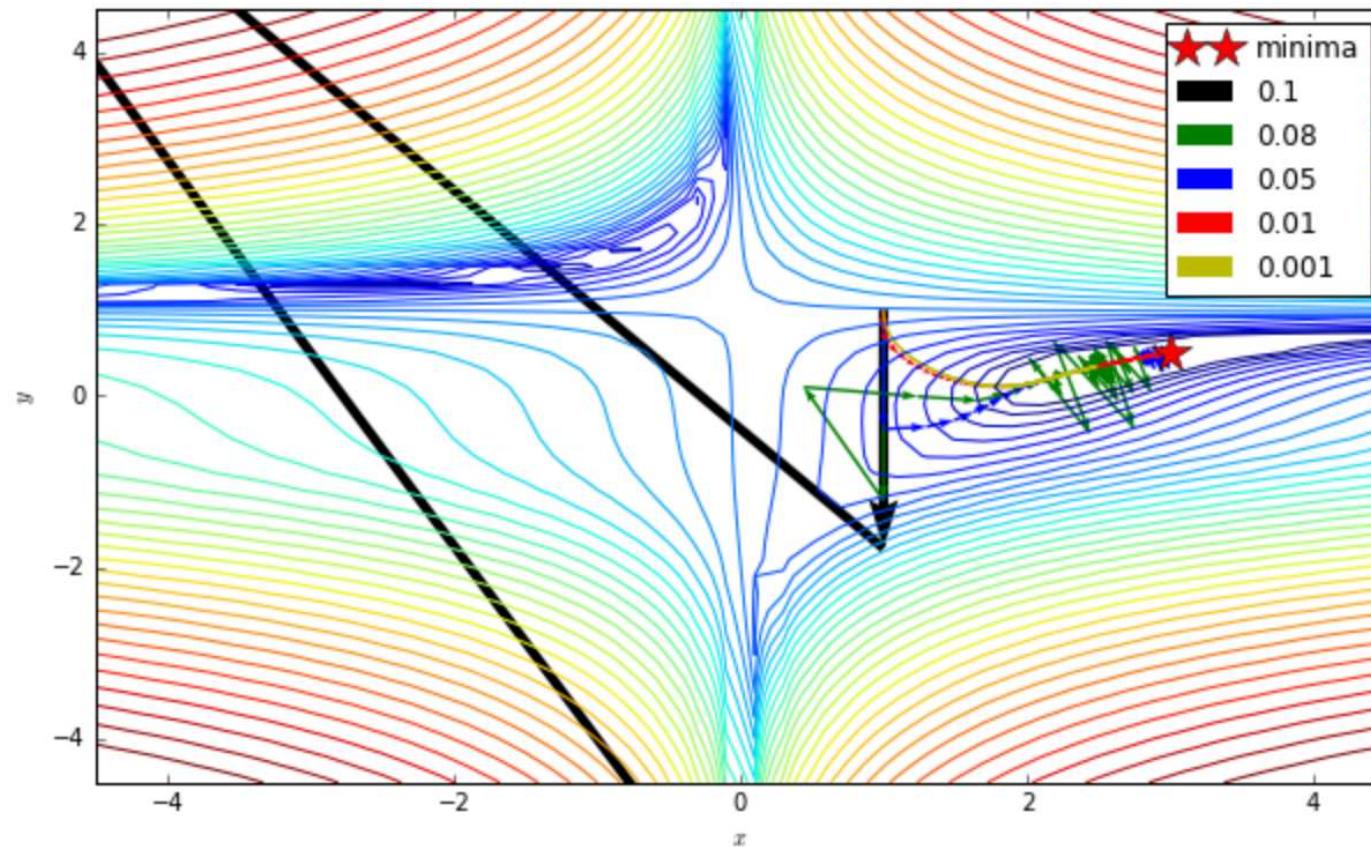
Video: http://seas.ucla.edu/~kao/opt_anim/2gd.mp4

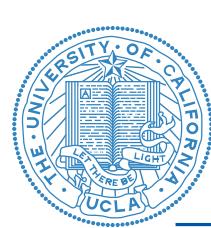




Finding the optimal weights through gradient descent

Varying the learning rate:

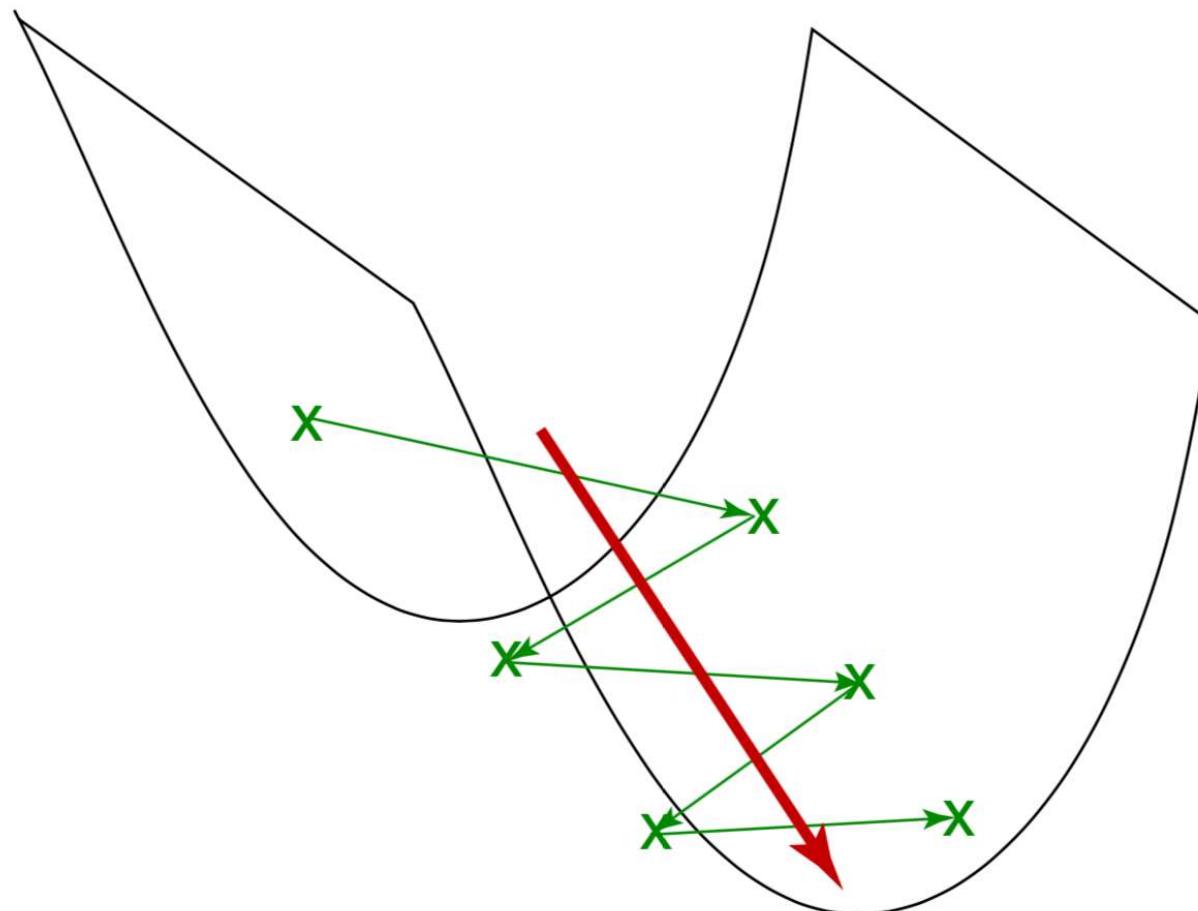




Momentum

Momentum (cont.)

An example of how momentum is useful is to consider a tilted surface with high curvature. Stochastic gradient descent may make steps that zigzag, although in general it proceeds in the right direction. Momentum will average away the zigzagging components.





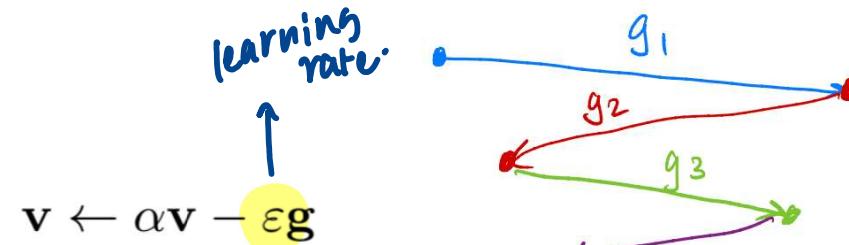
Momentum

Momentum

In momentum, we maintain the running mean of the gradients, which then updates the parameters.

Initialize $v = 0$. Set $\alpha \in [0, 1]$. Typical values are $\alpha = 0.9$ or 0.99 . Then, until stopping criterion is met:

- Compute gradient: g
- Update:
- Gradient step:



$$\theta \leftarrow \theta + v$$

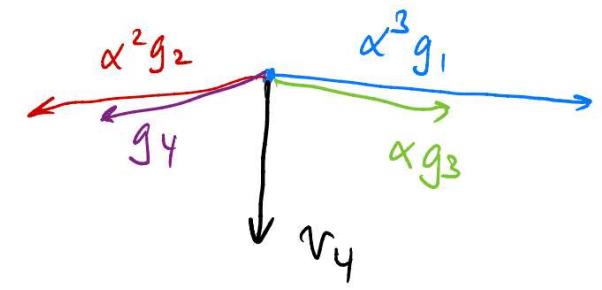
$$v_0 = 0$$

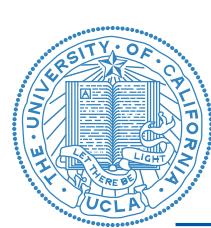
$$v_1 = -\varepsilon g_1$$

$$v_2 = \alpha v_1 - \varepsilon g_2 = -\alpha \varepsilon g_1 - \varepsilon g_2$$

$$v_3 = \alpha v_2 - \varepsilon g_3 = -\varepsilon (\alpha^2 g_1 + \alpha g_2 + g_3)$$

$$v_4 = -\varepsilon (\alpha^3 g_1 + \alpha^2 g_2 + \alpha g_3 + g_4)$$

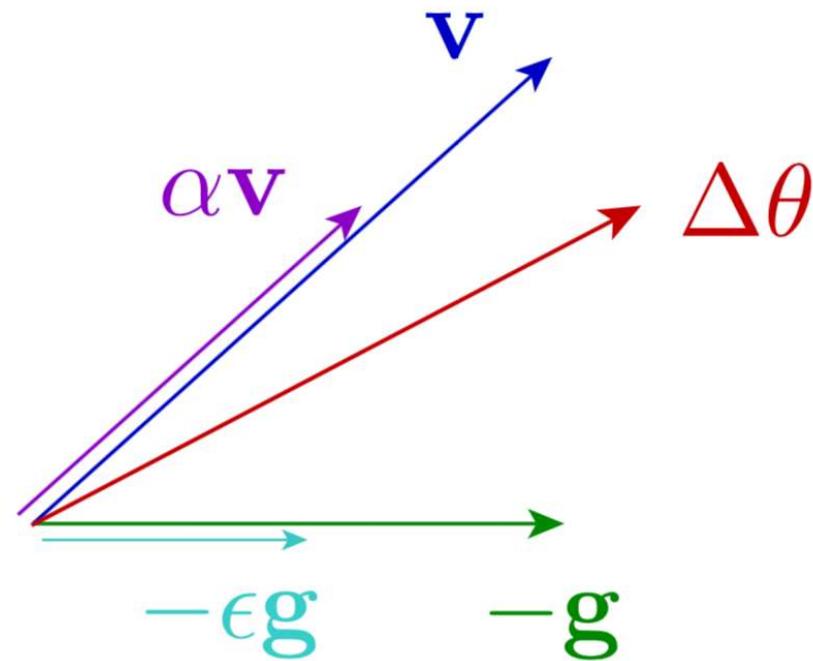


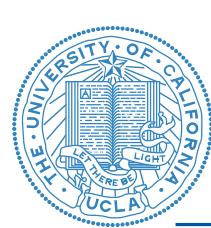


Momentum

Momentum (cont.)

This modification augments the gradient with the running average of previous gradients, which is analogous to a gradient “momentum.” The following image is appropriate to have in mind:

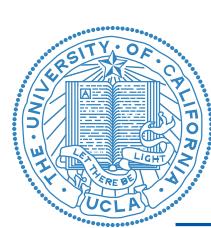




Momentum

```
alpha = 0.9
p = 0
while last_diff > tol:
    cost, g = func(x)
    p = alpha*p - eps*g
    x += p
    last_diff = np.linalg.norm(x - path[-1])
```

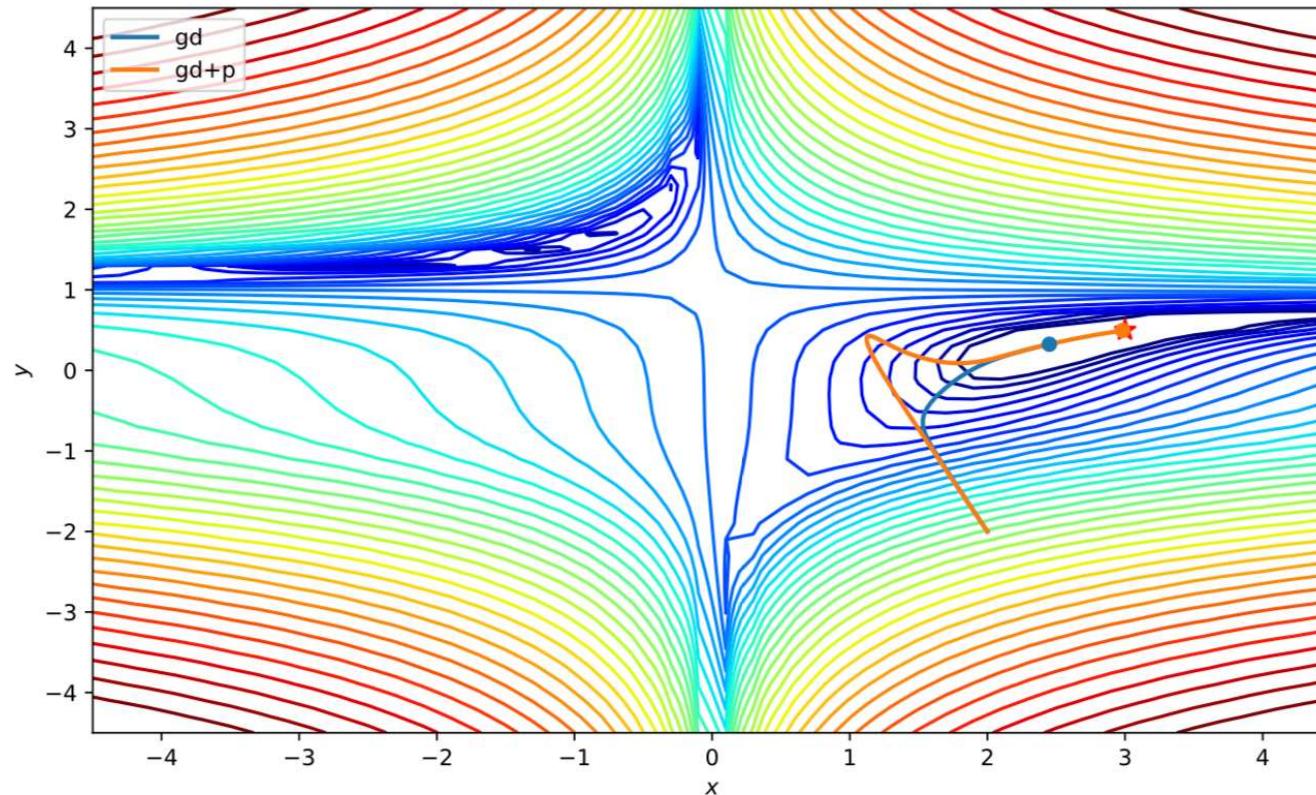




Momentum

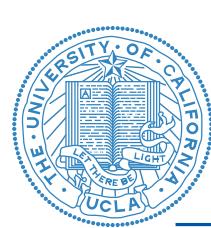
Momentum (opt 1)

gd+p denotes gradient descent with momentum.

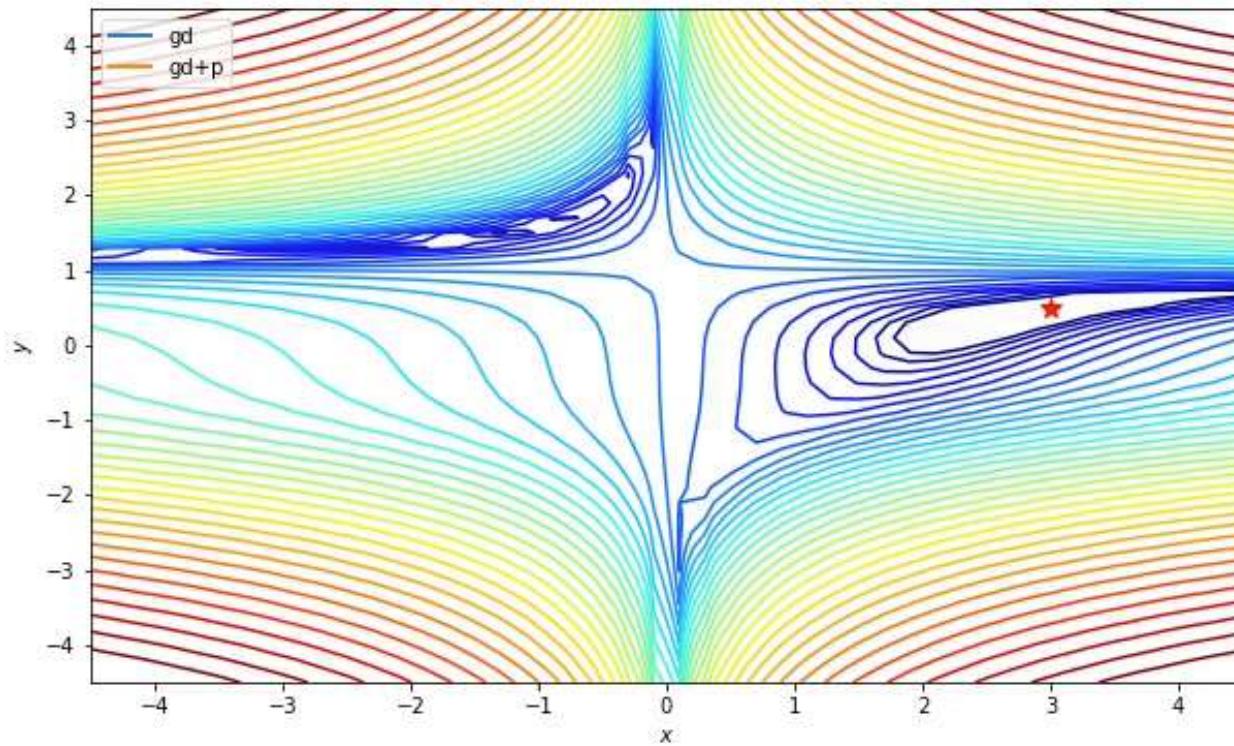


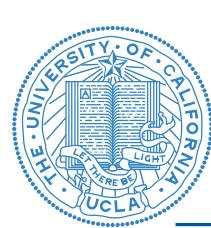
Video: http://seas.ucla.edu/~kao/opt_anim/1gd_gd+p.mp4





Momentum

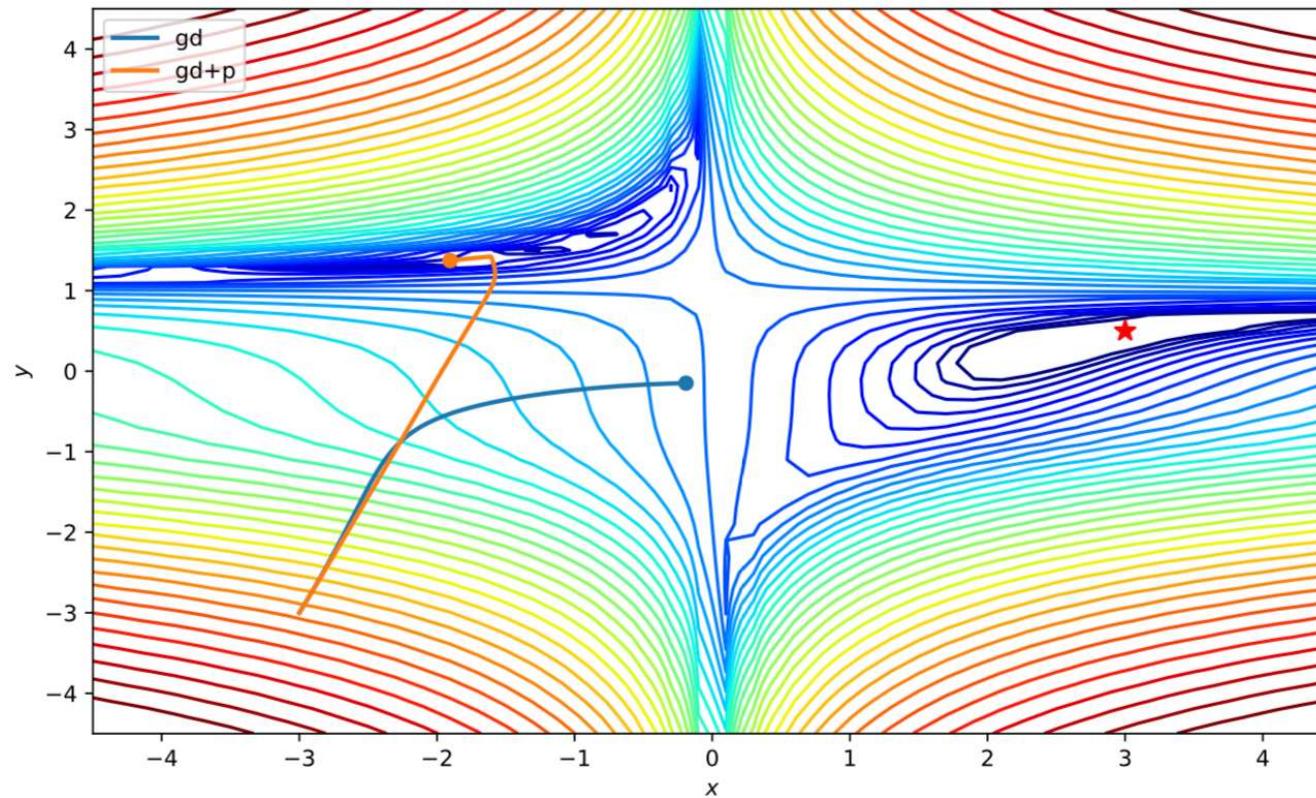




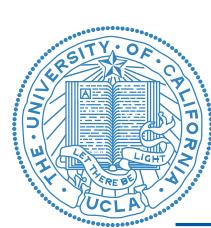
Momentum

Momentum (opt 2)

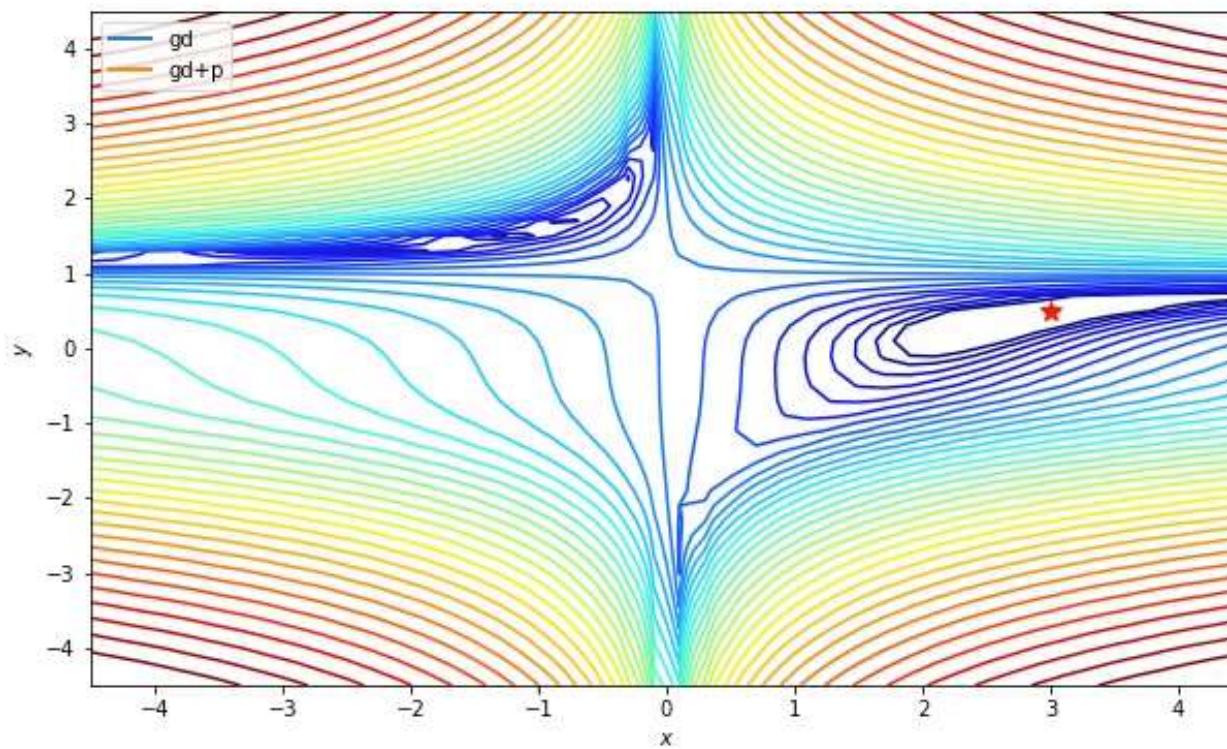
Notice how momentum pushes the descent to find a local, but not global, minimum.

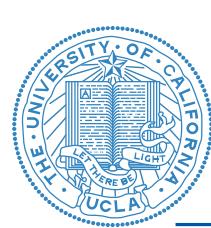


Video: http://seas.ucla.edu/~kao/opt_anim/2gd_gd+p.mp4



Momentum



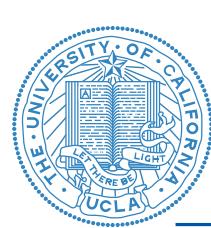


Does momentum help with local optima?

Does momentum help with local optima?

What kind of local optima does momentum tend to find?





Nesterov momentum

Nesterov momentum

Nesterov momentum is similar to momentum, except the gradient is calculated at the parameter setting after taking a step along the direction of the momentum.

Initialize $\mathbf{v} = 0$. Then, until stopping criterion is met:

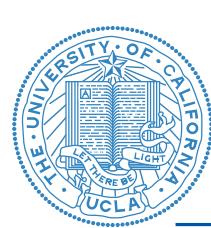
- Update:

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \varepsilon \nabla_{\theta} J(\theta + \alpha \mathbf{v})$$

Gradient¹ ($\theta + \alpha \mathbf{v}$)

- Gradient step:

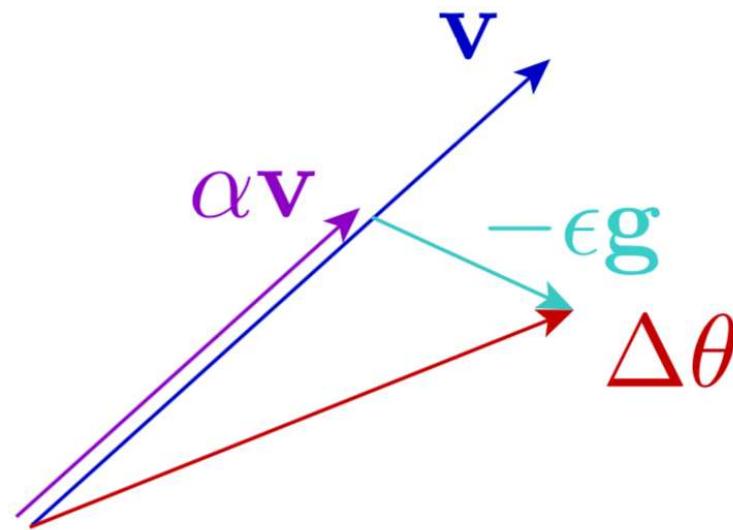
$$\theta \leftarrow \theta + \mathbf{v}$$

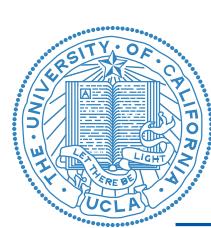


Nesterov momentum

Nesterov momentum (cont.)

The following image is appropriate for Nesterov momentum:





Nesterov momentum

By performing a change of variables with $\tilde{\theta}_{\text{old}} = \theta_{\text{old}} + \alpha \mathbf{v}_{\text{old}}$, it's possible to show that the following is equivalent to Nesterov momentum. (This representation doesn't require evaluating the gradient at $\theta + \alpha \mathbf{v}$.)

- Update:

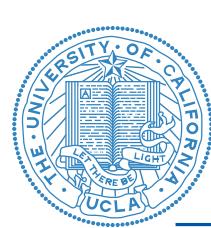
$$\mathbf{v}_{\text{new}} = \alpha \mathbf{v}_{\text{old}} - \varepsilon \nabla_{\tilde{\theta}_{\text{old}}} J(\tilde{\theta}_{\text{old}})$$

- Gradient step:

$$\tilde{\theta}_{\text{new}} = \tilde{\theta}_{\text{old}} + \mathbf{v}_{\text{new}} + \alpha (\mathbf{v}_{\text{new}} - \mathbf{v}_{\text{old}})$$

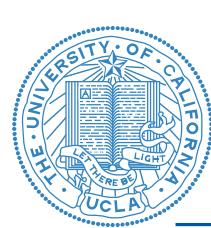
- Set $\mathbf{v}_{\text{new}} = \mathbf{v}_{\text{old}}$, $\tilde{\theta}_{\text{new}} = \tilde{\theta}_{\text{old}}$.





Nesterov momentum

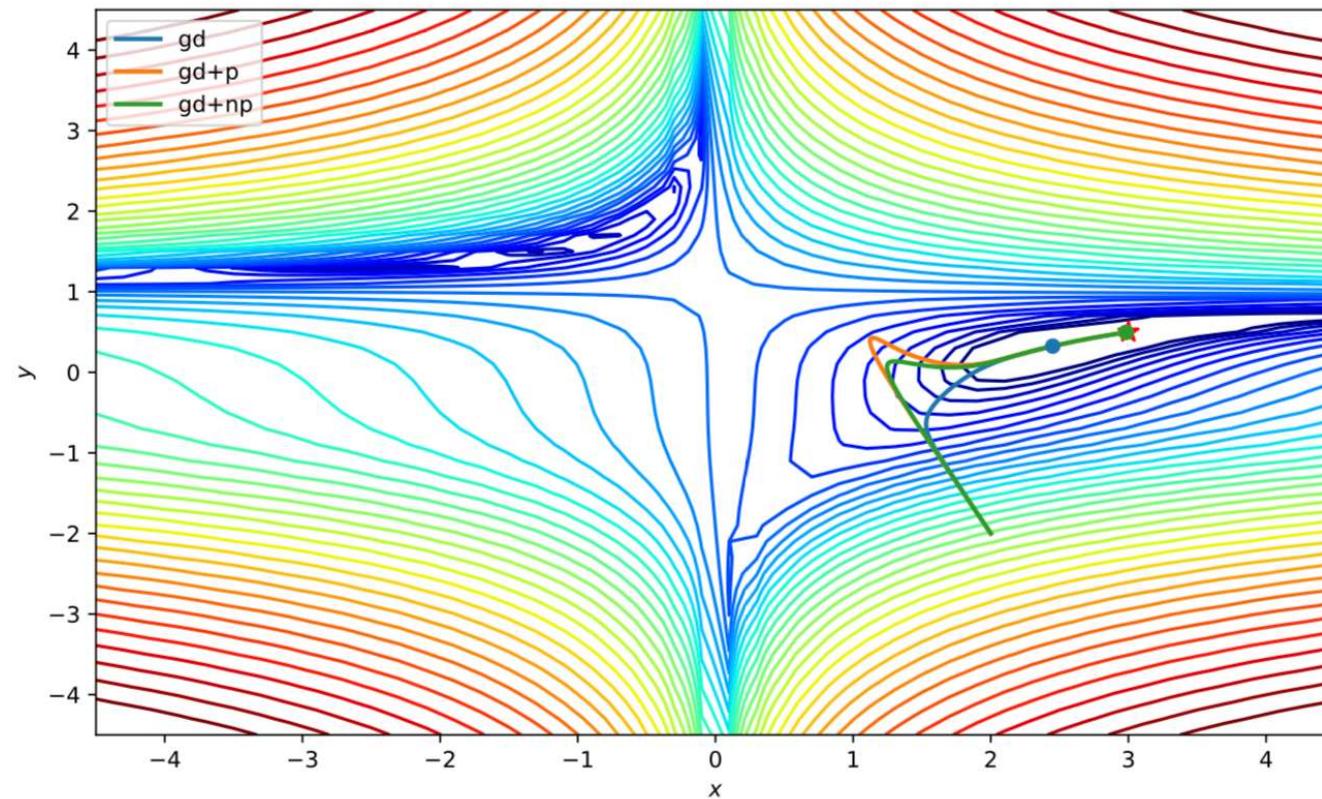
```
alpha = 0.9
p = np.zeros_like(x)
while last_diff > tol:
    cost, g = func(x)
    p_old = p
    p = alpha*p - eps*g
    x += p + alpha*(p-p_old)
    last_diff = np.linalg.norm(x - path[-1])
```



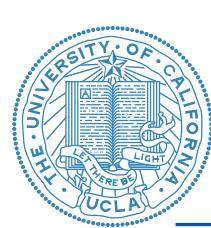
Nesterov momentum

Nesterov momentum (opt 1)

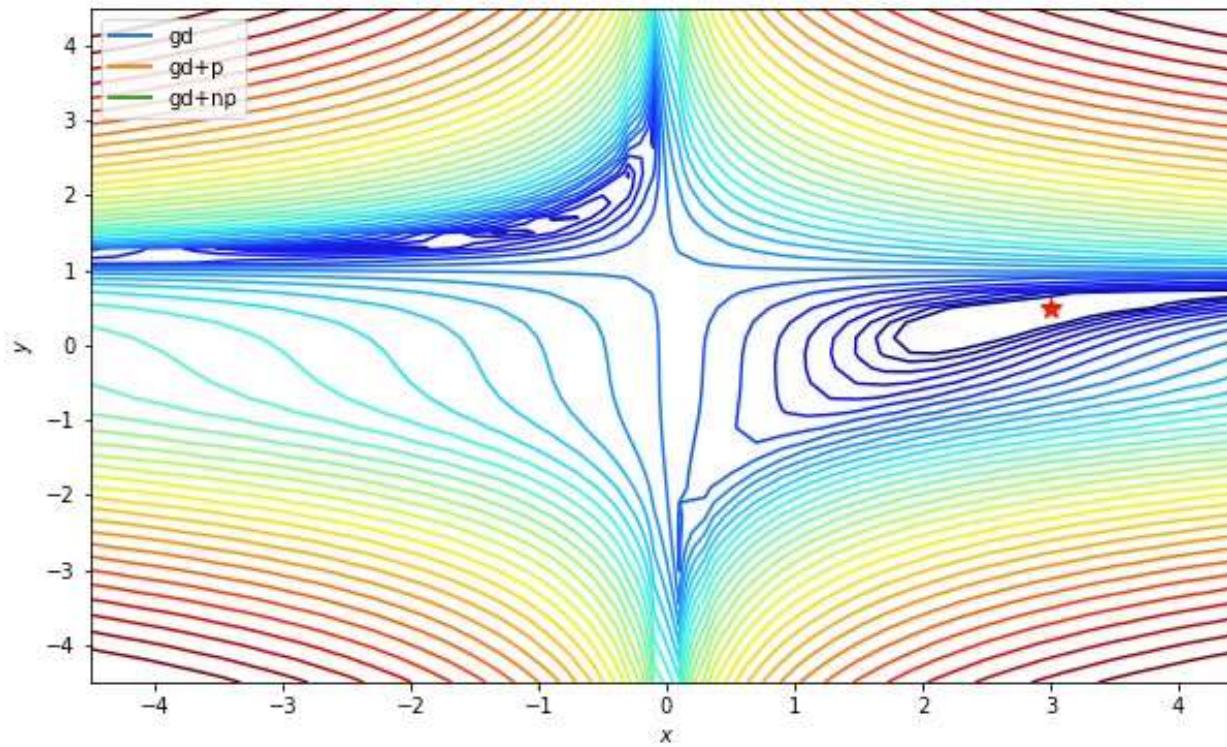
gd+np denotes gradient descent with Nesterov momentum.

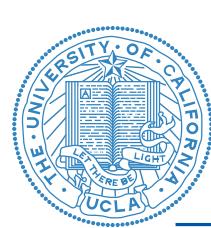


Video: http://seas.ucla.edu/~kao/opt_anim/1gd_gd+p_gd+np.mp4



Nesterov momentum

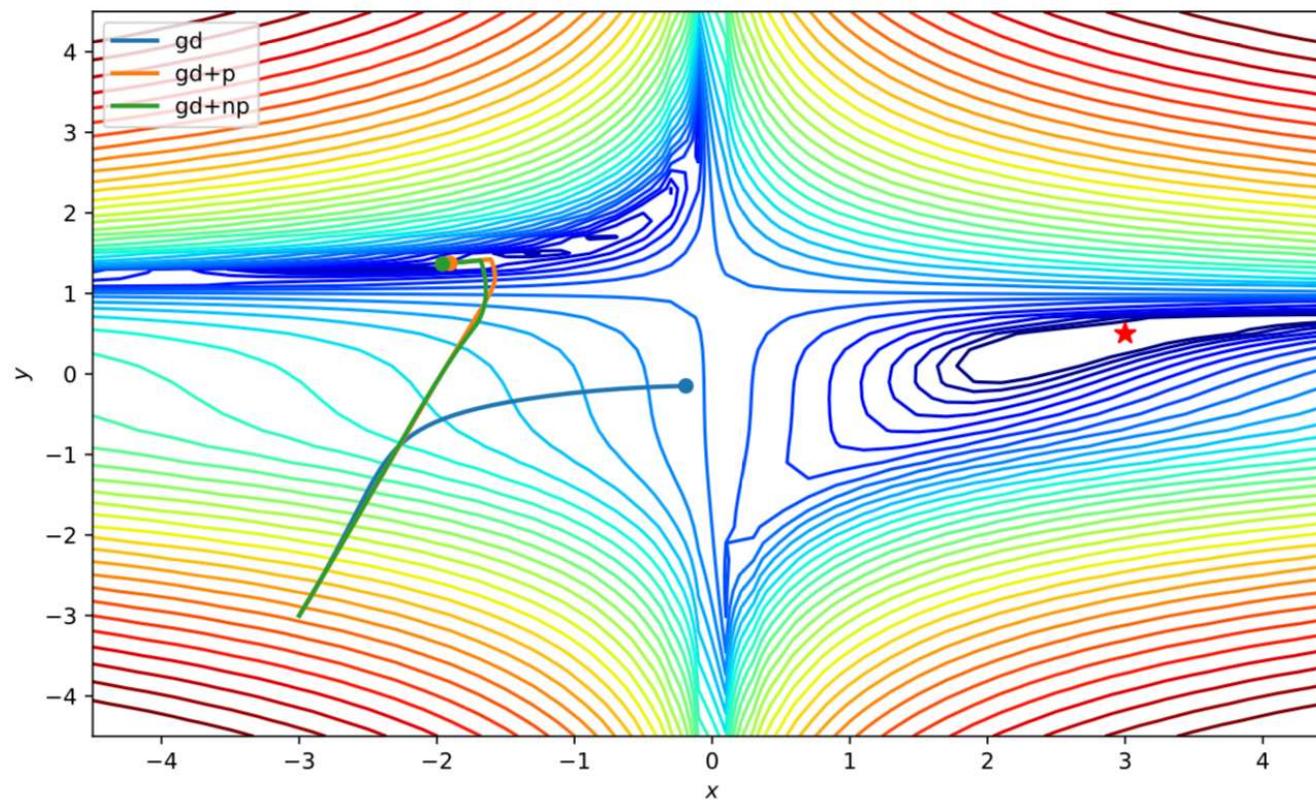




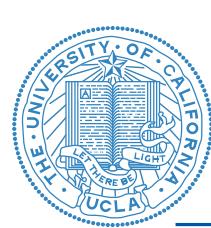
Nesterov momentum

Nesterov momentum (opt 2)

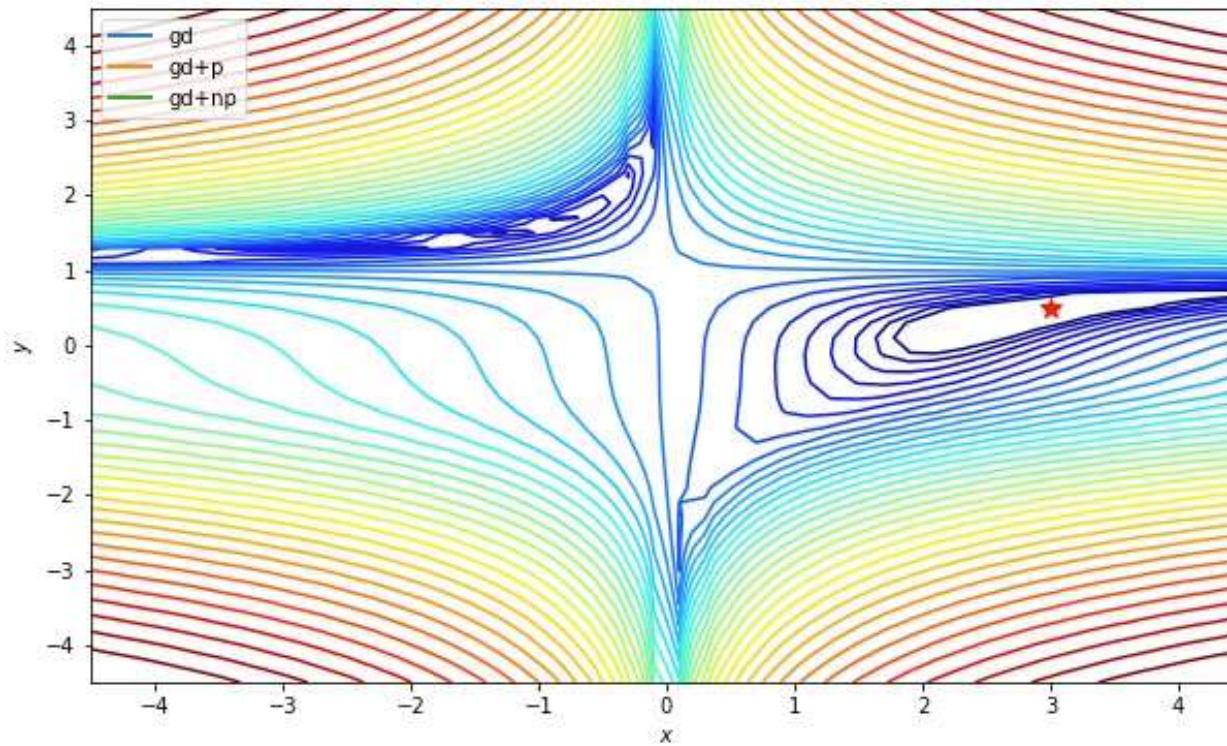
Notice how Nesterov momentum finds the same local minimum as momentum.

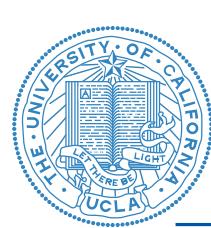


Video: http://seas.ucla.edu/~kao/opt_anim/2gd_gd+p_gd+np.mp4



Nesterov momentum





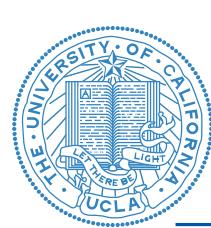
Is there a good way to adapt the learning rule?

Techniques to adapt the learning rate

Choosing ε judiciously can be important for learning. In the beginning, a larger learning rate is typically better, since bigger updates in the parameters may accelerate learning. However, as time goes on, ε may need to be small to be able to make appropriate updates to the parameters. We mentioned before that often times, one applies a decay rule to the learning rate. This is called *annealing*. A common form to anneal the learning rate is to do so manually when the loss plateaus, or to anneal it after a set number of epochs of gradient descent.

Another approach is to update the learning rate based off of the history of gradients.





Adagrad

Adaptive gradient (Adagrad)

Adaptive gradient (Adagrad) is a form of stochastic gradient descent where the learning rate is decreased through division by the historical gradient norms. We will let the variable a denote a running sum of squares of gradient norms.

Initialize $\mathbf{a} = 0$. Set ν at a small value to avoid division by zero (e.g., $\nu = 1e - 7$). Then, until stopping criterion is met:

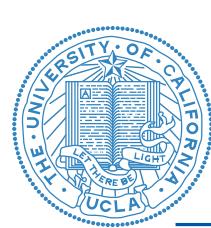
- Compute the gradient: \mathbf{g}
- Update:

$$\mathbf{a} \leftarrow \mathbf{a} + \mathbf{g} \odot \mathbf{g}$$

- Gradient step:

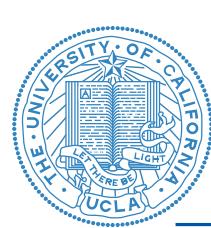
$$\theta \leftarrow \theta - \frac{\varepsilon}{\sqrt{\mathbf{a}} + \nu} \odot \mathbf{g}$$





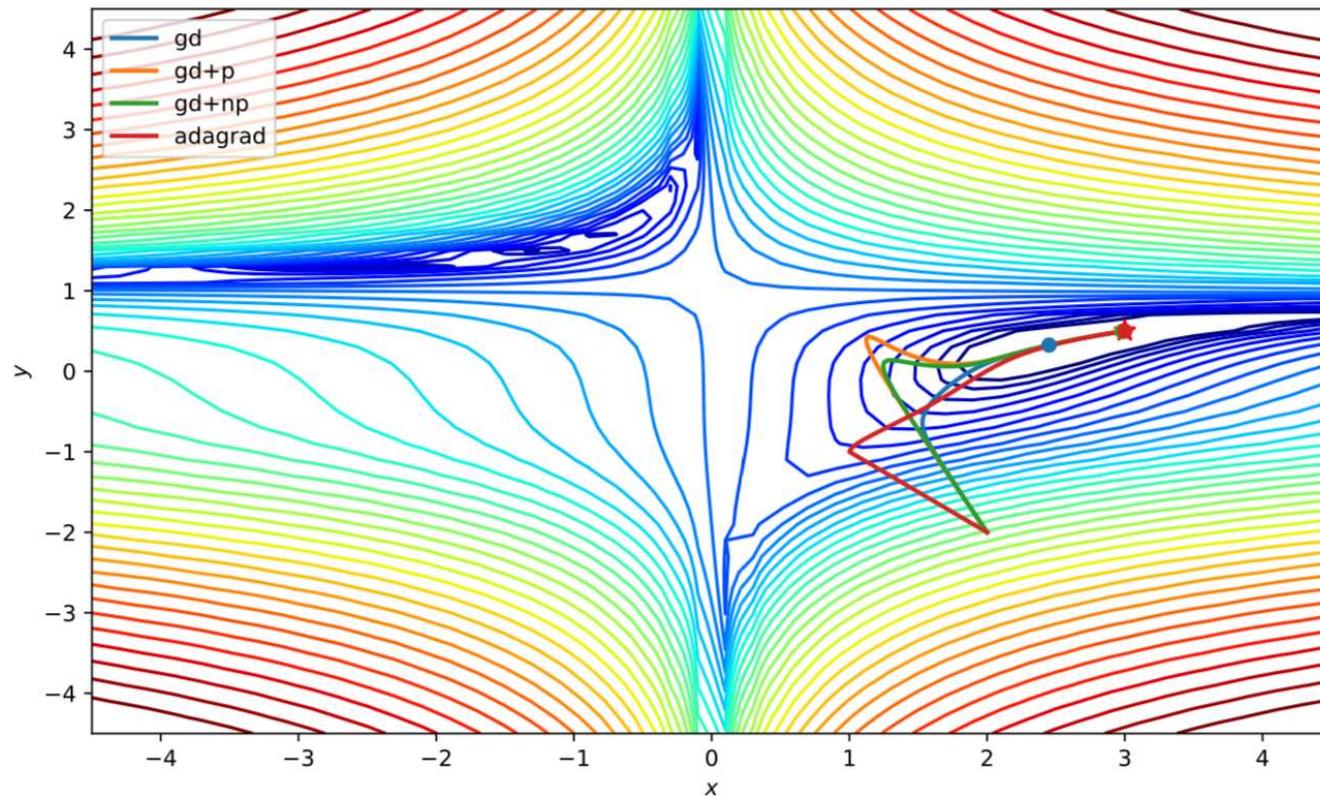
Adagrad

```
a = 0
nu = 1e-7
while last_diff > tol:
    cost, g = func(x)
    a += g*g
    x -= eps * g / (np.sqrt(a) + nu)
    last_diff = np.linalg.norm(x - path[-1])
```

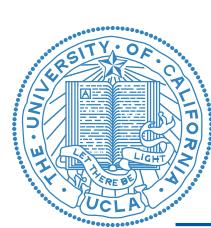


Adagrad

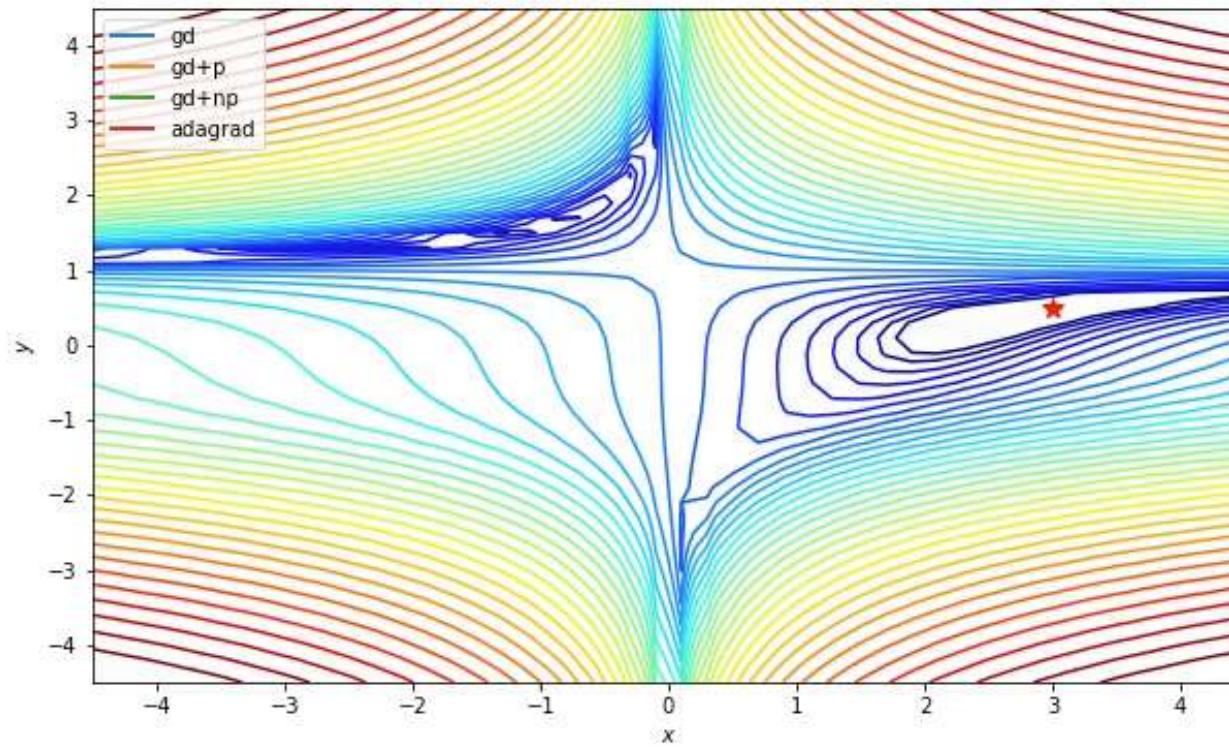
Adagrad (opt 1)

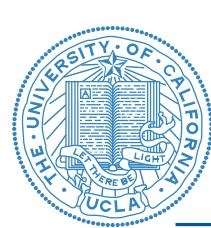


Video: http://seas.ucla.edu/~kao/opt_anim/1gd_gd+p_gd+np_adagrad.mp4



Adagrad

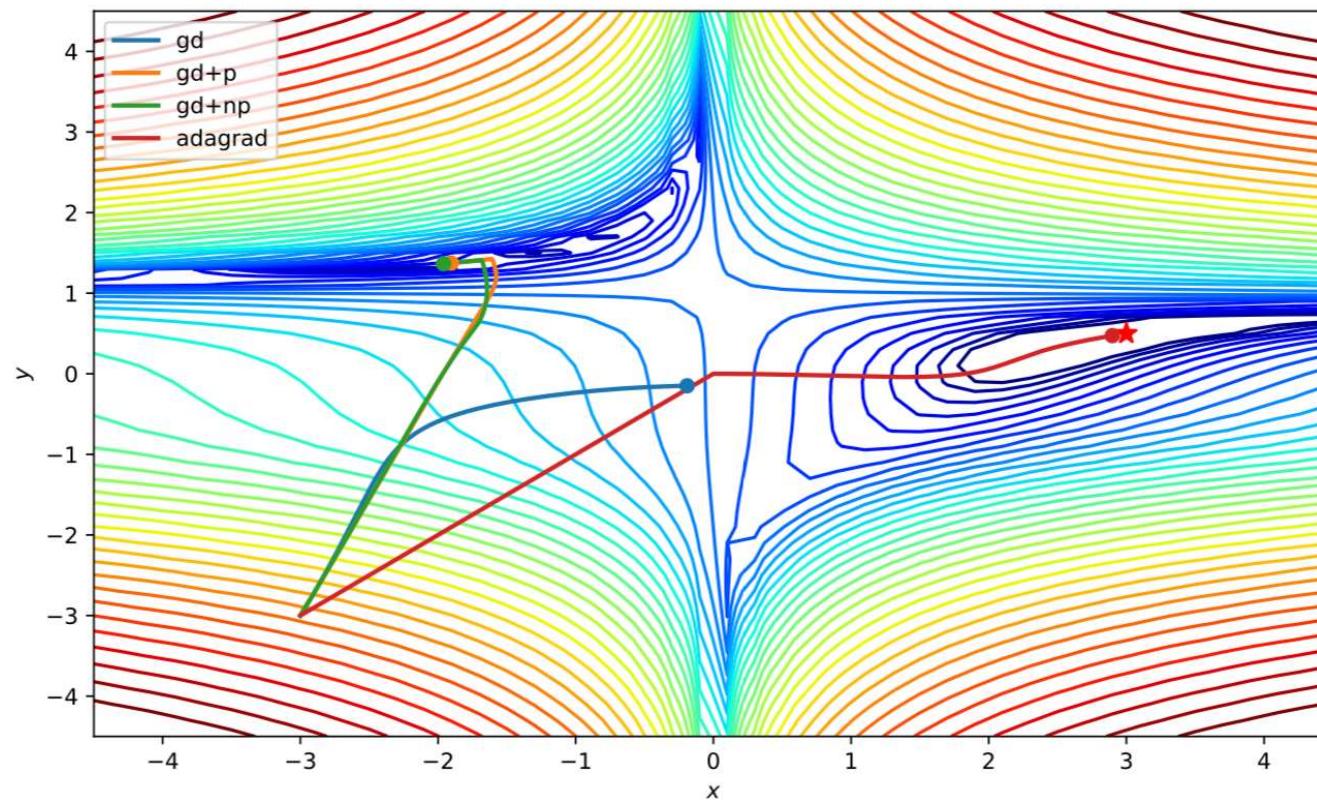




Adagrad

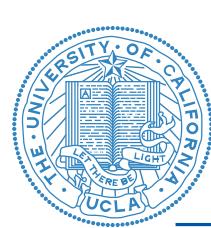
Adagrad (opt 2)

Adagrad proceeds to the global minimum.

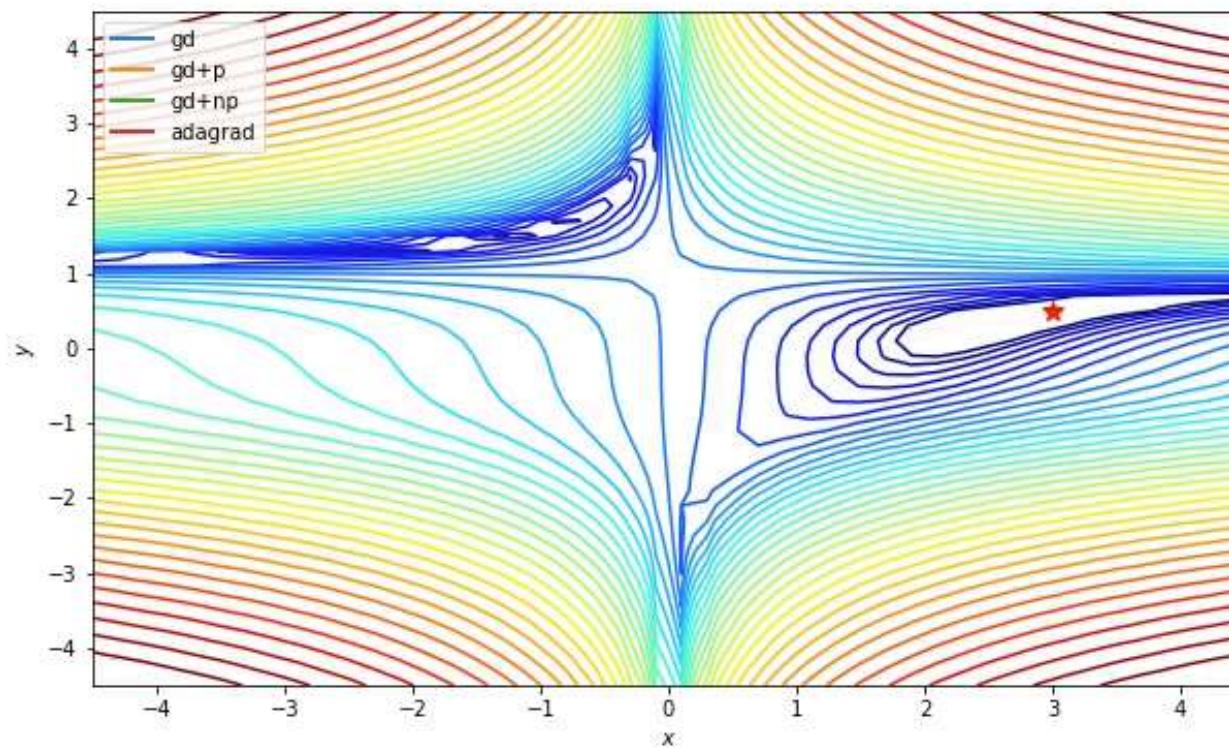


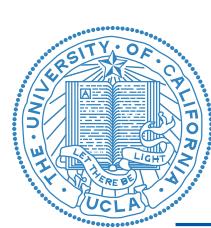
Video: http://seas.ucla.edu/~kao/opt_anim/2gd_gd+p_gd+np_adagrad.mp4





Adagrad





Adagrad

Is there a problem with adagrad?

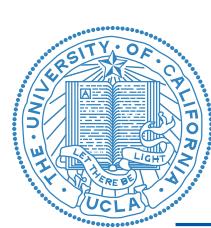
- Compute the gradient: \mathbf{g}
- Update:

$$\mathbf{a} \leftarrow \mathbf{a} + \mathbf{g} \odot \mathbf{g}$$

- Gradient step:

$$\theta \leftarrow \theta - \frac{\varepsilon}{\sqrt{\mathbf{a}} + \nu} \odot \mathbf{g}$$





RMSProp

RMSProp

RMSProp augments Adagrad by making the gradient accumulator an exponentially weighted moving average.

Initialize $\mathbf{a} = 0$ and set ν to be sufficiently small. Set β to be between 0 and 1 (typically a value like 0.99). Then, until stopping criterion is met:

- Compute the gradient: \mathbf{g}
- Update:

$$\mathbf{a} \leftarrow \beta\mathbf{a} + (1 - \beta)\mathbf{g} \odot \mathbf{g}$$

- Gradient step:

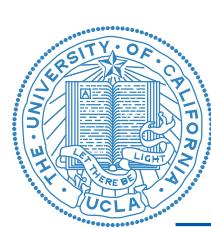
$$\theta \leftarrow \theta - \frac{\varepsilon}{\sqrt{\mathbf{a}} + \nu} \odot \mathbf{g}$$





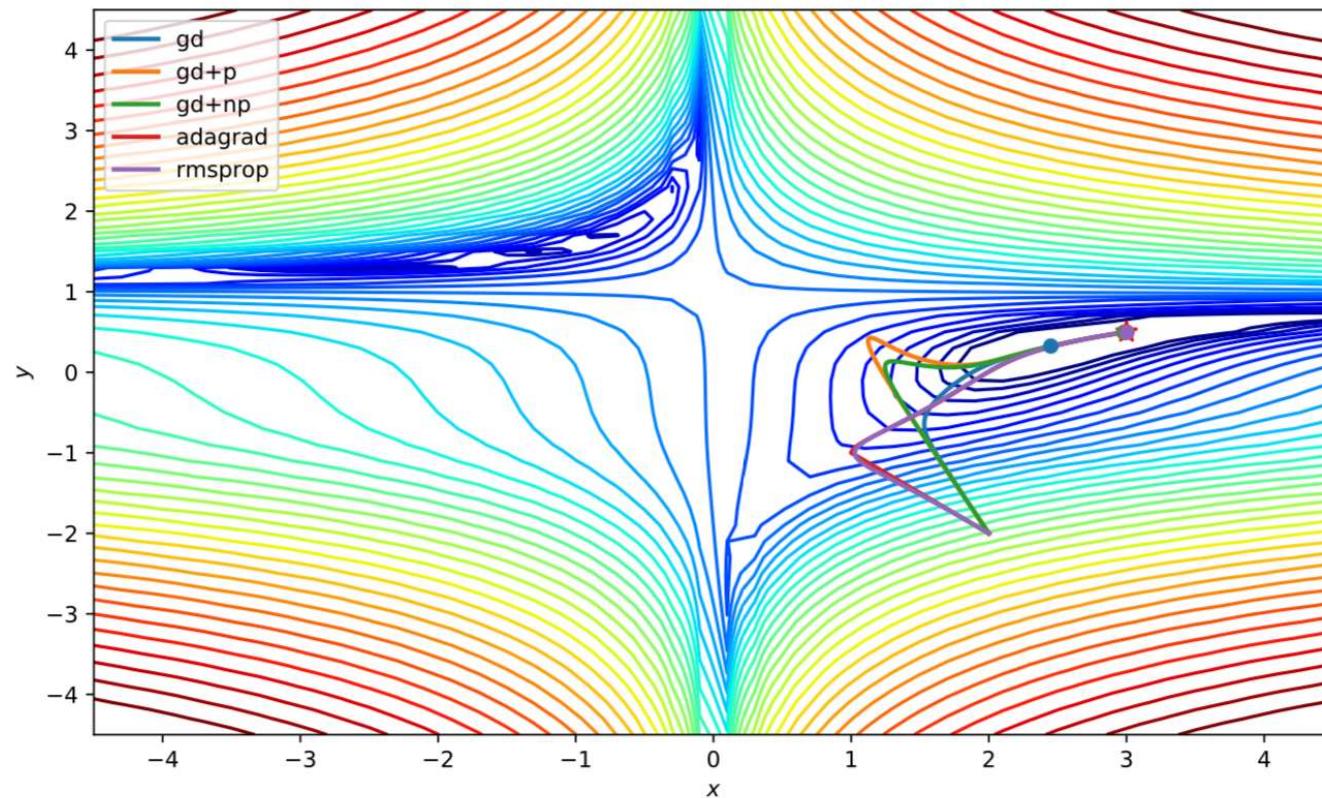
RMSProp

```
a = 0
beta = 0.99
nu = 1e-7
while last_diff > tol:
    cost, g = func(x)
    a = beta * a + (1-beta) * g * g
    x -= eps * g / (np.sqrt(a + nu))
    last_diff = np.linalg.norm(x - path[-1])
```

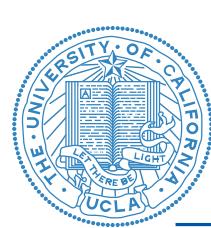


RMSProp

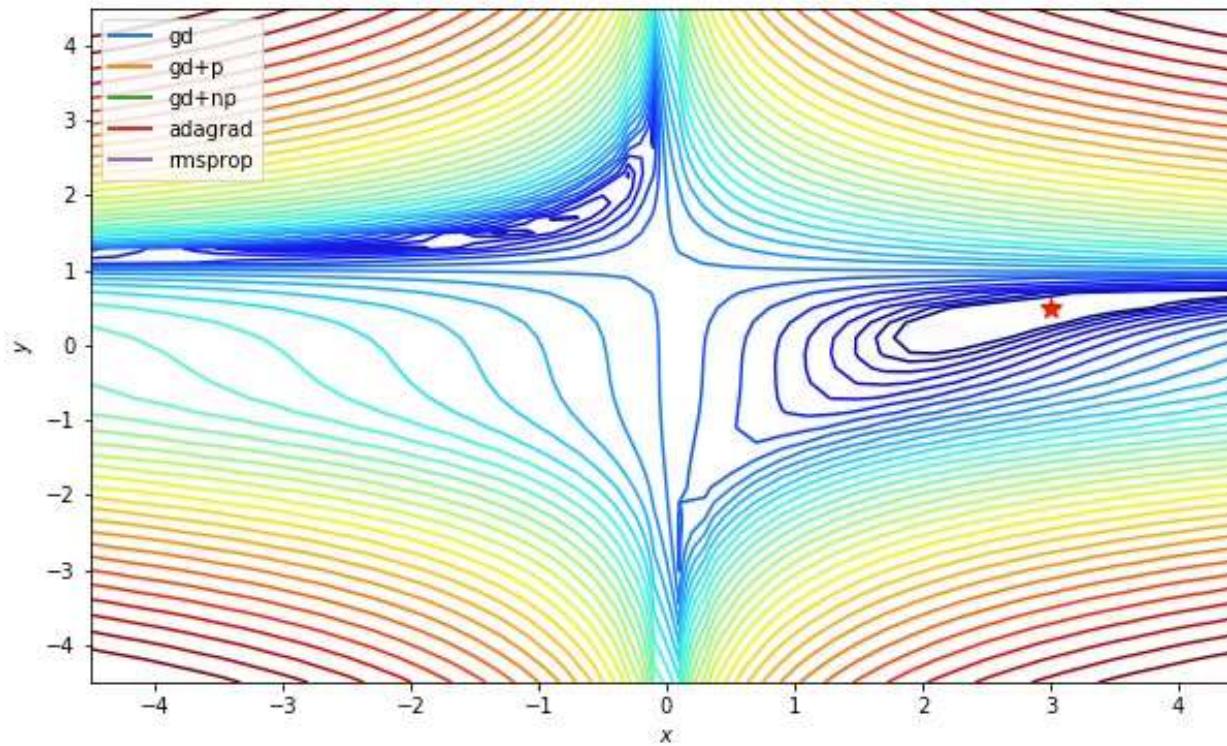
RMSProp (opt 1)

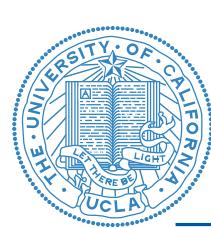


Video: http://seas.ucla.edu/~kao/opt_anim/1gd_gd+p_gd+np_adagrad_rmsprop.mp4



RMSProp

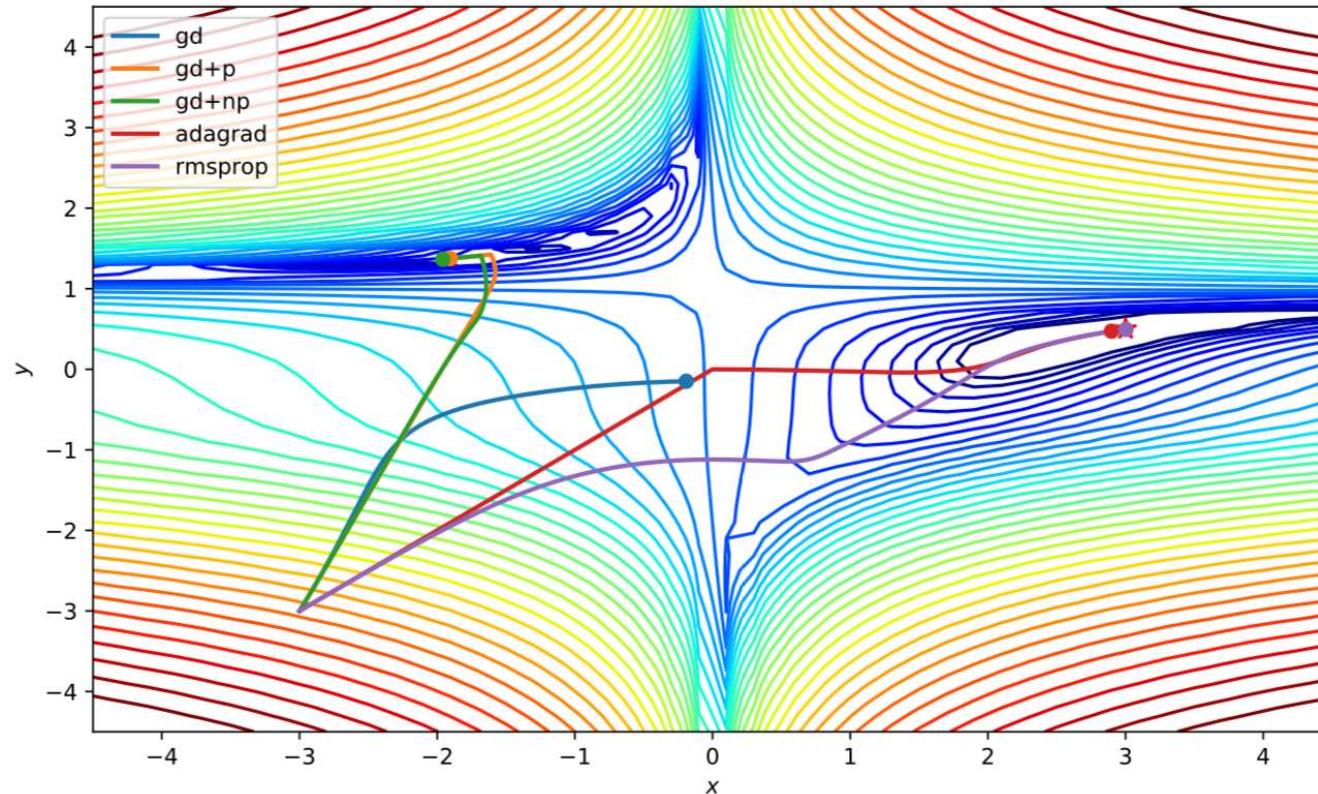




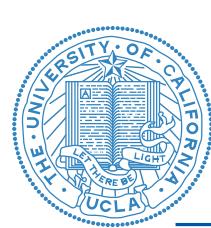
RMSProp

RMSProp (opt 2)

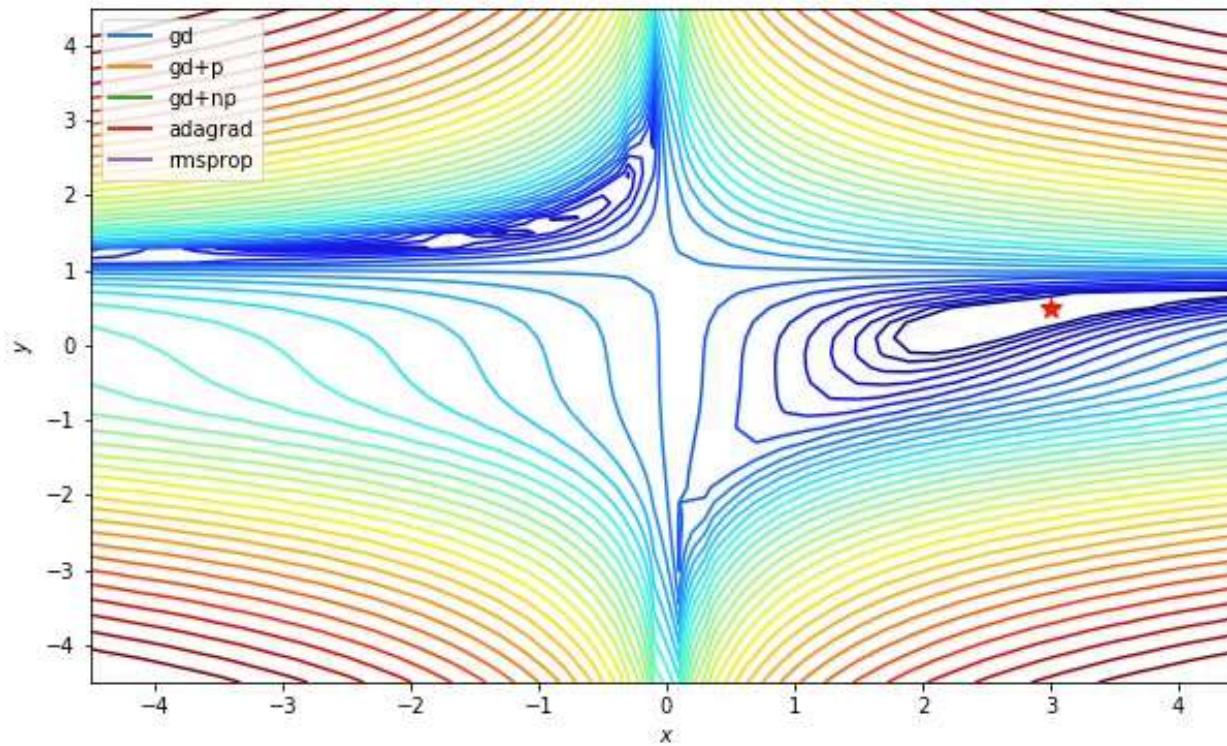
RMSProp proceeds to the global minimum, and in the video you can see it does so more quickly than Adagrad.

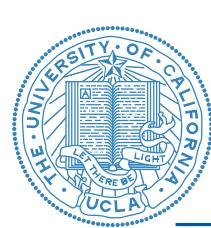


Video: http://seas.ucla.edu/~kao/opt_anim/2gd_gd+p_gd+np_adagrad_rmsprop.mp4



RMSProp





RMSProp + momentum

RMSProp with momentum

RMSProp can be combined with momentum as follows.

Initialize $\mathbf{a} = 0$. Set α, β to be between 0 and 1. Set $\nu = 1e - 7$. Until stopping criterion is met:

- Compute gradient: \mathbf{g}
- Accumulate gradient:

$$\mathbf{a} \leftarrow \beta\mathbf{a} + (1 - \beta)\mathbf{g} \odot \mathbf{g}$$

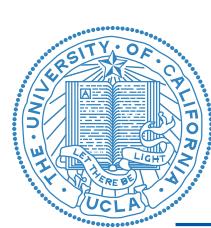
- Momentum:

$$\mathbf{v} \leftarrow \alpha\mathbf{v} - \frac{\varepsilon}{\sqrt{\mathbf{a}} + \nu} \odot \mathbf{g}$$

- Gradient step:

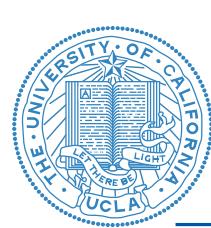
$$\theta \leftarrow \theta + \mathbf{v}$$

It is also possible to RMSProp with Nesterov momentum.



RMSProp + momentum

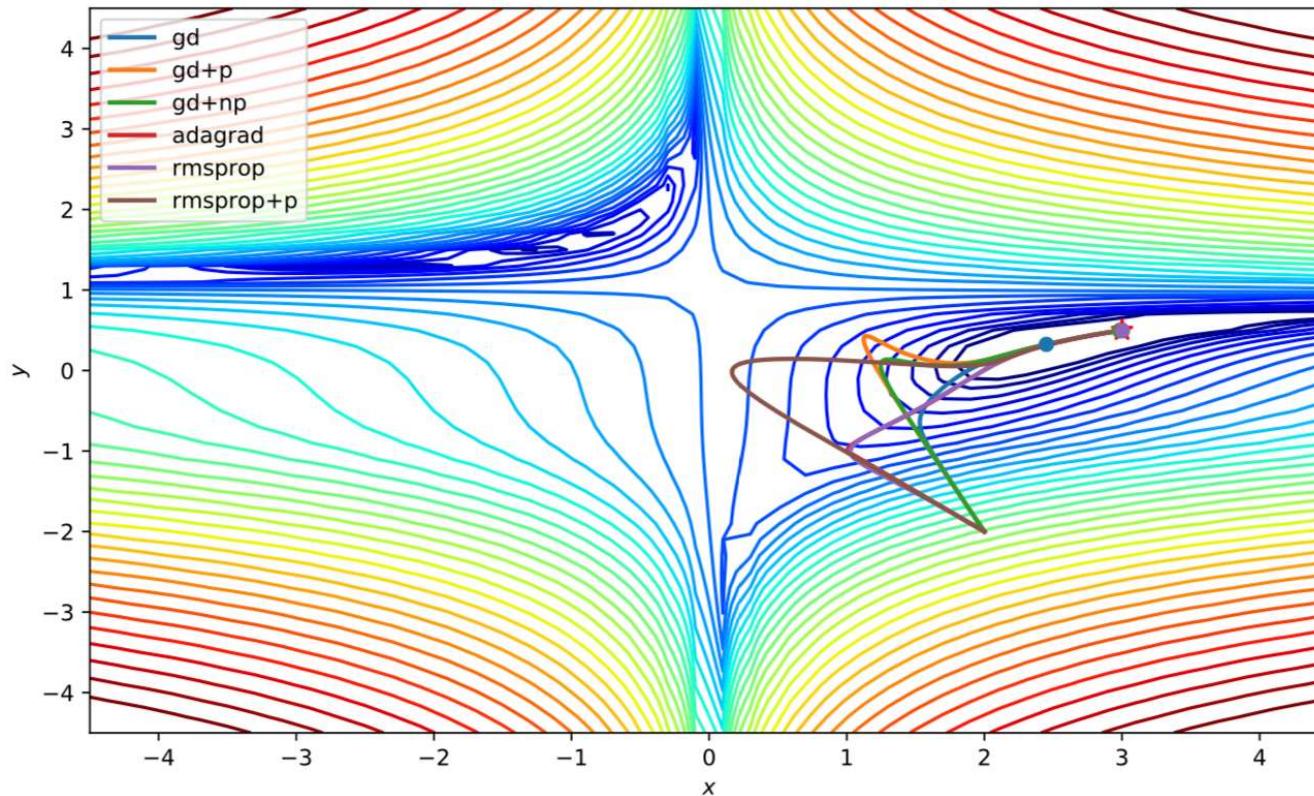
```
a = 0
p = 0
alpha = 0.9
beta = 0.99
nu = 1e-7
while last_diff > tol:
    cost, g = func(x)
    a = beta * a + (1-beta) * g * g
    p = alpha * p - eps * g / (np.sqrt(a) + nu)
    x += p
    last_diff = np.linalg.norm(x - path[-1])
```



RMSProp + momentum

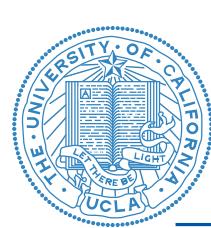
RMSProp with momentum (opt 1)

rmsprop+p denotes RMSProp with momentum. Though it makes a larger excursion out of the way, it gets to the optimum more quickly than all other optimizers. This is more apparent in the 2nd optimizer.

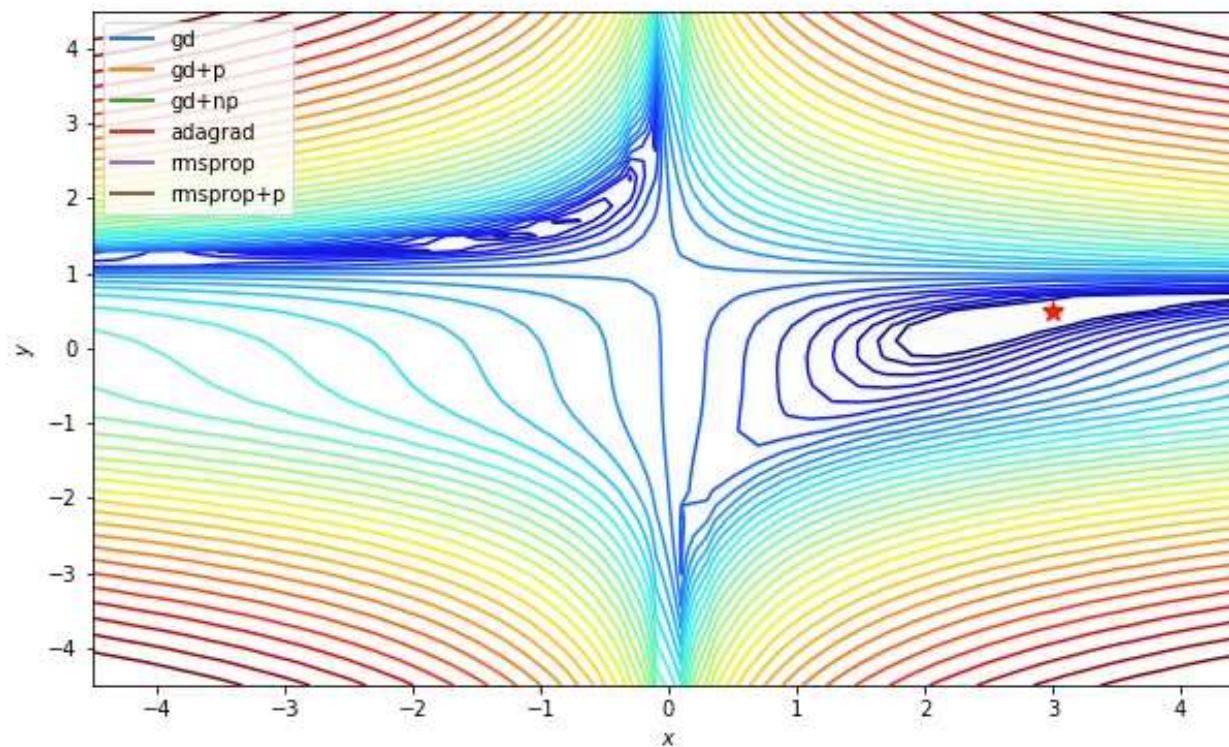


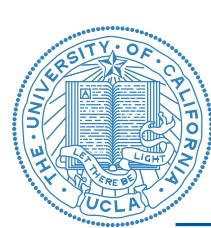
Video: http://seas.ucla.edu/~kao/opt_anim/1gd_gd+p_gd+np_adagrad_rmsprop+p.mp4

Prof J.C. Kao, UCLA ECE



RMSProp + momentum

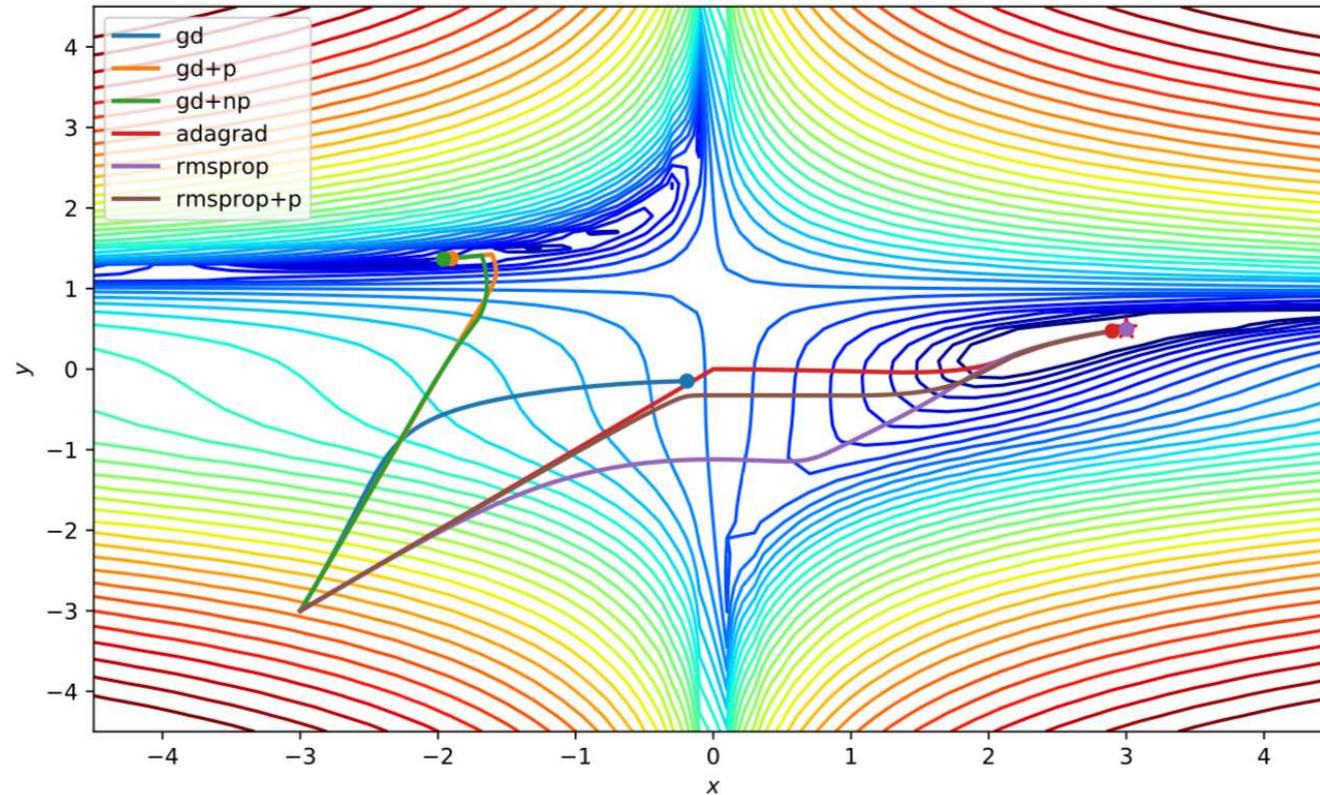




RMSProp + momentum

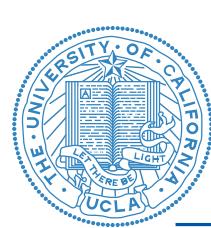
RMSProp (opt 2)

RMSProp proceeds to the global minimum, and in the video you can see it does so more quickly than Adagrad.

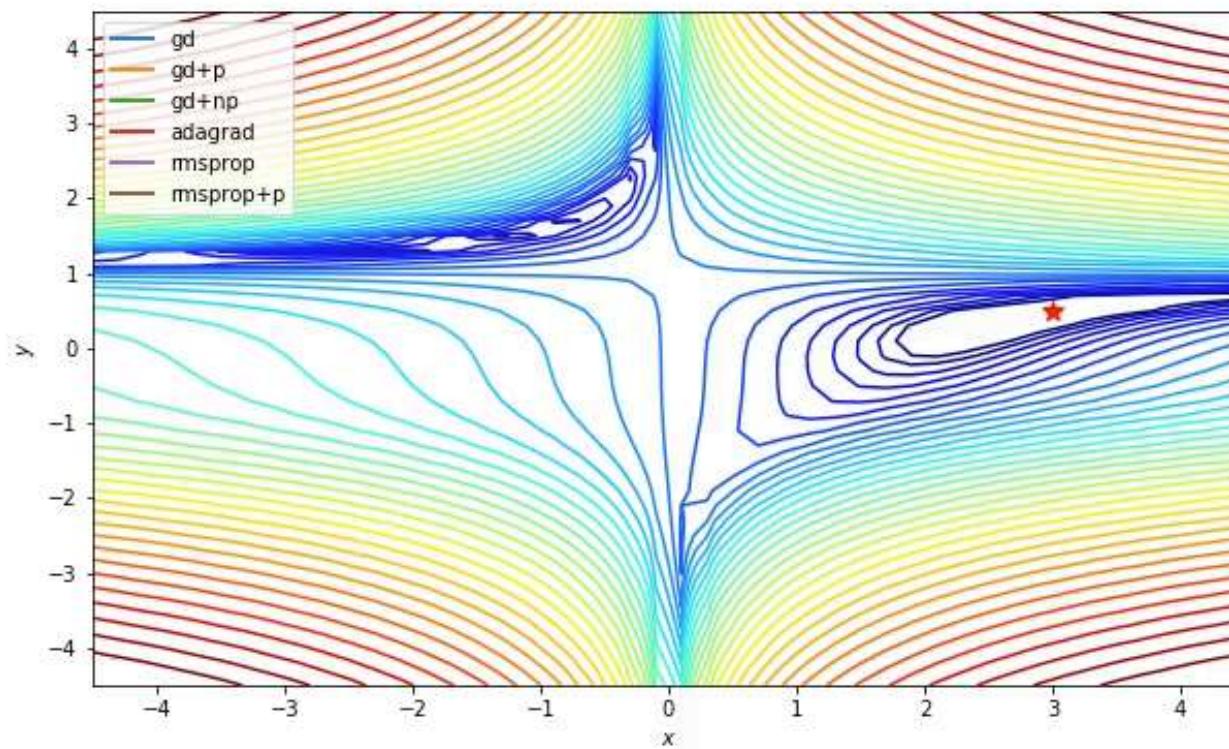


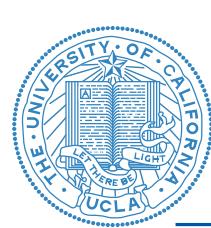
Video: http://seas.ucla.edu/~kao/opt_anim/2gd_gd+p_gd+np_adagrad_rmsprop+p.mp4





RMSProp + momentum



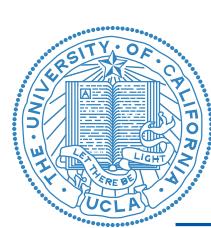


Adam

Adaptive moments without bias correction

The adaptive moments optimizer (Adam) is one of the most commonly used (and robust to e.g., hyperparameter choice) optimizers. Adam is composed of a momentum-like step, followed by an Adagrad/RMSProp-like step. For intuition, we first present Adam without a bias correction step.





Adam with no bias correction

Initialize $\mathbf{v} = 0$ as the “first moment”, and $\mathbf{a} = 0$ as the “second moment.” Set β_1 and β_2 to be between 0 and 1. (Suggested defaults are $\beta_1 = 0.9$ and $\beta_2 = 0.999$.) Initialize ν to be sufficiently small. Until stopping criterion is met:

- Compute gradient: \mathbf{g}
- First moment update (momentum-like):

$$\mathbf{v} \leftarrow \beta_1 \mathbf{v} + (1 - \beta_1) \mathbf{g}$$

- Second moment update (gradient normalization):

$$\mathbf{a} \leftarrow \beta_2 \mathbf{a} + (1 - \beta_2) \mathbf{g} \odot \mathbf{g}$$

- Gradient step:

$$\theta \leftarrow \theta - \frac{\varepsilon}{\sqrt{\mathbf{a}} + \nu} \odot \mathbf{v}$$



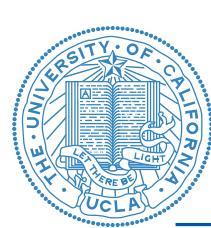


Adam

Adaptive moments (Adam)

Adam incorporates a bias correction on the moments. The intuition for the bias correction is to account for initialization; these bias corrections amplify the second moments, so that extremely large steps are not taken at the start of the optimization.





Adam

Adam (cont.)

Initialize $\mathbf{v} = 0$ as the “first moment”, and $\mathbf{a} = 0$ as the “second moment.” Set β_1 and β_2 to be between 0 and 1. (Suggested defaults are $\beta_1 = 0.9$ and $\beta_2 = 0.999$.) Initialize ν to be sufficiently small. Initialize $t = 0$. Until stopping criterion is met:

- Compute gradient: \mathbf{g}
- Time update: $t \leftarrow t + 1$
- First moment update (momentum-like):

$$\mathbf{v} \leftarrow \beta_1 \mathbf{v} + (1 - \beta_1) \mathbf{g}$$

- Second moment update (gradient normalization):

$$\mathbf{a} \leftarrow \beta_2 \mathbf{a} + (1 - \beta_2) \mathbf{g} \odot \mathbf{g}$$

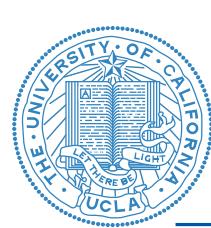
- Bias correction in moments:

$$\tilde{\mathbf{v}} = \frac{1}{1 - \beta_1^t} \mathbf{v}$$

$$\tilde{\mathbf{a}} = \frac{1}{1 - \beta_2^t} \mathbf{a}$$

- Gradient step:

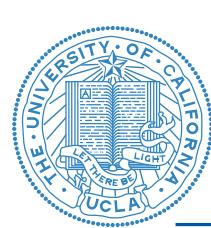
$$\theta \leftarrow \theta - \frac{\varepsilon}{\sqrt{\tilde{\mathbf{a}}} + \nu} \odot \tilde{\mathbf{v}}$$



Adam

```
a = 0
p = 0
beta1 = 0.9
beta2 = 0.99
nu = 1e-7
t = 0
while last_diff > tol:
    cost, g = func(x)
    t += 1
    p = beta1 * p + (1-beta1) * g
    a = beta2 * a + (1-beta2) * g * g
    p_u = p / (1 - beta1**t)
    a_u = a / (1 - beta2**t)
    x -= eps * p_u / (np.sqrt(a_u) + nu)
    last_diff = np.linalg.norm(x - path[-1])
```

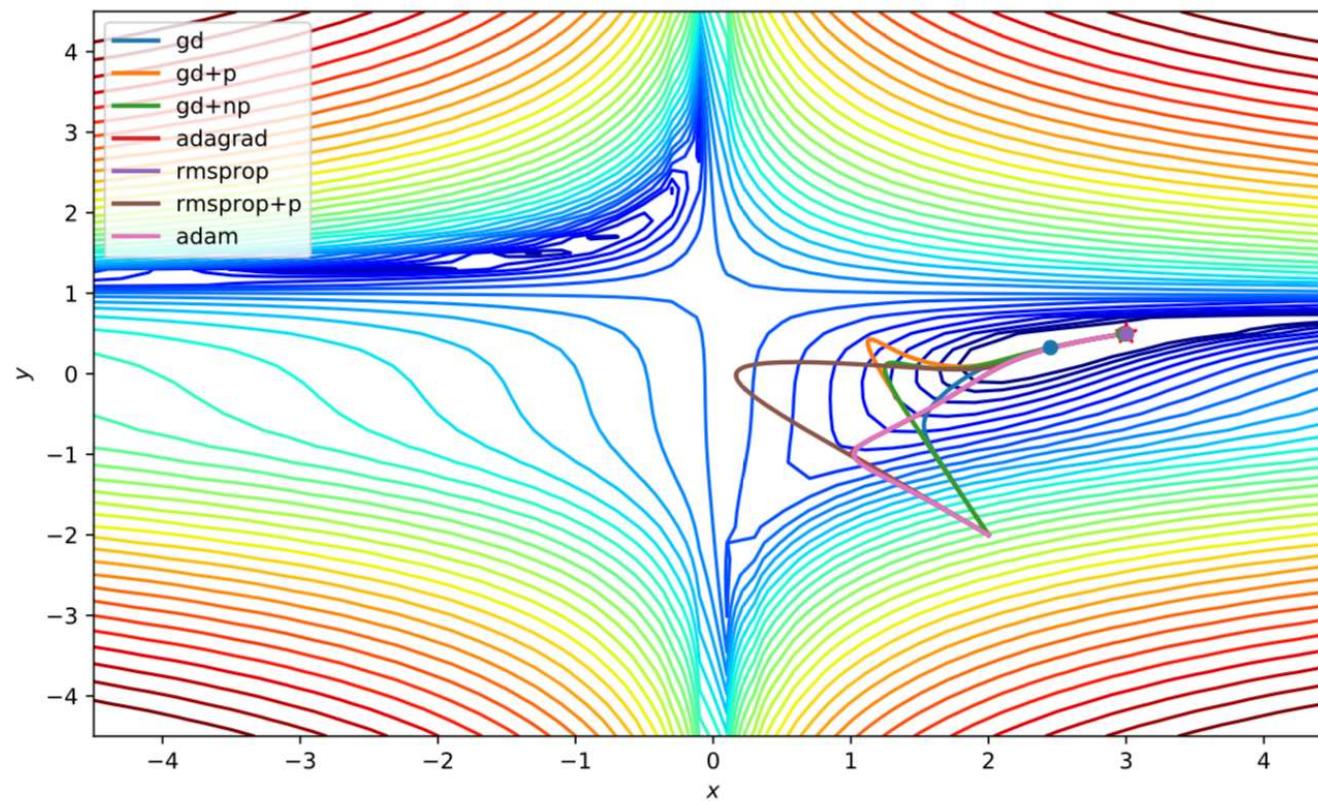




Adam

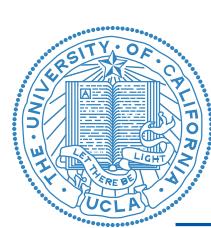
Adam (opt 1)

In both example optimizations, Adam is just slightly slower than RMSProp.

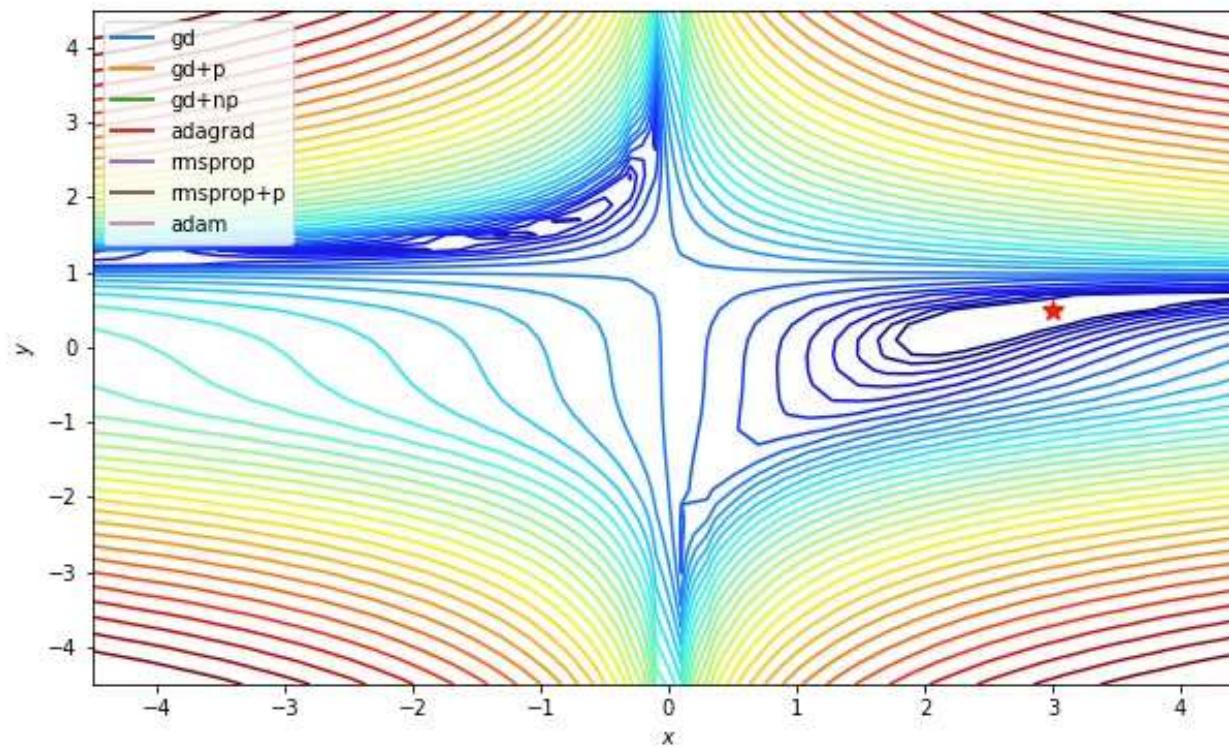


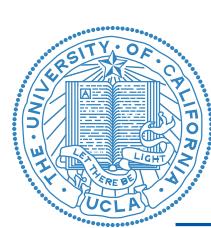
Video: http://seas.ucla.edu/~kao/opt_anim/1gd_gd+p_gd+np_adagrad_rmsprop+p_adam.mp4





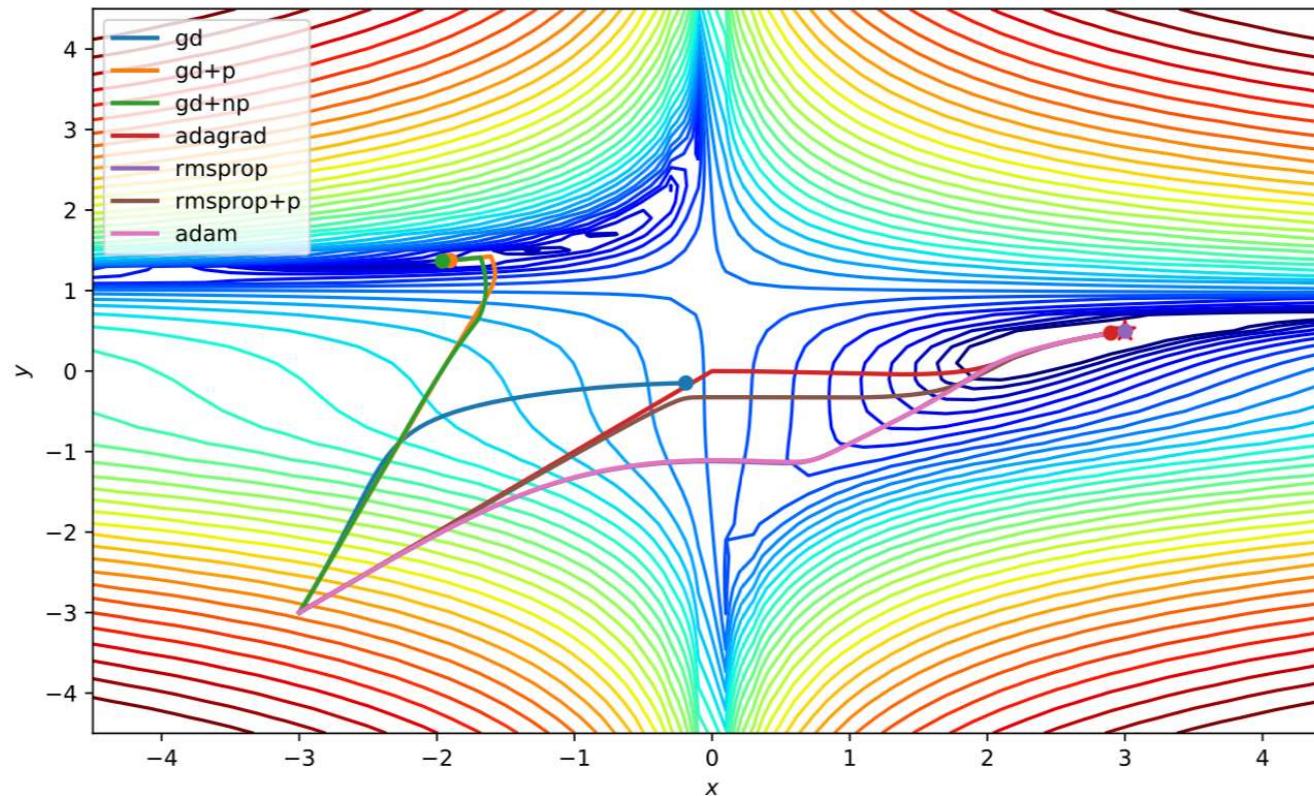
Adam





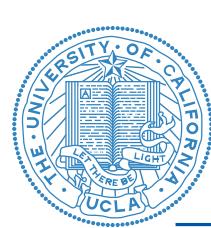
Adam

Adam (opt 2)

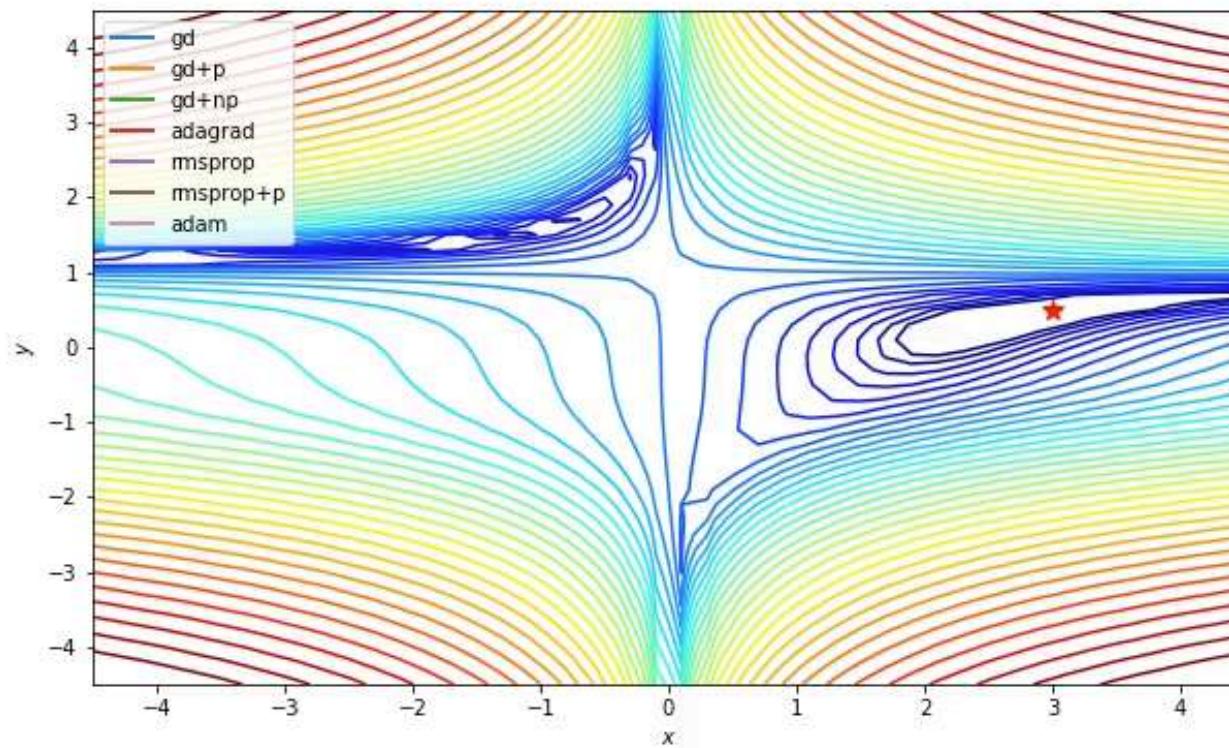


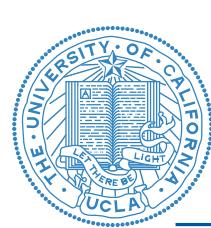
Video: http://seas.ucla.edu/~kao/opt_anim/2gd_gd+p_gd+np_adagrad_rmsprop+p_adam.mp4





Adam



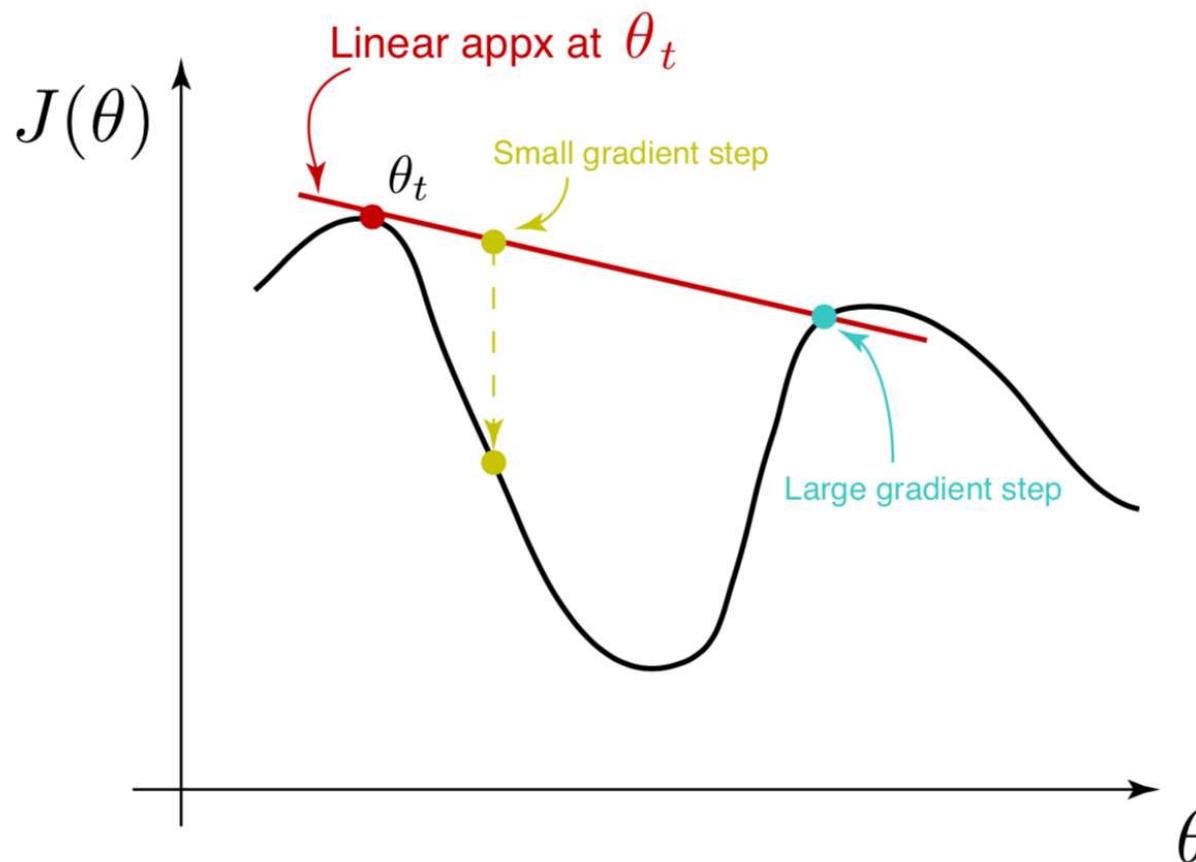


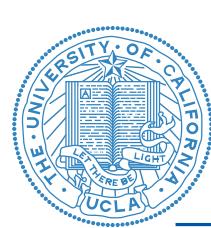
First order methods

First order vs second order methods

SGD, and optimizers used to augment the learning rate (Adagrad, RMSProp, and Adam), are all *first order* methods and have the learning rate ε as a hyperparameter.

- First-order refers to the fact that we only use the first derivative, i.e., the gradient, and take linear steps along the gradient.
- The following picture of a first order method is appropriate.

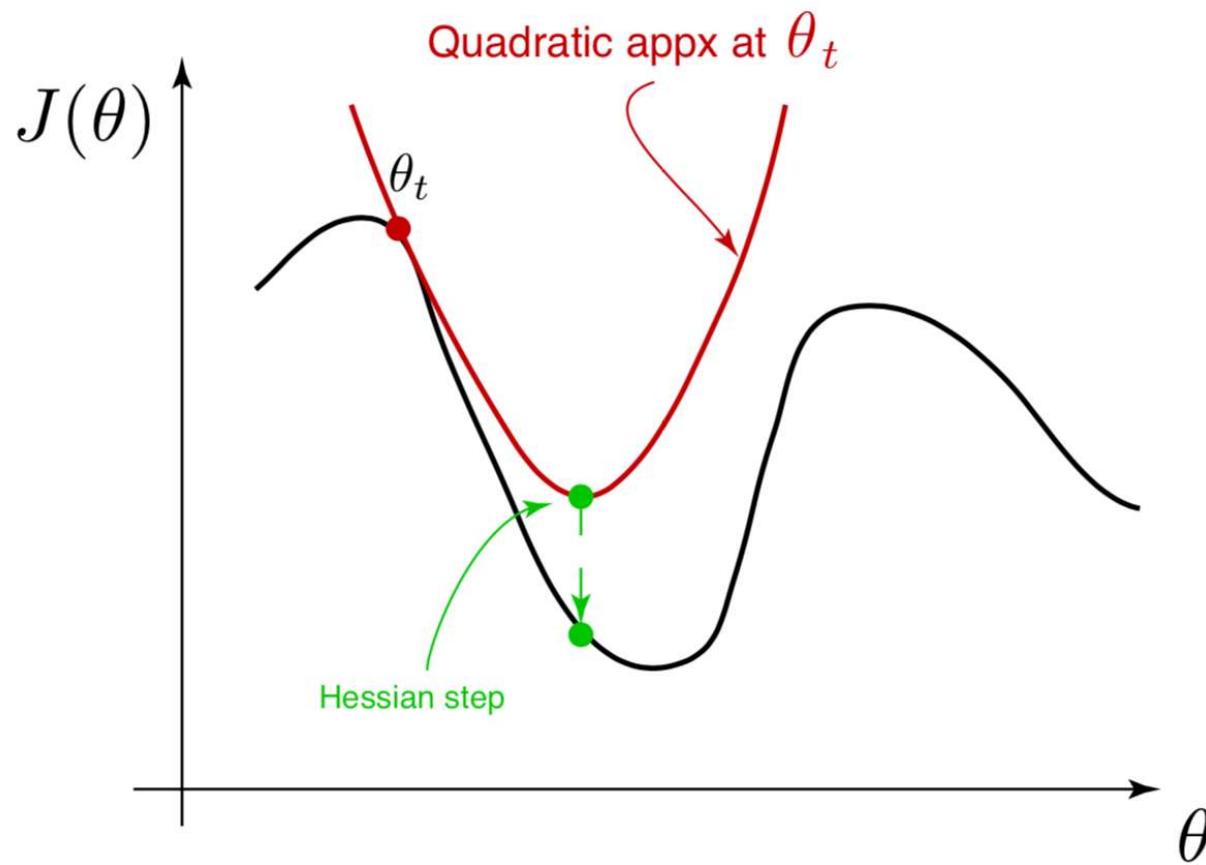


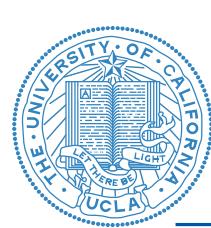


Second order methods

First order vs second order methods (cont)

It is possible to also use the *curvature* of the cost function to know how to take steps. These are called second-order methods, because they use the second derivative (or Hessian) to assess the curvature and thus take appropriate sized steps in each dimension. See following picture for intuition:





Newton's method

Second order optimization

The most widely used second order method is Newton's method. To arrive at it, consider the Taylor series expansion of $J(\theta)$ around θ_0 up to the second order terms, i.e.,

$$J(\theta) \approx J(\theta_0) + (\theta - \theta_0)^T \nabla_{\theta} J(\theta_0) + \frac{1}{2}(\theta - \theta_0)^T \mathbf{H}(\theta - \theta_0)$$

If this were just a second order function, we could minimize it by taking its derivative and setting it to zero:

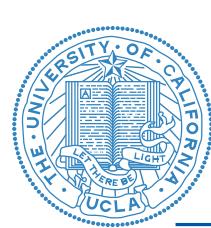
$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} J(\theta_0) + \mathbf{H}(\theta - \theta_0) \\ &= \mathbf{0}\end{aligned}$$

This results in the Newton step,

$$\theta = \theta_0 - \mathbf{H}^{-1} \nabla_{\theta} J(\theta_0)$$

If the function is quadratic, this step takes us to the minimum. If not, this approximates the function as quadratic at θ_0 and goes to the minimum of the quadratic approximation.

Does this form of step make intuitive sense?



Newton's method

Newton's method

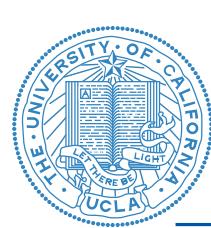
Newton's method, by using the curvature information in the Hessian, does not require a learning rate.

Until stopping criterion is met:

- Compute gradient: \mathbf{g}
- Compute Hessian: \mathbf{H}
- Gradient step:

$$\theta \leftarrow \theta - \mathbf{H}^{-1} \mathbf{g}$$





Newton's method

Newton's method (cont.)

A few notes about Newton's method:

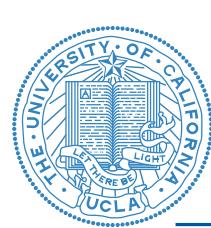
- When the Hessian has negative eigenvalues, then steps along the corresponding eigenvectors are gradient ascent steps. To counteract this, it is possible to regularize the Hessian, so that the updates become:

$$\theta \leftarrow \theta - (\mathbf{H} + \alpha \mathbf{I})^{-1} \nabla_{\theta} J(\theta_0)$$

As α becomes larger, this turns into first order gradient descent with learning rate $1/\alpha$.

- Newton's method requires, at every iteration, calculating and then inverting the Hessian. If the network has N parameters, then inverting the Hessian is $\mathcal{O}(N^3)$. This renders Newton's method impractical for many types of deep neural networks.





Quasi-Newton methods

Quasi-Newton methods

To get around the problem of having to compute and invert the Hessian, quasi-Newton methods are often used. Amongst the most well-known is the BFGS (Broyden Fletcher Goldfarb Shanno) update.

- The idea is that instead of computing and inverting the Hessian at each iteration, the inverse Hessian \mathbf{H}_0^{-1} is initialized at some value, and it is recursively updated via:

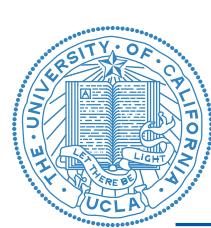
$$\mathbf{H}_k^{-1} \leftarrow \left(\mathbf{I} - \frac{\mathbf{s}\mathbf{y}^T}{\mathbf{y}^T\mathbf{s}} \right) \mathbf{H}_{k-1}^{-1} \left(\mathbf{I} - \frac{\mathbf{y}\mathbf{s}^T}{\mathbf{y}^T\mathbf{s}} \right) + \frac{\mathbf{s}\mathbf{s}^T}{\mathbf{y}^T\mathbf{s}}$$

for

$$\mathbf{s} = \theta_k - \theta_{k-1} \text{ and } \mathbf{y} = \nabla J(\theta_k) - \nabla J(\theta_{k-1})$$

- The proof of this result is beyond the scope of this class; if you'd like to learn more about this (and about optimization in general), consider taking ECE 236C.



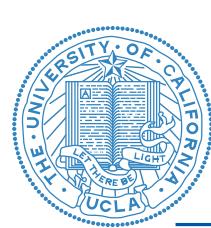


Quasi-Newton methods

Quasi-Newton methods (cont.)

- An important aspect of this update is that the inverse of any Hessian can be reconstructed from the sequence of s_k , y_k , and the initial \mathbf{H}_0^{-1} . Thus a recurrence relationship can be written to calculate $\mathbf{H}_k^{-1}\mathbf{x}$ without explicitly having to calculate \mathbf{H}_k^{-1} . However, it does require iterating over $i = 0, \dots, k$ examples.
- A way around this is to use limited memory BFGS (L-BFGS), where you calculate the inverse Hessian using just the last m examples assuming \mathbf{H}_{k-m}^{-1} is some \mathbf{H}_0^{-1} (e.g., it could be the identity matrix).
- Quasi-Newton methods usually require a full batch (or very large minibatches) since errors in estimating the inverse Hessian can result in poor steps.





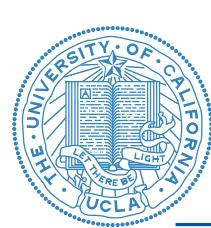
Conjugate gradients

Conjugate gradient methods

CG methods are also beyond the scope of this class, but we bring it up here in case helpful to look into further. Again, ECE 236C is recommended if you'd like to learn more about these techniques.

- CG methods find search directions that are *conjugate* with respect to the Hessian, i.e., that $\mathbf{g}_k^T \mathbf{H} \mathbf{g}_{k-1} = 0$.
- It turns out that these derivatives can be calculated iteratively through a recurrence relation.
- Implementations of “Hessian-free” CG methods have been demonstrated to converge well (e.g., Martens et al., ICML 2011).



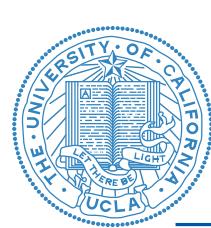


Challenges in gradient descent

- **Exploding gradients.**

Sometimes the cost function can have “cliffs” whereby small changes in the parameters can drastically change the cost function. (This usually happens if parameters are repeatedly multiplied together, as in recurrent neural networks.) Because the gradient at a cliff is large, an update can result in going to a completely different parameter space. This can be ameliorated via gradient clipping, which upper bounds the maximum gradient norm.





Challenges in gradient descent

- **Vanishing gradients.**

Like in exploding gradients, repeated multiplication of a matrix \mathbf{W} can cause vanishing gradients. Say that each time step can be thought of as a layer of a feedforward network where each layer has connectivity \mathbf{W} to the next layer. By layer t , there have been \mathbf{W}^t multiplications. If $\mathbf{W} = \mathbf{U}\Lambda\mathbf{U}^{-1}$ is its eigendecomposition, then $\mathbf{W}^t = \mathbf{U}\Lambda^t\mathbf{U}^{-1}$, and hence the gradient along eigenvector \mathbf{u}_i is shrunk (or grown) by the factor λ_i^t . Architectural decisions, as well as appropriate regularization, can deal with vanishing gradients.